

# Android IPC Mechanism

Jim Huang (黃敬群)

Developer, 0xlab

[jserv@0xlab.org](mailto:jserv@0xlab.org)

March 19, 2012 / 南台科大

# Rights to copy

© Copyright 2012 0xlab

<http://0xlab.org/>

[contact@0xlab.org](mailto:contact@0xlab.org)



## **Attribution – ShareAlike 3.0**

### **You are free**

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Corrections, suggestions, contributions and translations are welcome!

Latest update: Mar 21, 2012

### **Under the following conditions**

 **Attribution.** You must give the original author credit.

 **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.

- For any reuse or distribution, you must make clear to others the license terms of this work.
- Any of these conditions can be waived if you get permission from the copyright holder.

**Your fair use and other rights are in no way affected by the above.**

License text: <http://creativecommons.org/licenses/by-sa/3.0/legalcode>



# Agenda

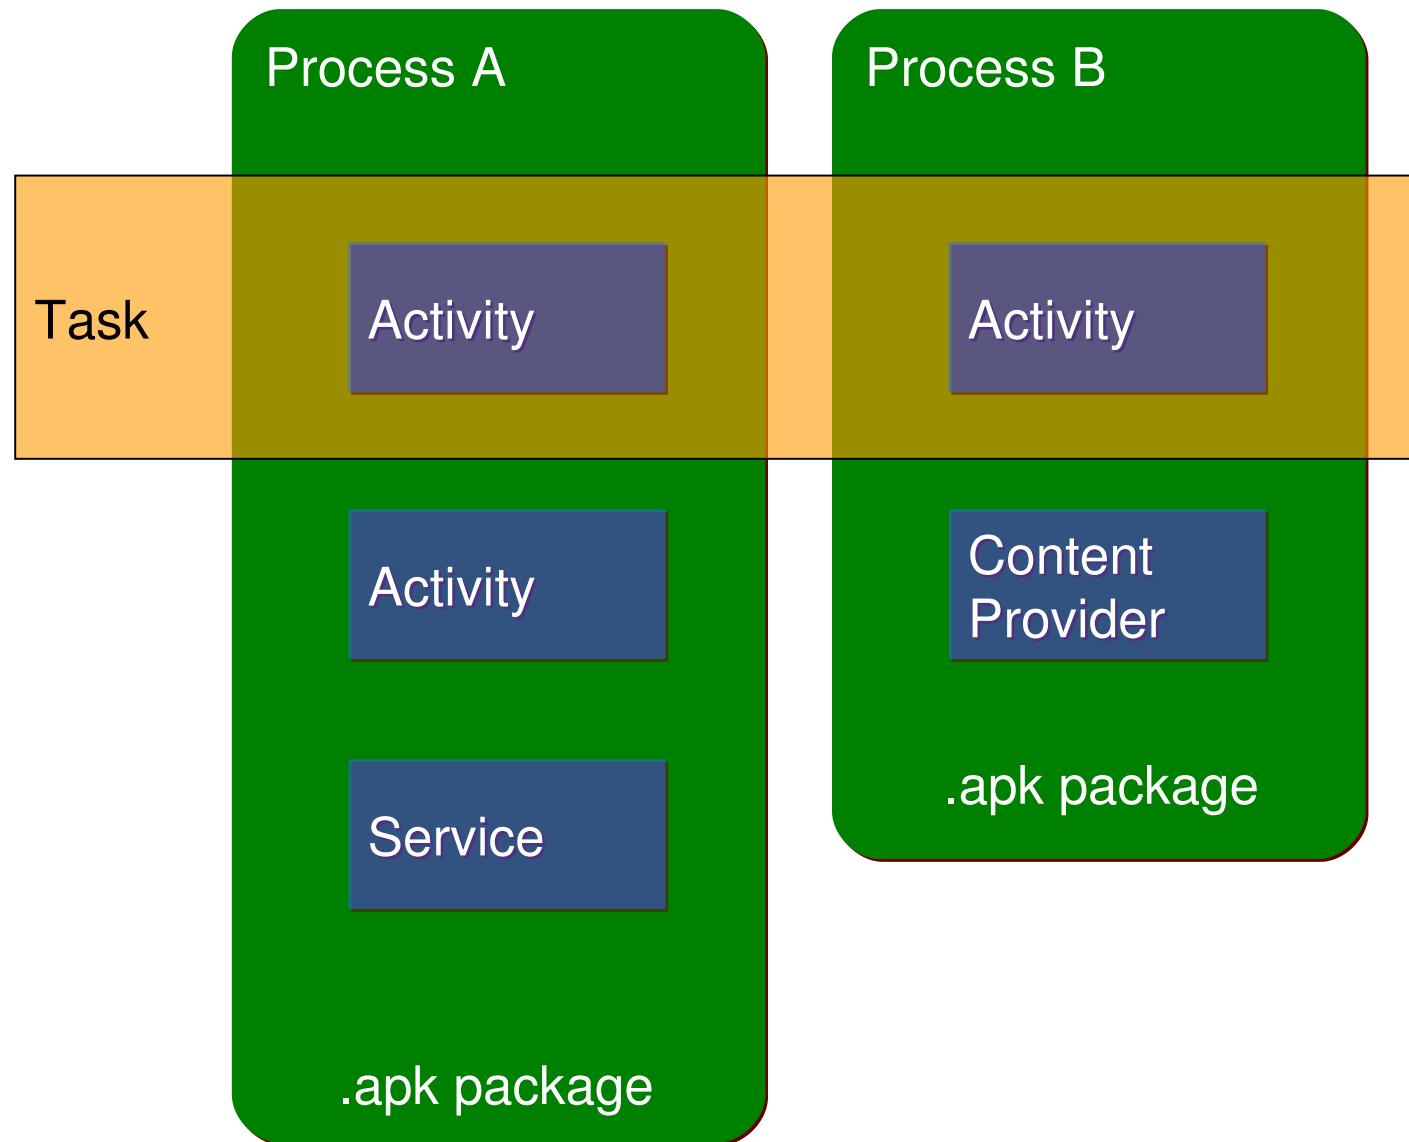
- (0) IPC: The heart of Android
- (1) Design Patterns
- (2) Binder IPC Internals
- (3) Use case: Graphics



# Binder IPC: The heart of Android

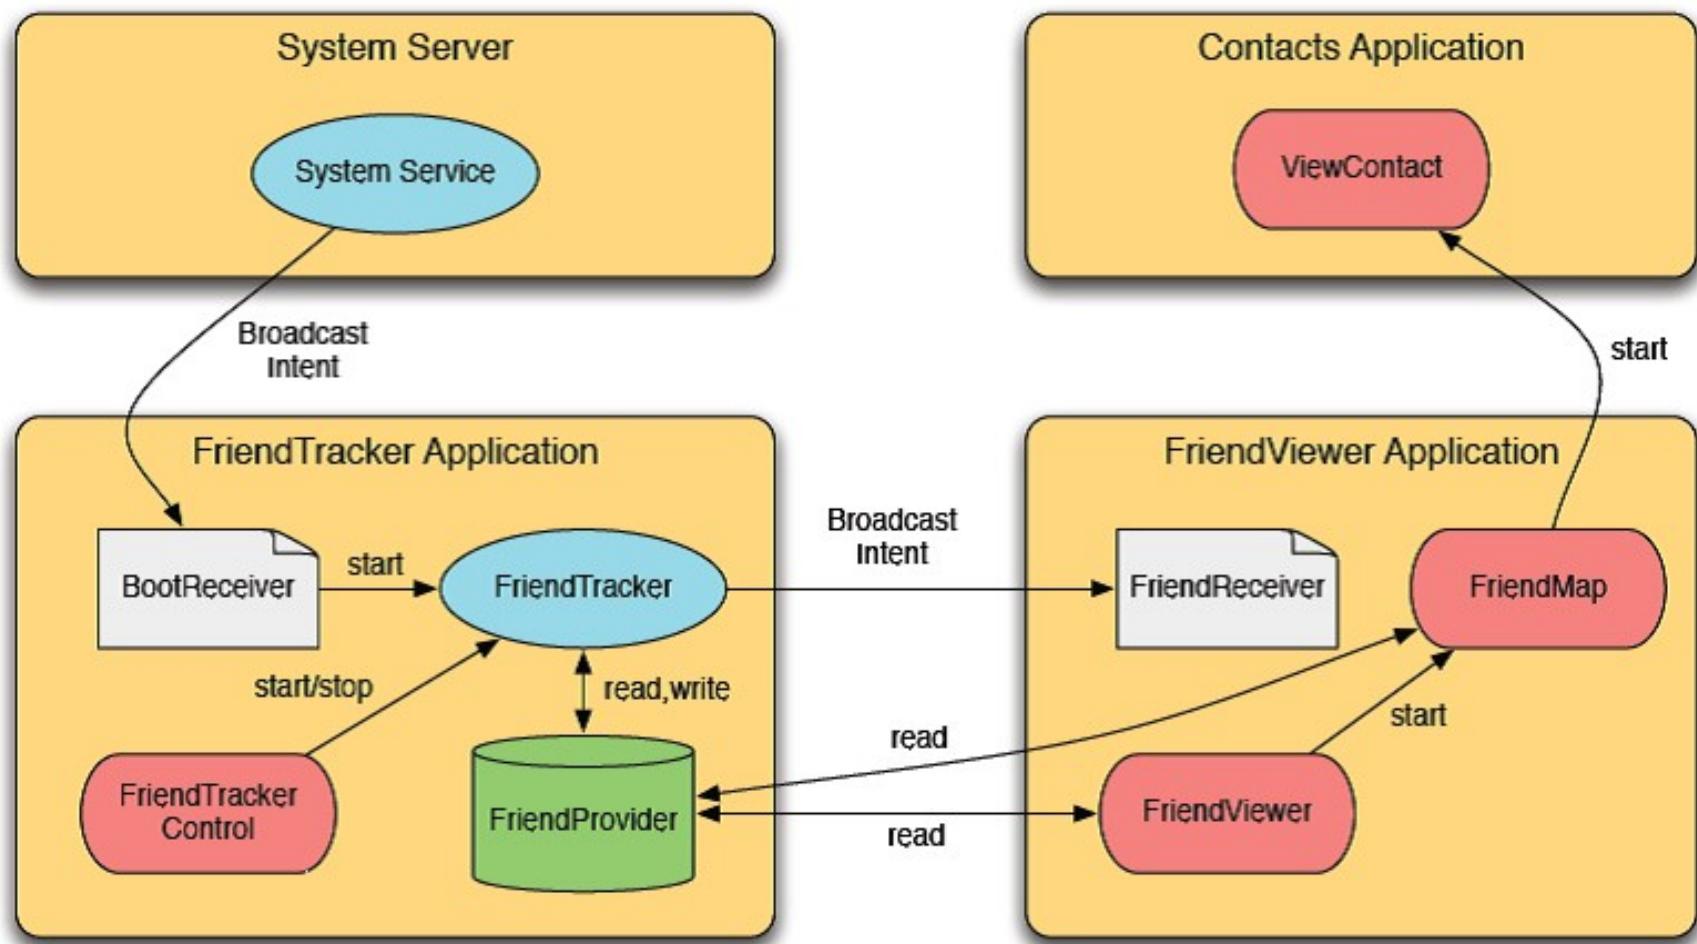


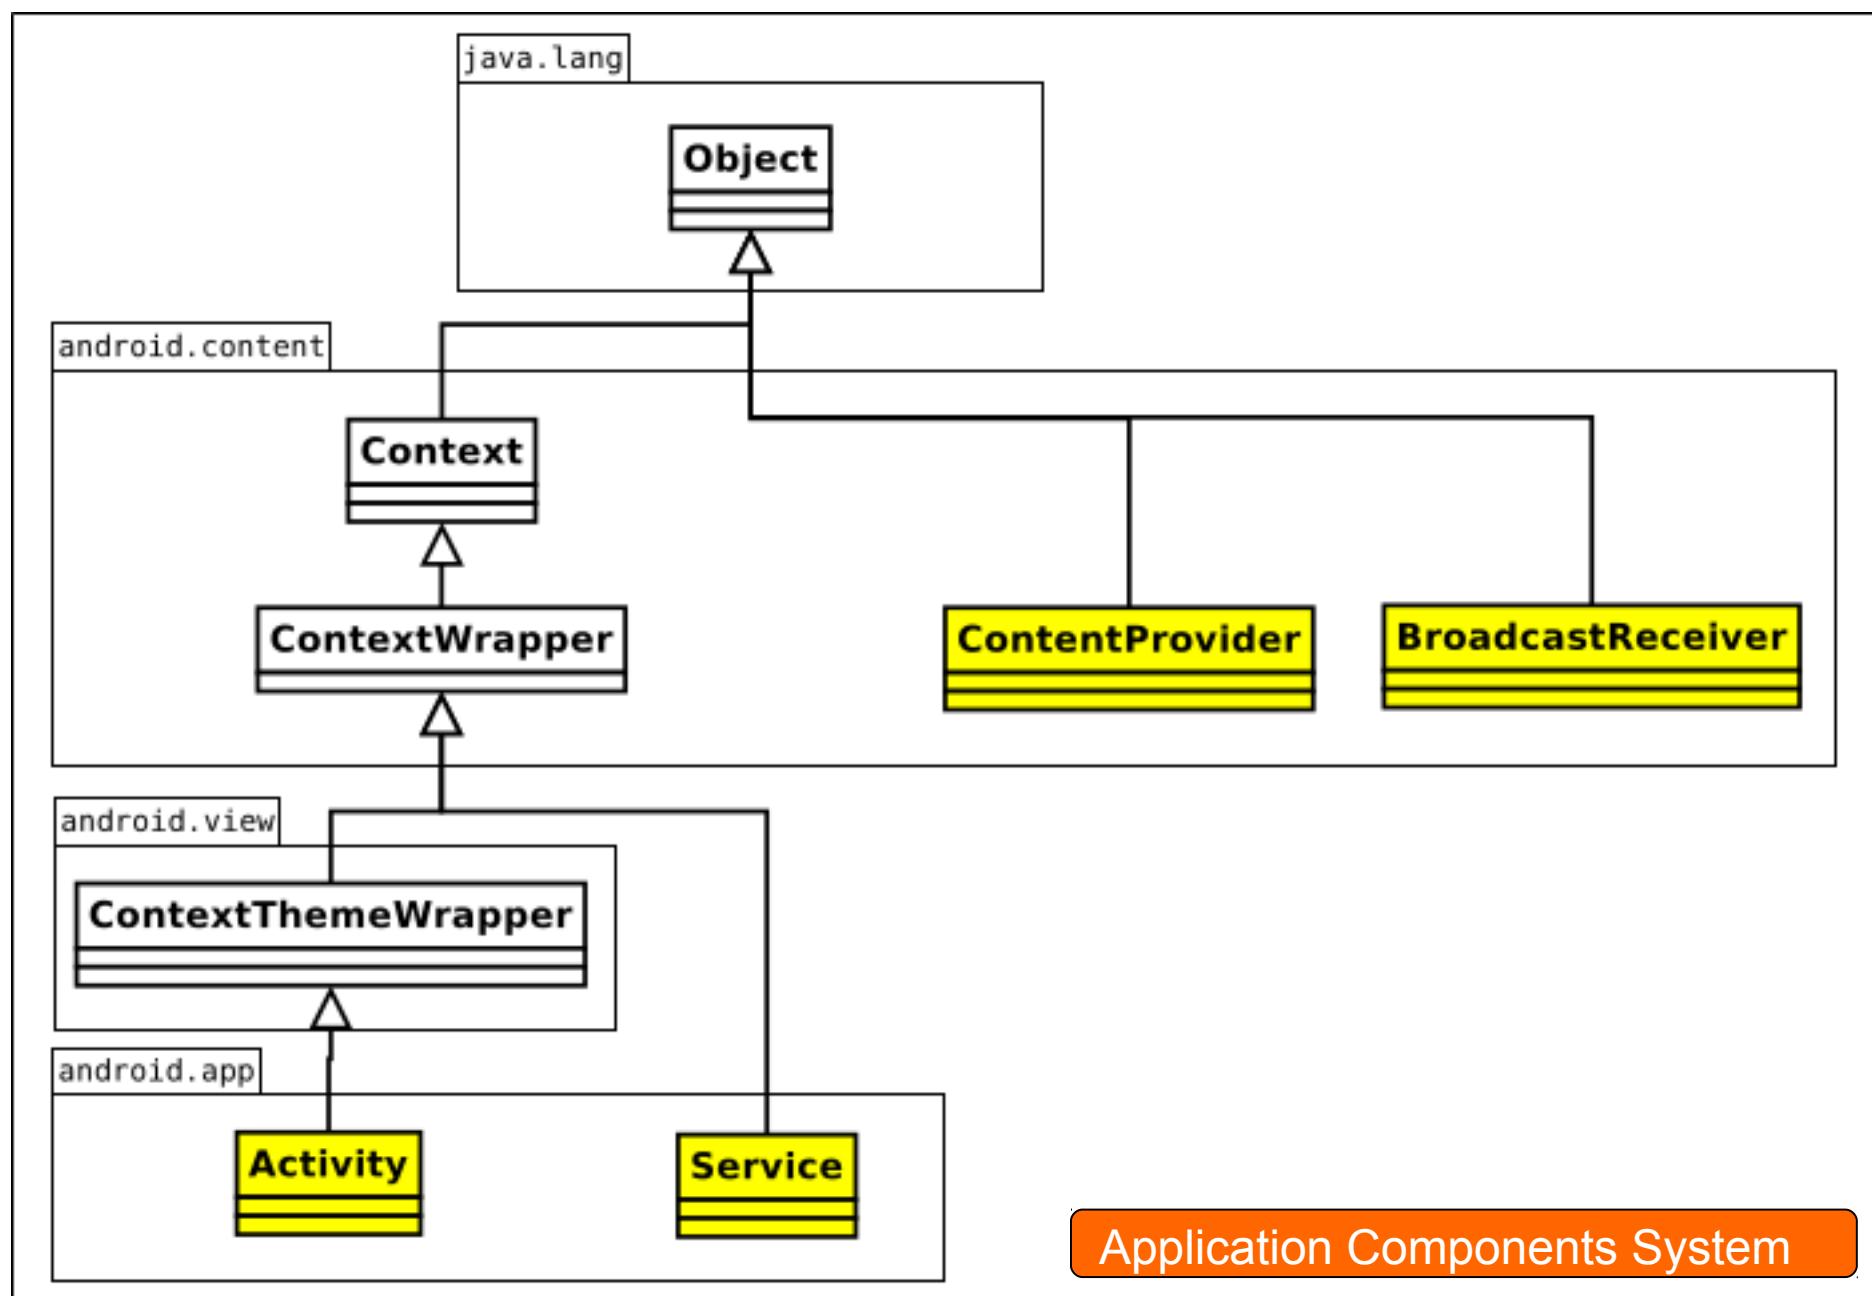
# Android Tasks



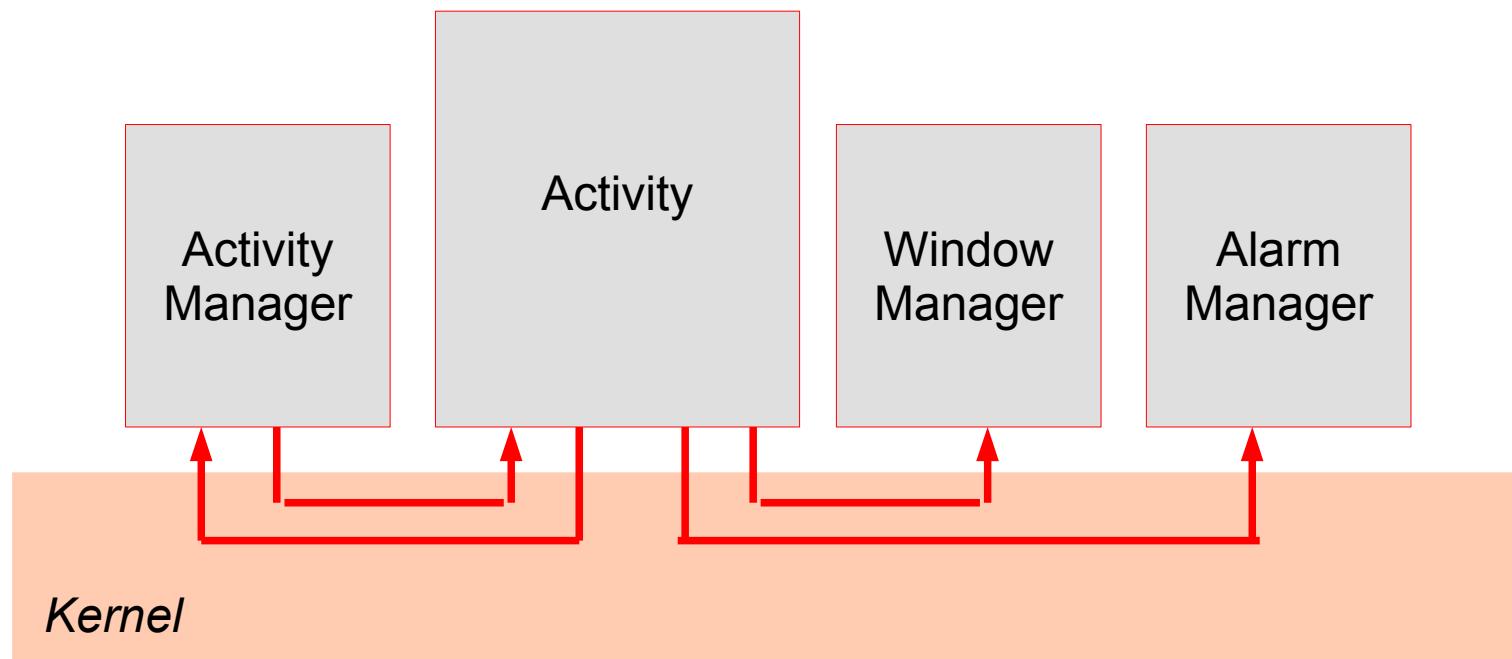
- Different component types
  - Activity
  - Service
  - Content Provider
  - Broadcast Receiver

# Component View





# IPC = Inter-Process Communication



# Why IPC?

- Each process has its own address space
- Provides data isolation
- Prevents harmful direct interaction between two different processes
  - Sometimes, communication between processes is required for modularization



# IPC Mechanisms

- In GNU/Linux
  - Signal
  - Pipe
  - Socket
  - Semaphore
  - Message queue
  - Shared memory
- In Android
  - Binder: lightweight RPC (Remote Procedure Communication) mechanism



# Binder History

- Developed under the name OpenBinder by Palm Inc. under the leadership of Dianne Hackborn
- Android Binder is the customized re-implementation of OpenBinder, which provides bindings to functions and data from one execution environment to another



# Background Problems

- Applications and Services may run in separate processes but must communicate and share data
- IPC can introduce significant processing overhead and security holes



# Binder: Android's Solution

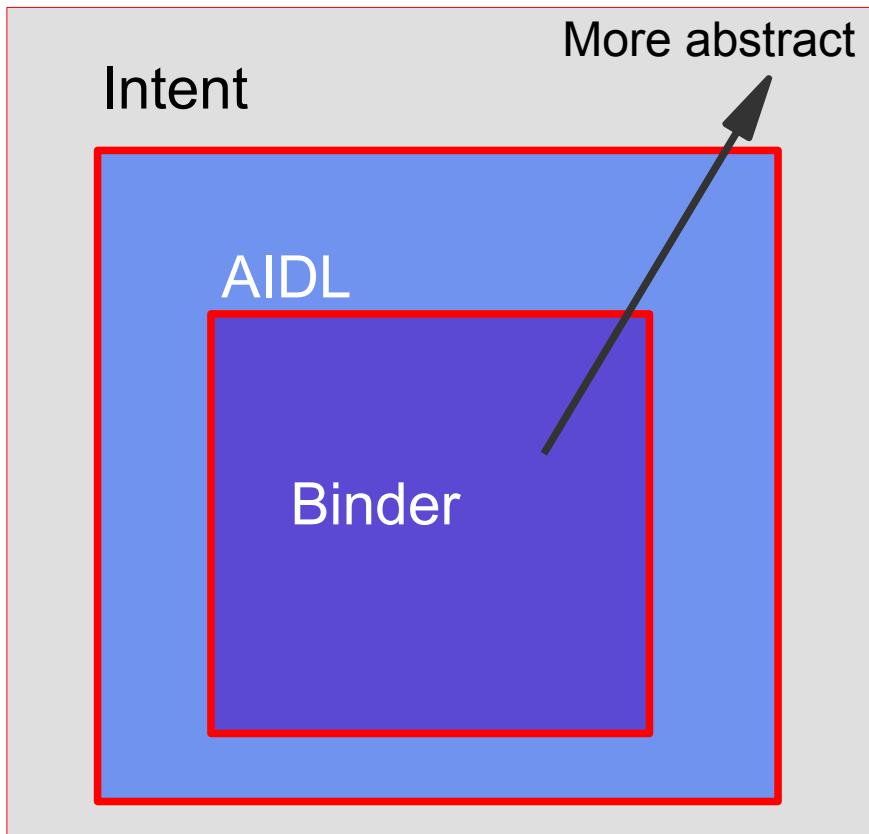
- Driver to facilitate inter-process communication
- High performance through shared memory
- Per-process thread pool for processing requests
- Reference counting, and mapping of object references across processes
- Synchronous calls between processes

“In the Android platform, the binder is used for nearly everything that happens across processes in the core platform. ” – Dianne Hackborn

<https://lkml.org/lkml/2009/6/25/3>



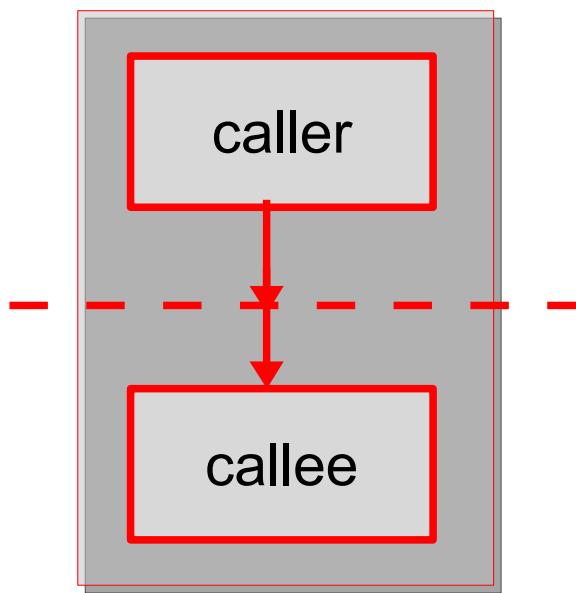
# IPC Abstraction



- Intent
  - The highest level abstraction
- Inter process method invocation
  - **AIDL**: Android Interface Definition Language
- binder: kernel driver
- ashmem: shared memory



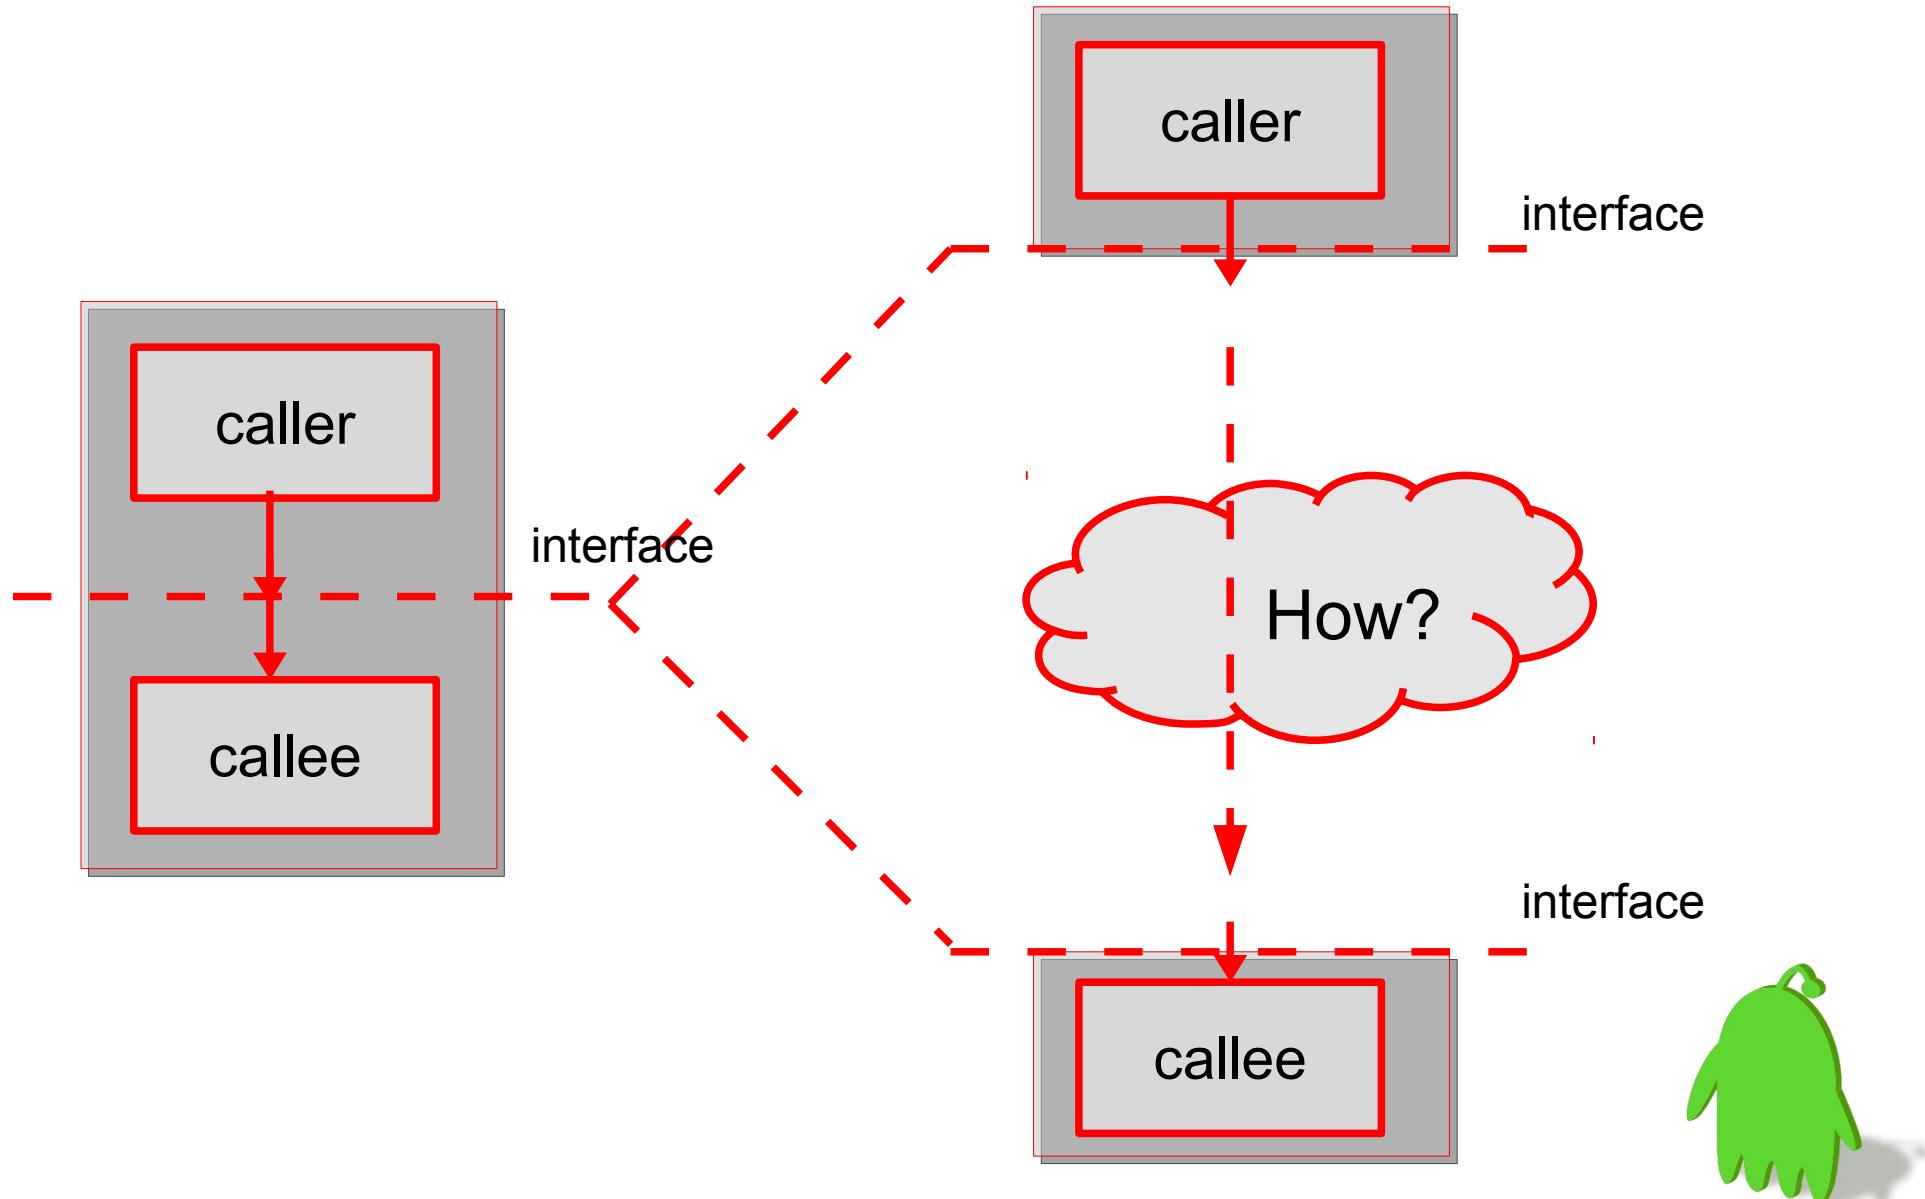
# Method invocation



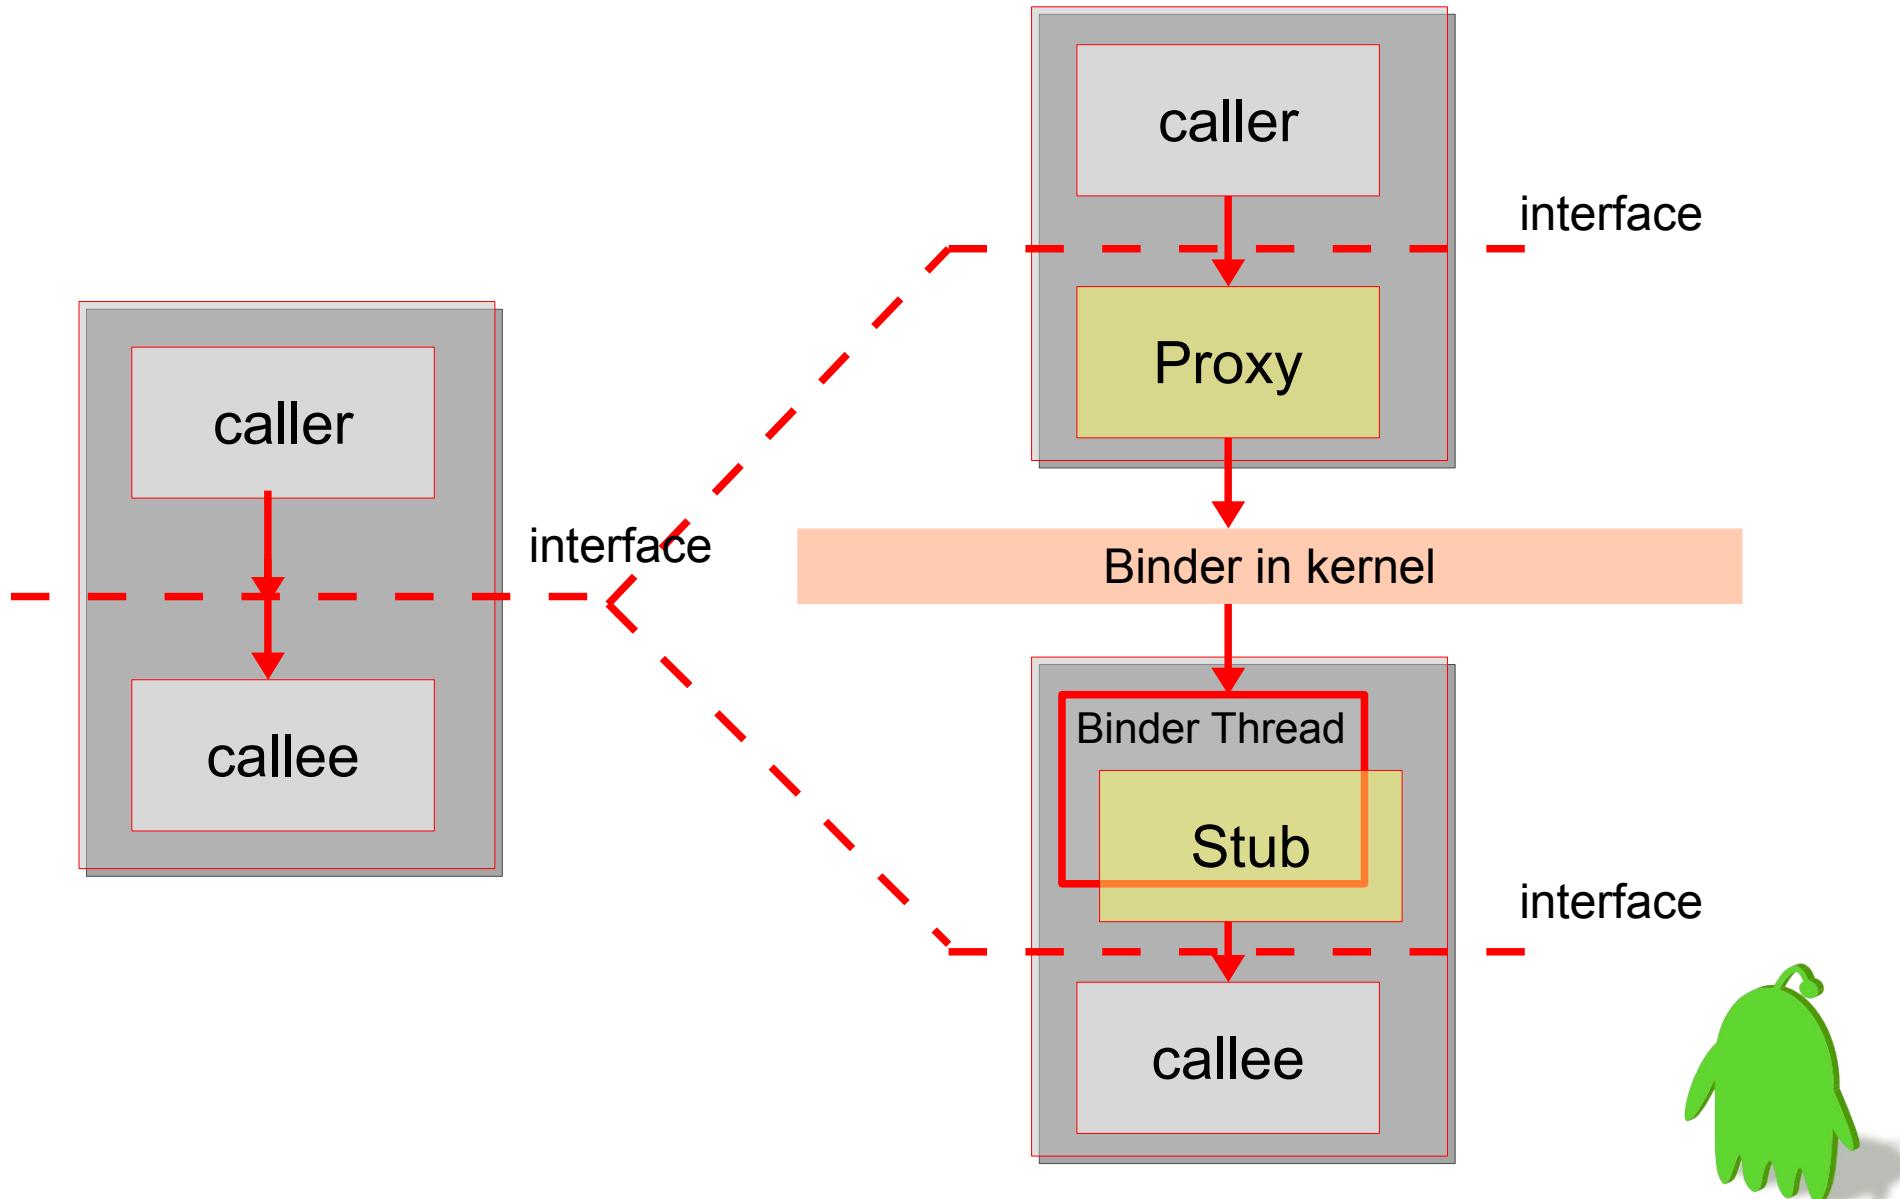
In the same process



# Inter-process method invocation



# Inter-process method invocation

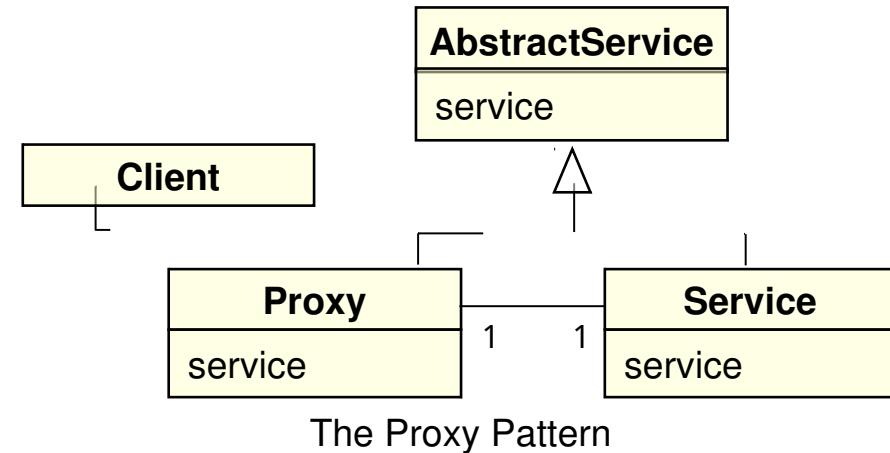


# Design Patterns



# Pattern

- Abstracts and names a recurring design structure
- Comprises class and/or object
  - Dependencies
  - Structures
  - Interactions
  - Conventions
- Specifies the design structure explicitly
- is distilled from actual design experience
- Android itself follows object oriented design



# Design Patterns used in Binder

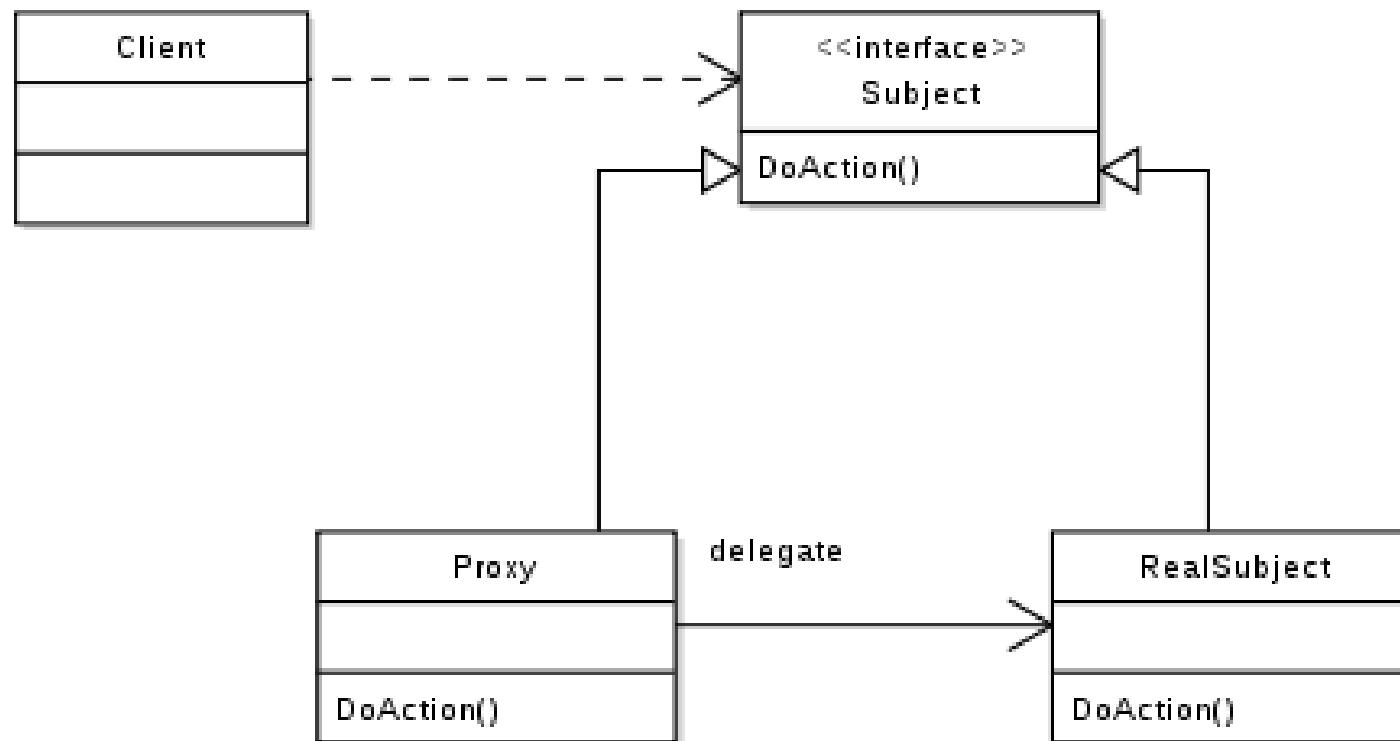
(incomplete)

- Proxy Pattern
- Mediator Pattern
- Bridge Pattern



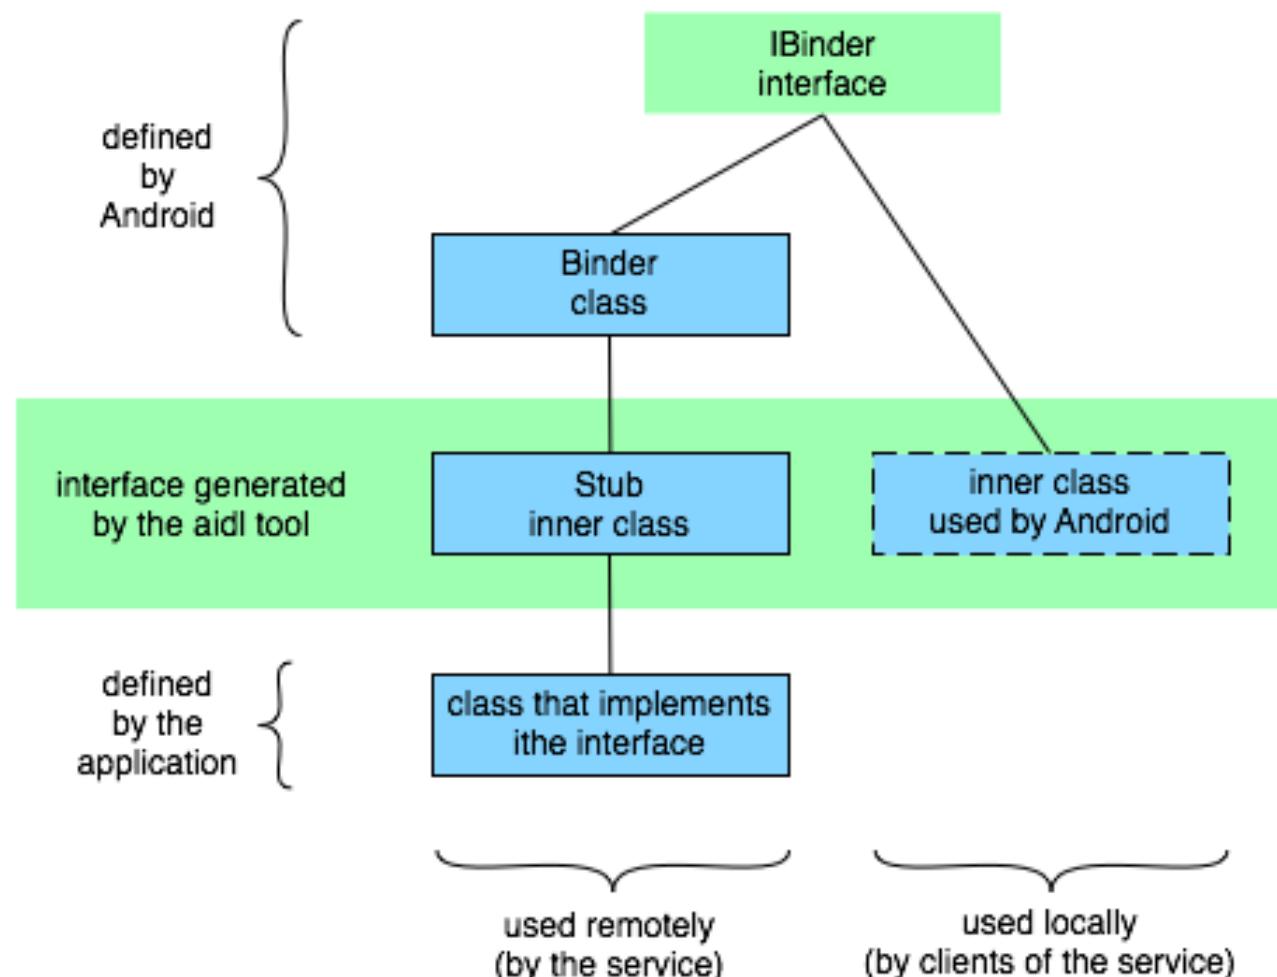
# Proxy Pattern

- The proxy could interface to anything: a network connection, a large object in memory, a file, or some other resource that is expensive or impossible to duplicate.



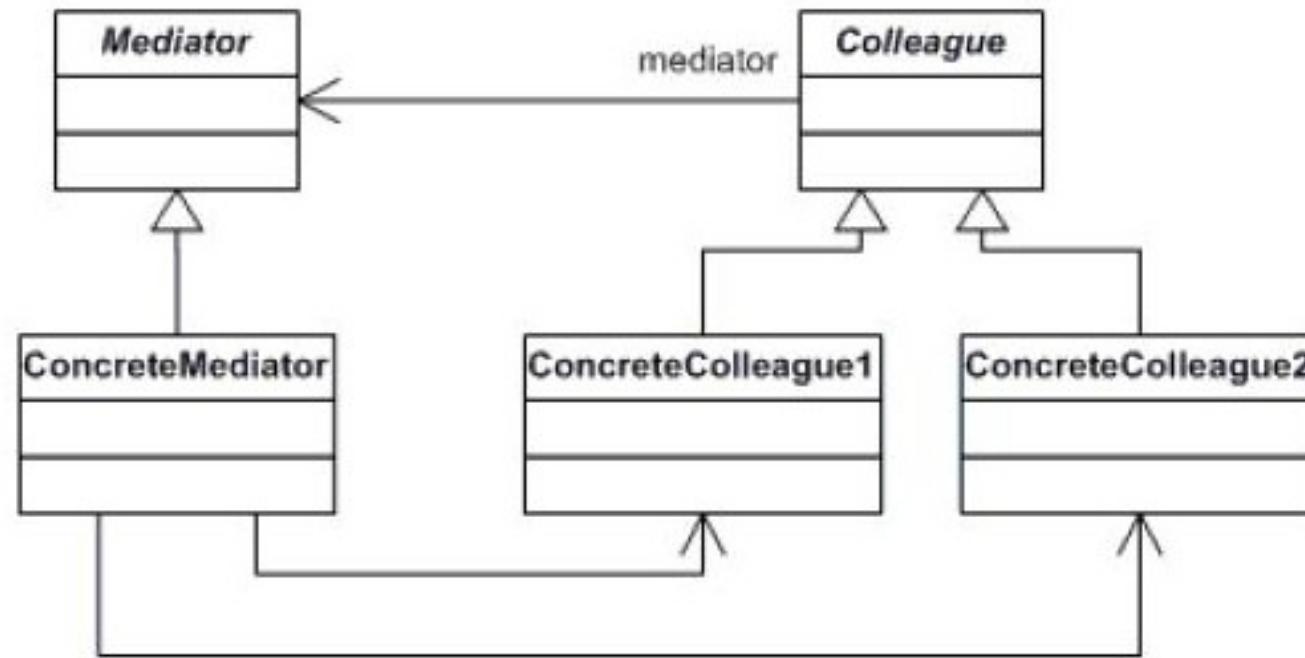
# Proxy Pattern in Android

- Binder decomposes the method call and all its corresponding data to a level that Linux can understand, transmitting it from the local process and address space to the remote process and address space, and reassembling and reenacting the call there.



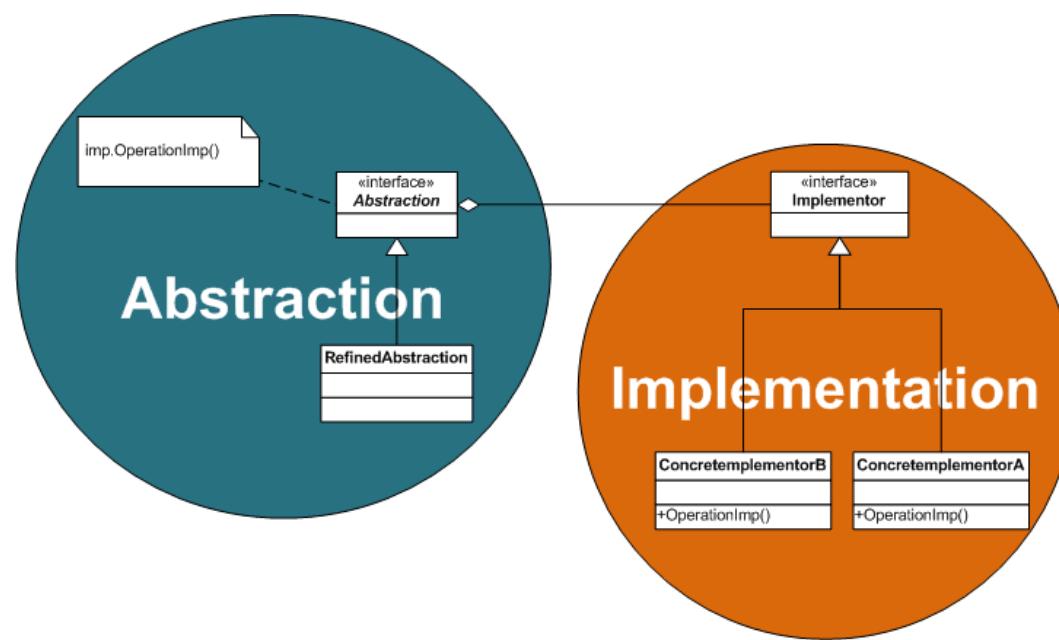
# Mediator Pattern

- With the **mediator pattern**, communication between objects is encapsulated with a **mediator** object.



# Bridge Pattern

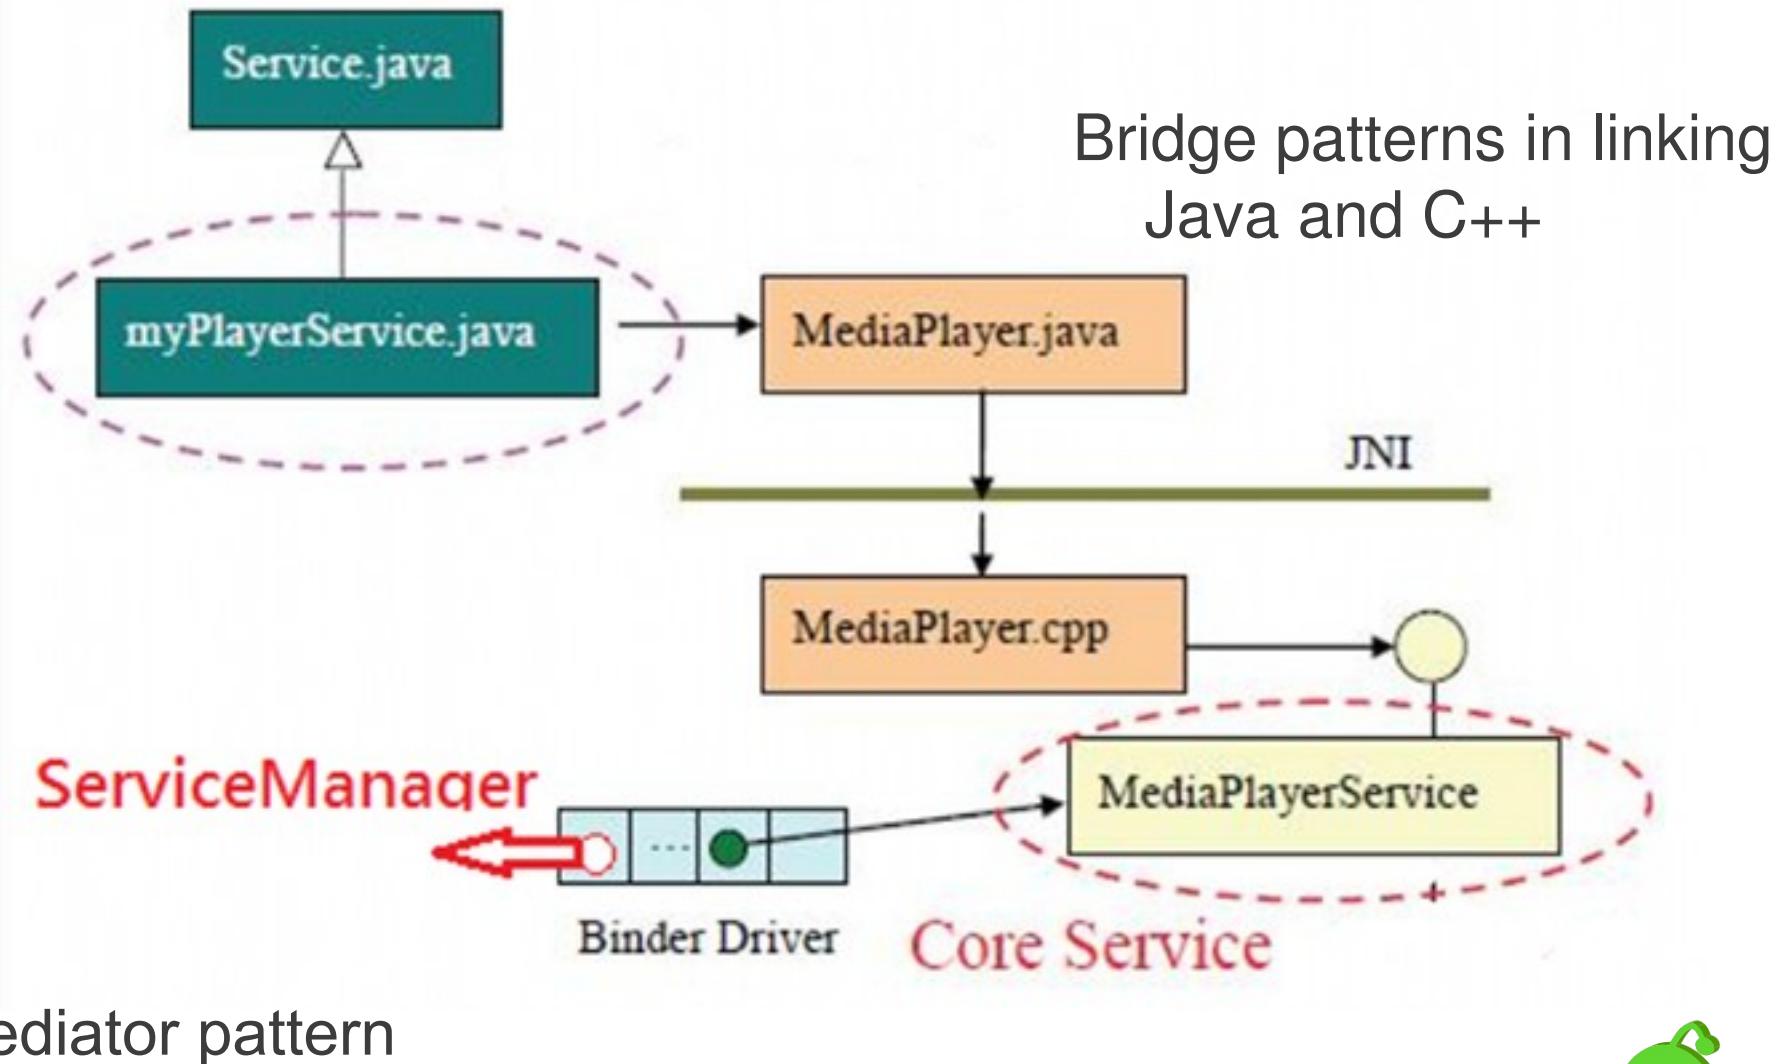
- decouple an abstraction from its implementation so that the two can vary independently



Bridge Pattern - Generic Structure



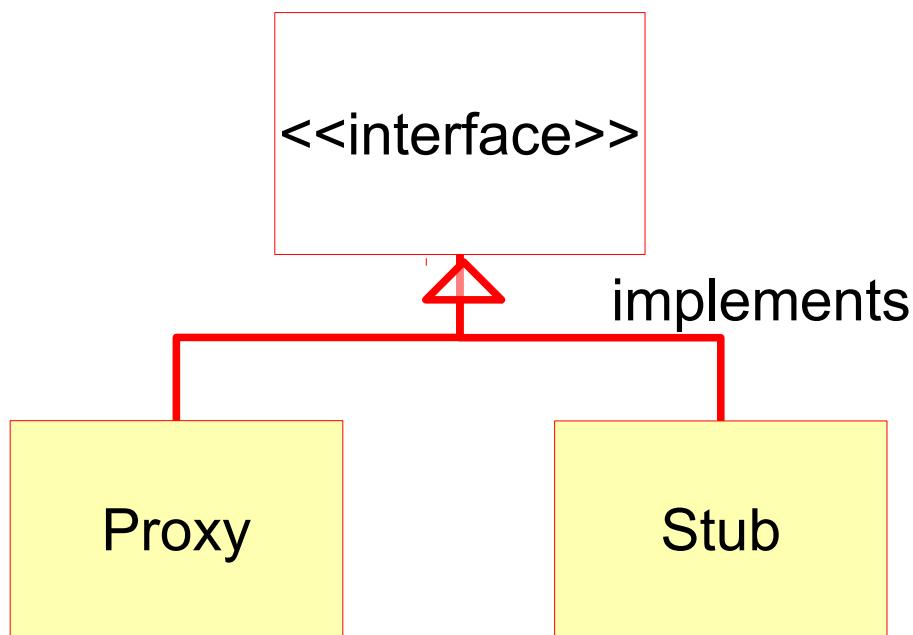
# Bridge and Mediator Pattern in Android



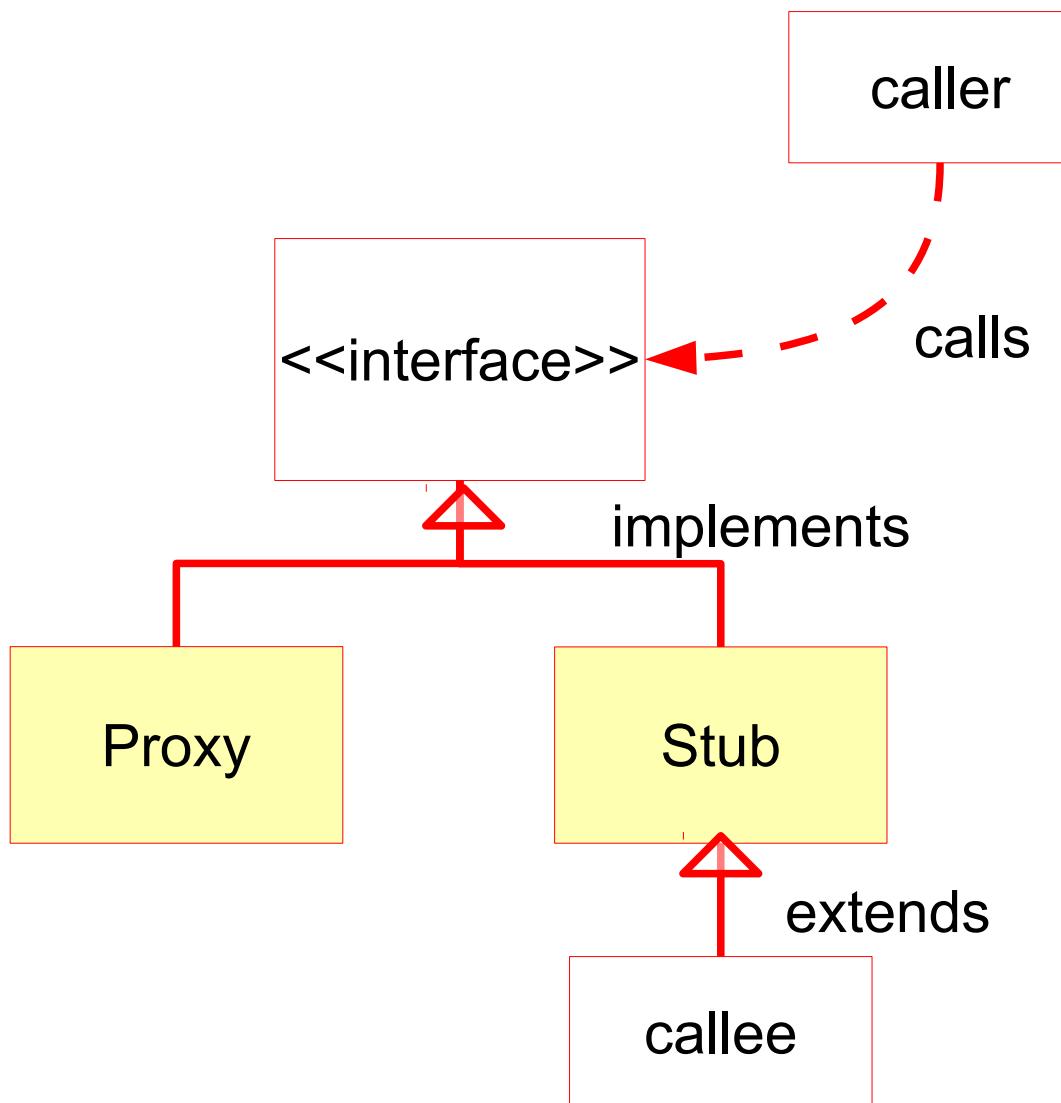
Mediator pattern



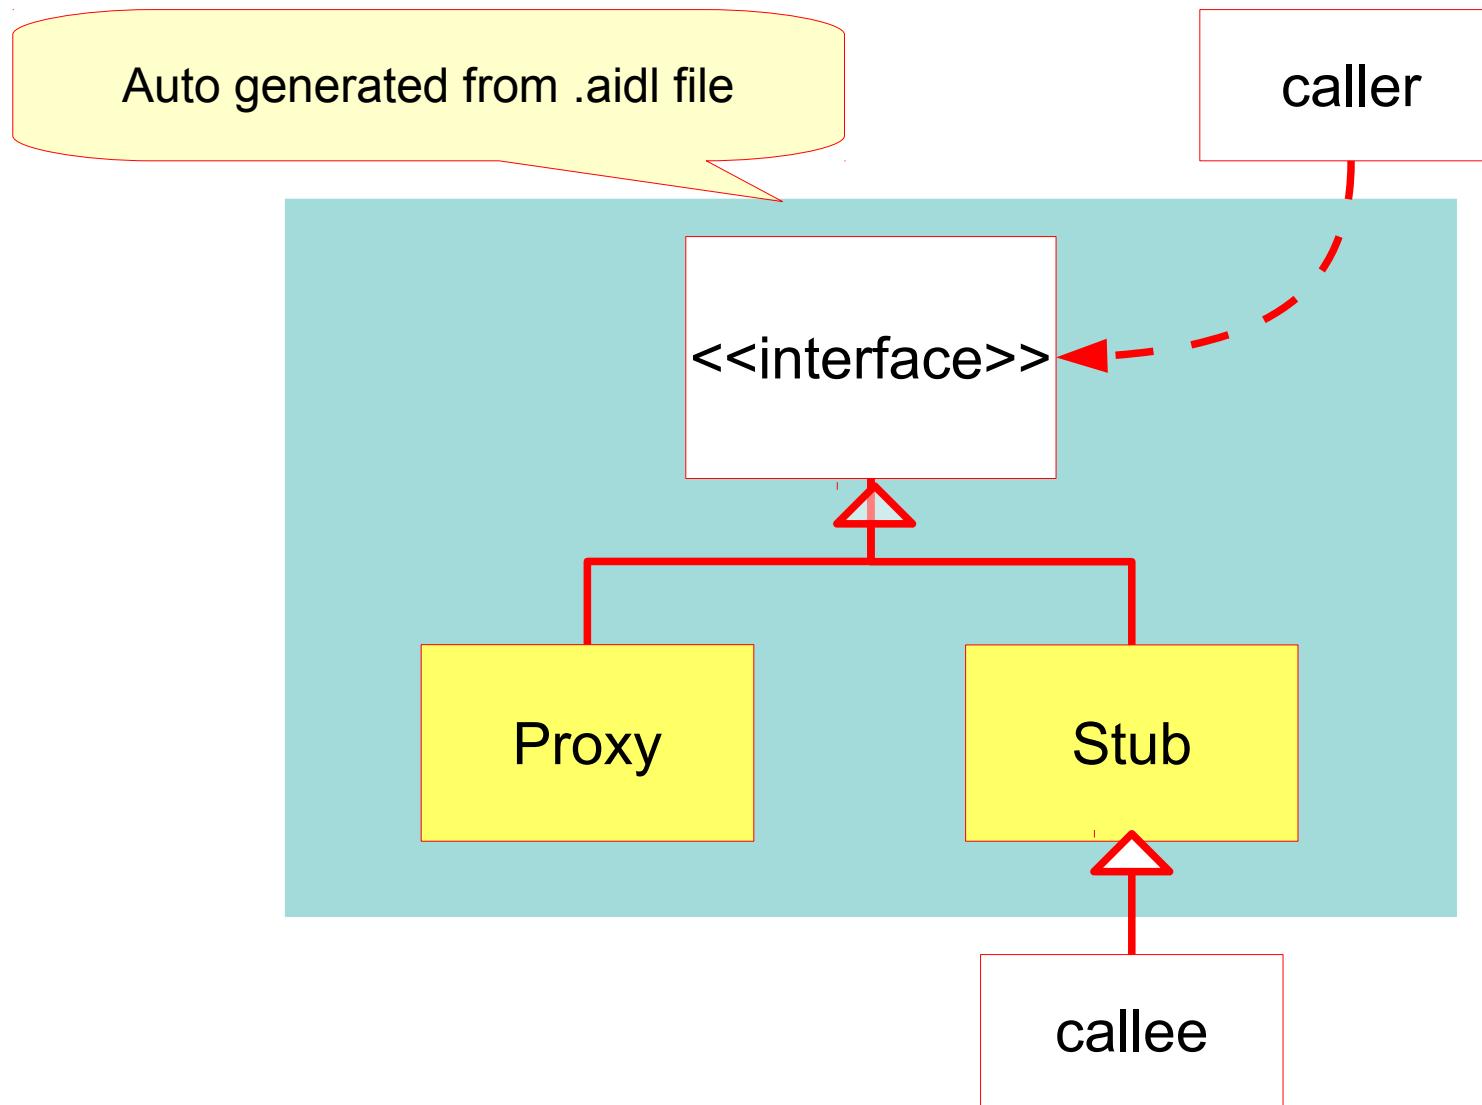
# UML Representation



# UML Representation

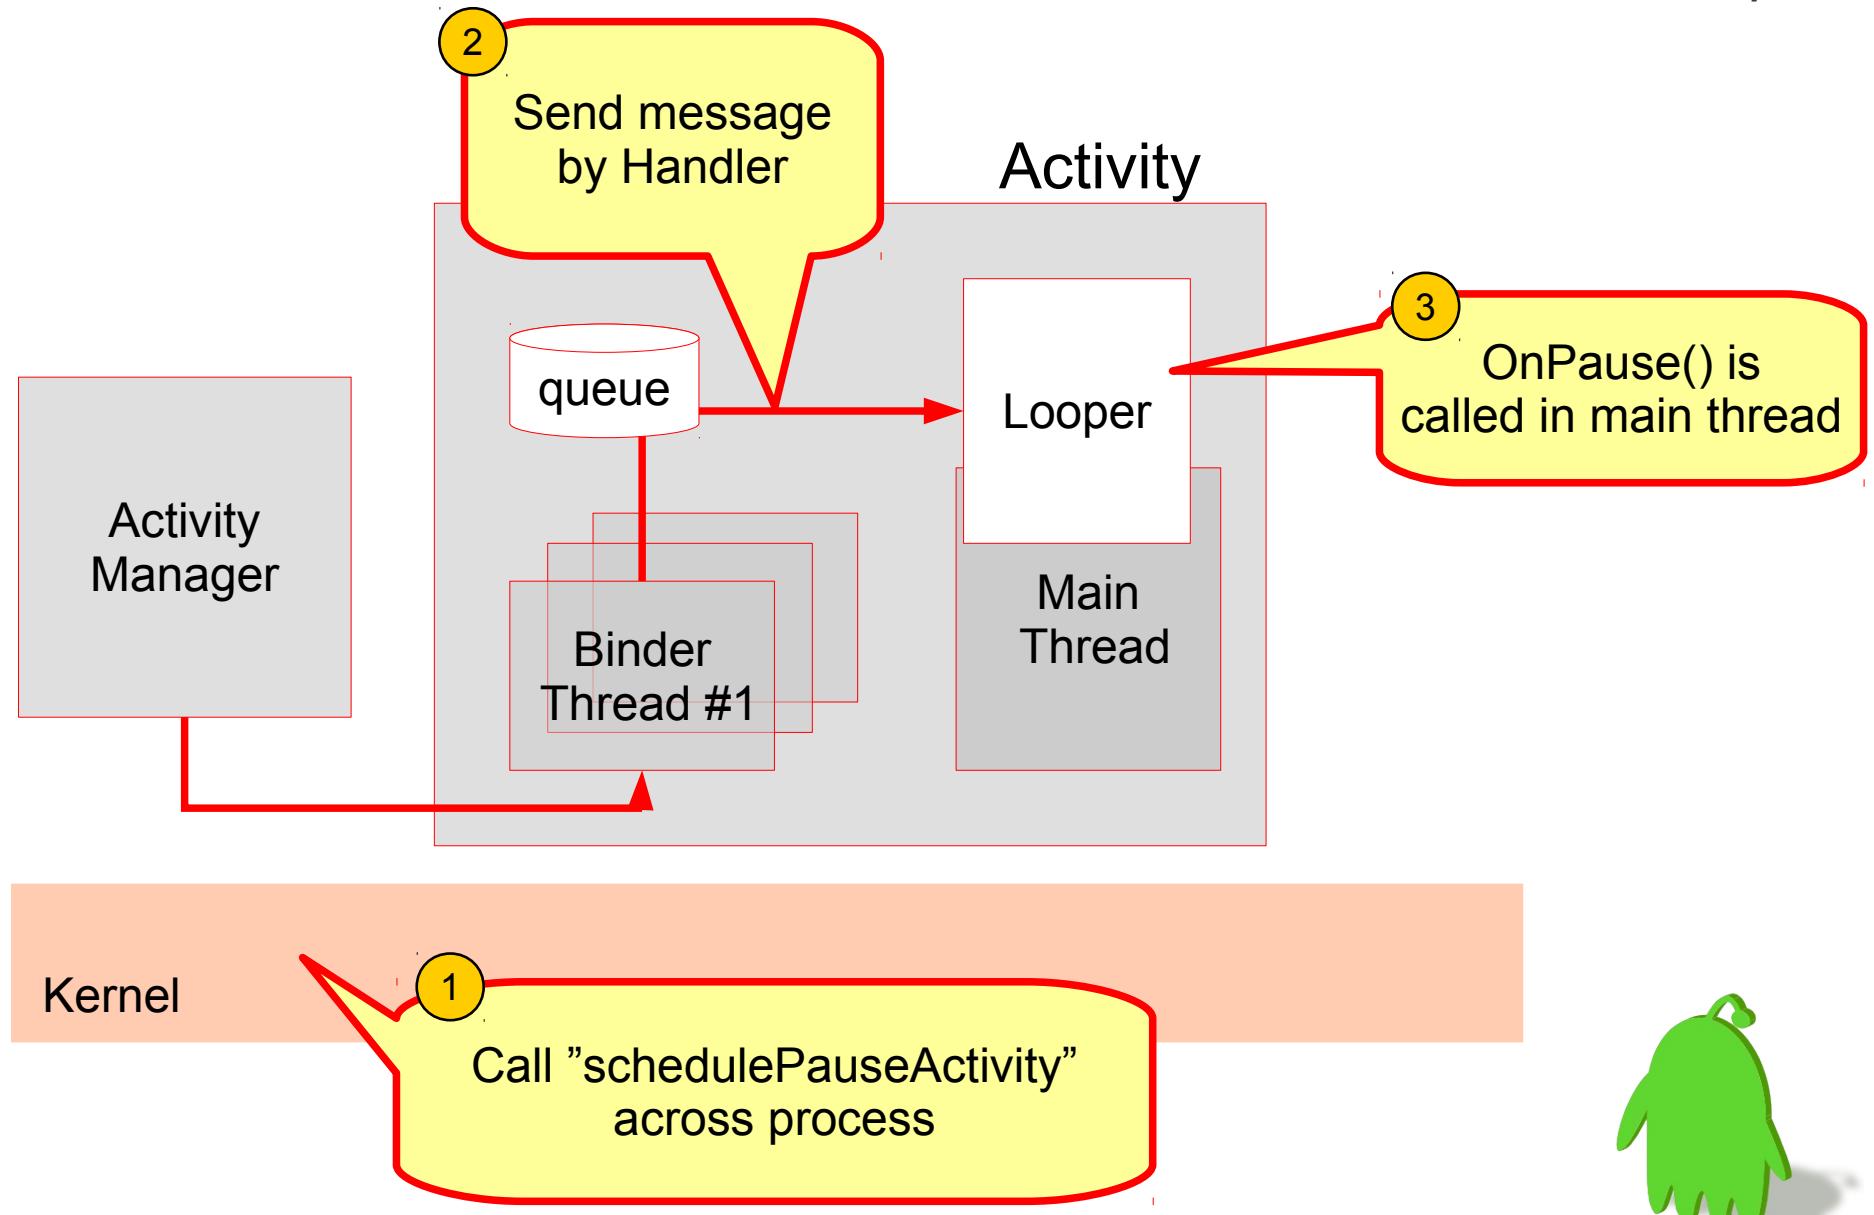


# AIDL

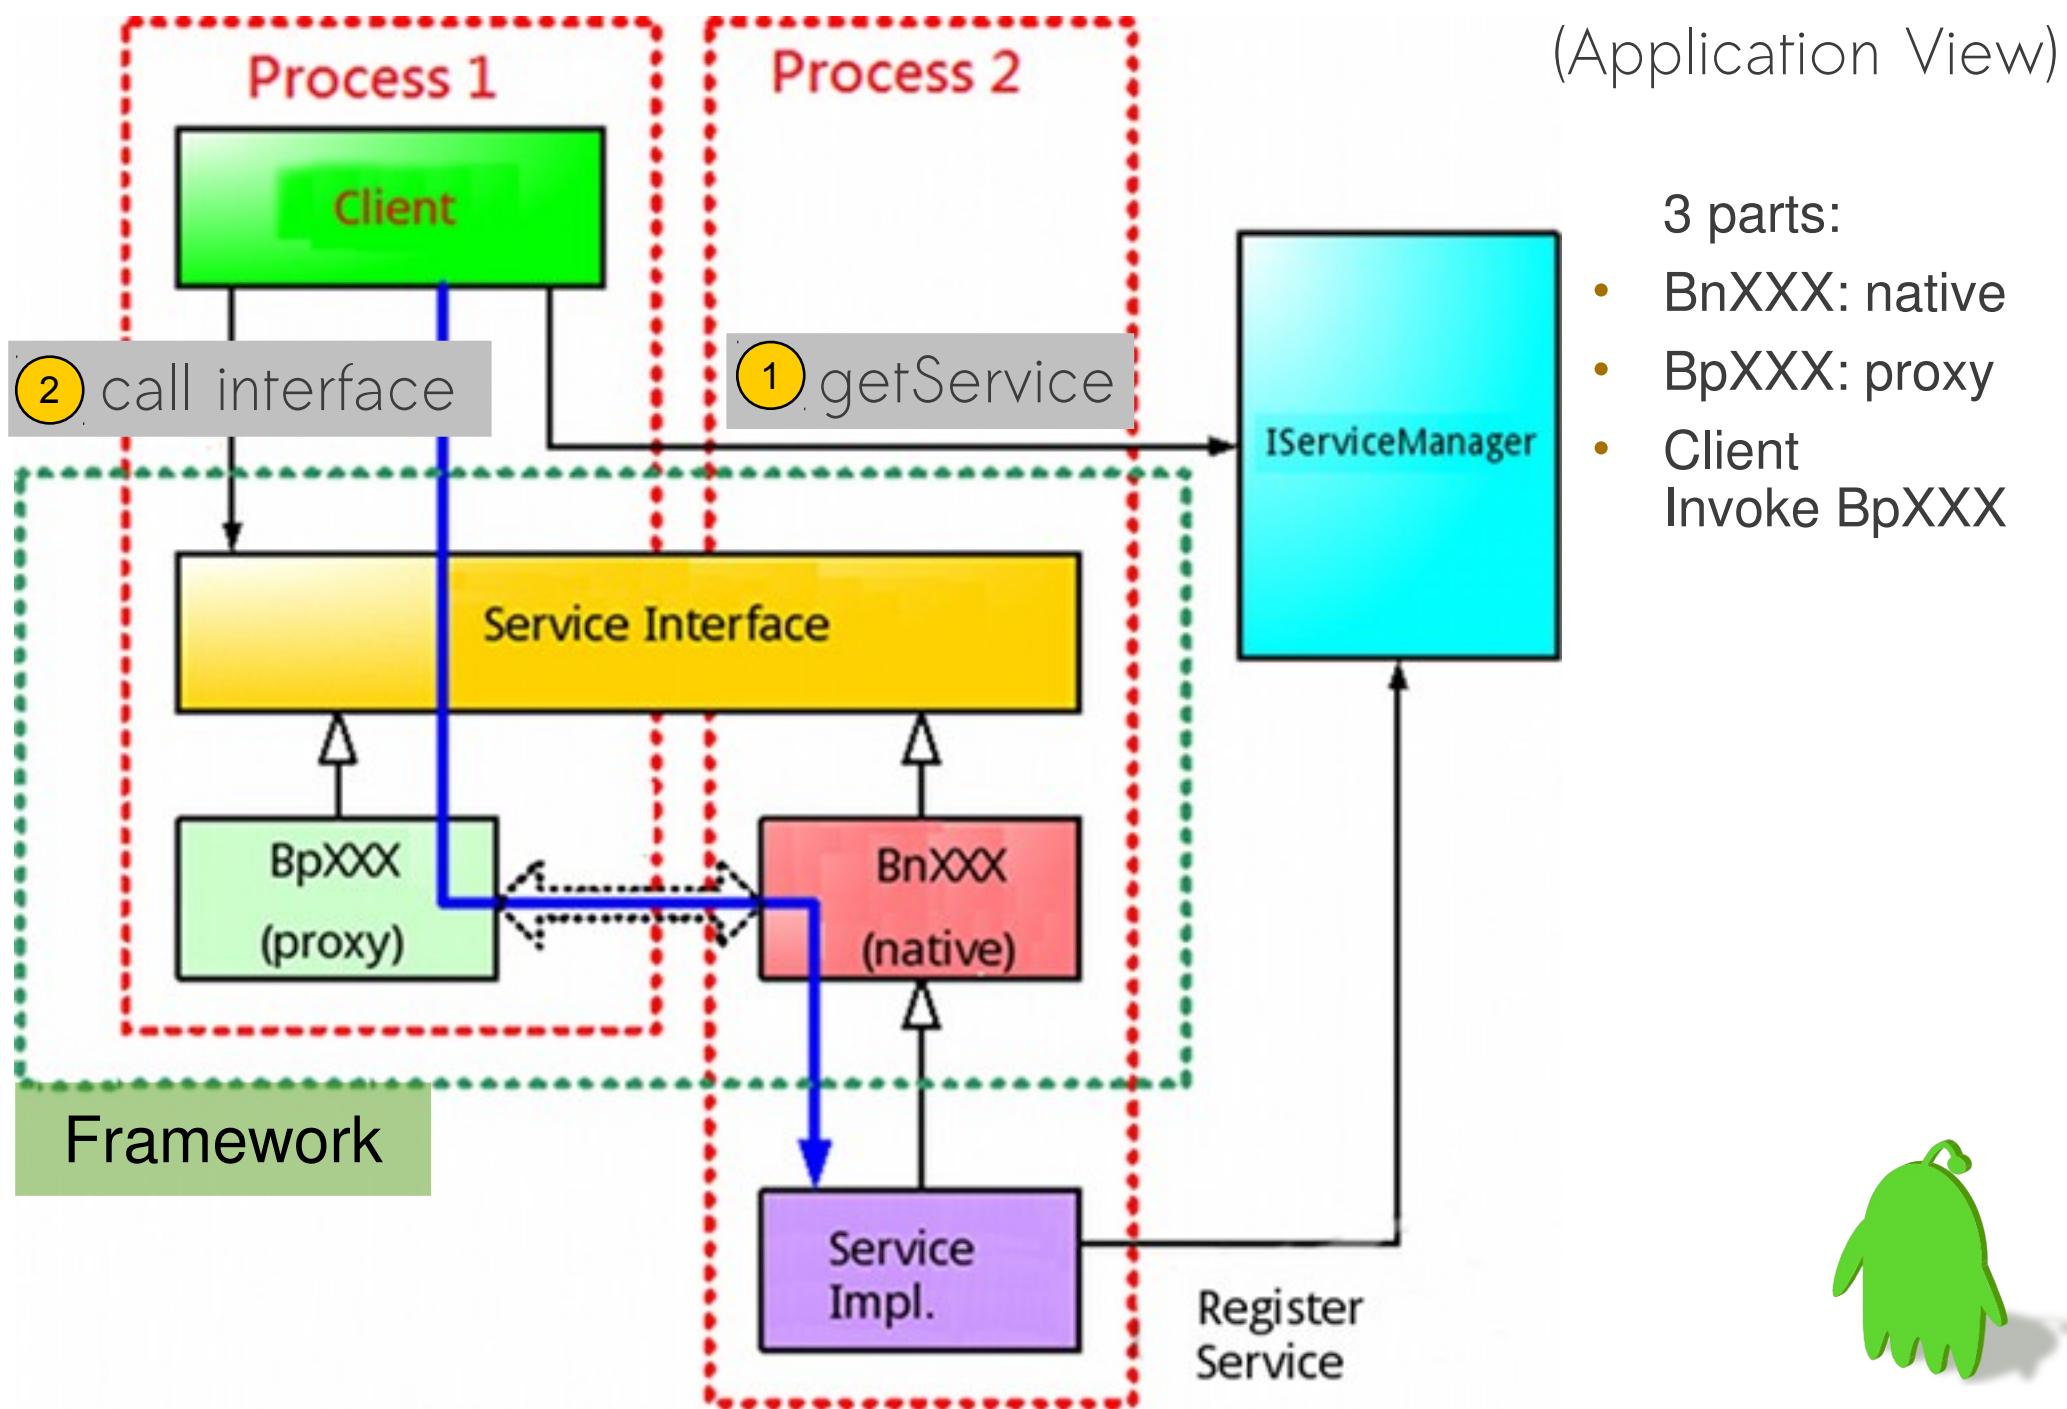


# Use Case:

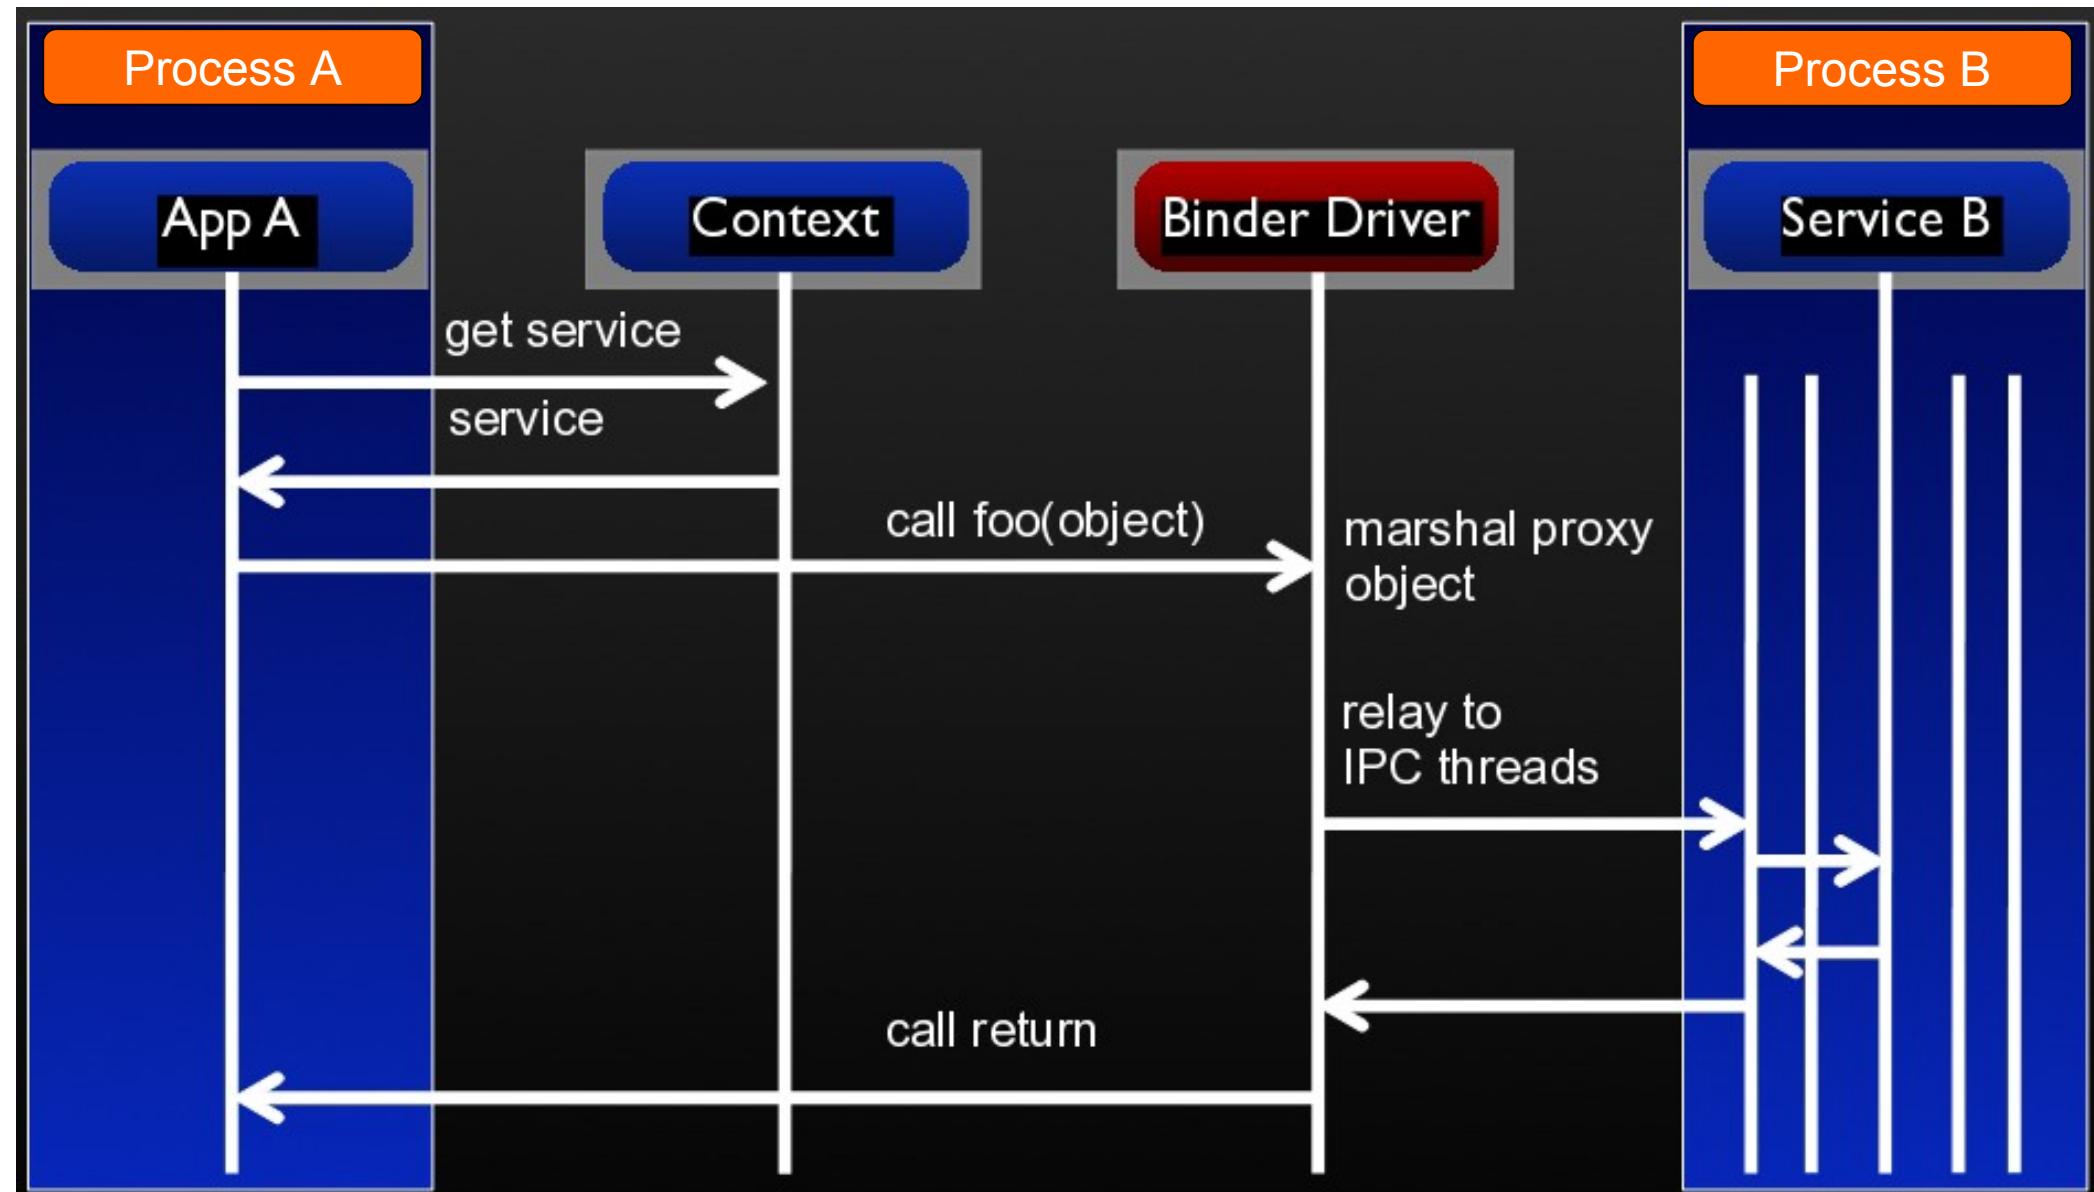
## Who calls onPause() in Activity?



# IPC Interaction in Android



# Binder in Action



# Binder Internals



# Binder Terminology

- Binder
- Binder Object
  - an instance of a class that implements the Binder interface.
  - One Binder object can implement multiple Binders
- Binder Protocol
- IBinder Interface
  - is a well-defined set of methods, properties and events that a Binder can implement.
- Binder Token
  - A numeric value that uniquely identifies a Binder

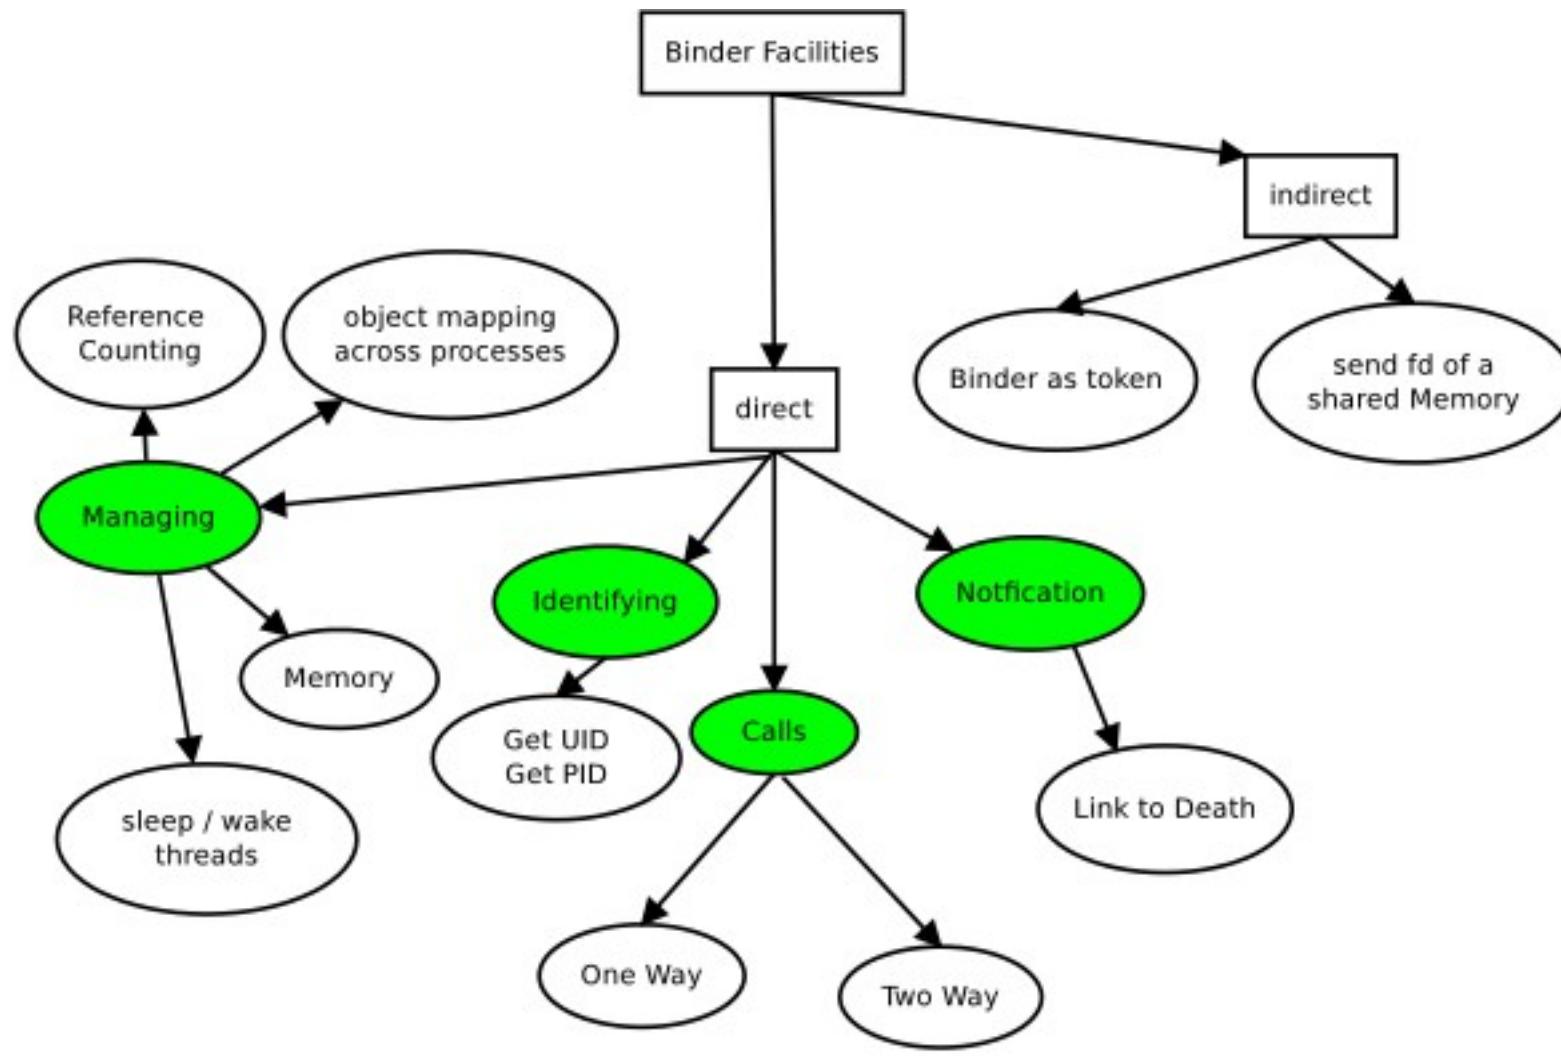


# Facilities

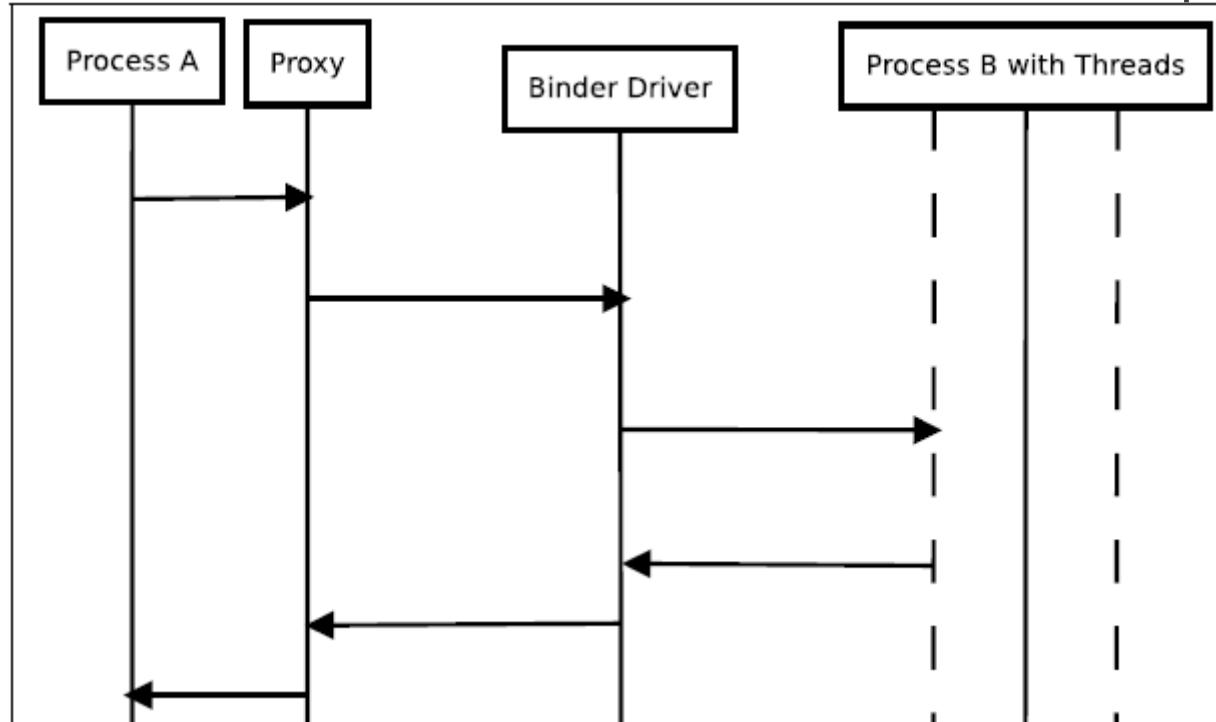
- Simple inter process messaging system
- Managing
- Identifying
- Calls
- Notification
- Binder as a security access token



- Binder framework provides more than a simple interprocess messaging system.
- Methods on remote objects can be called as if they were local object methods.



# Communication protocol



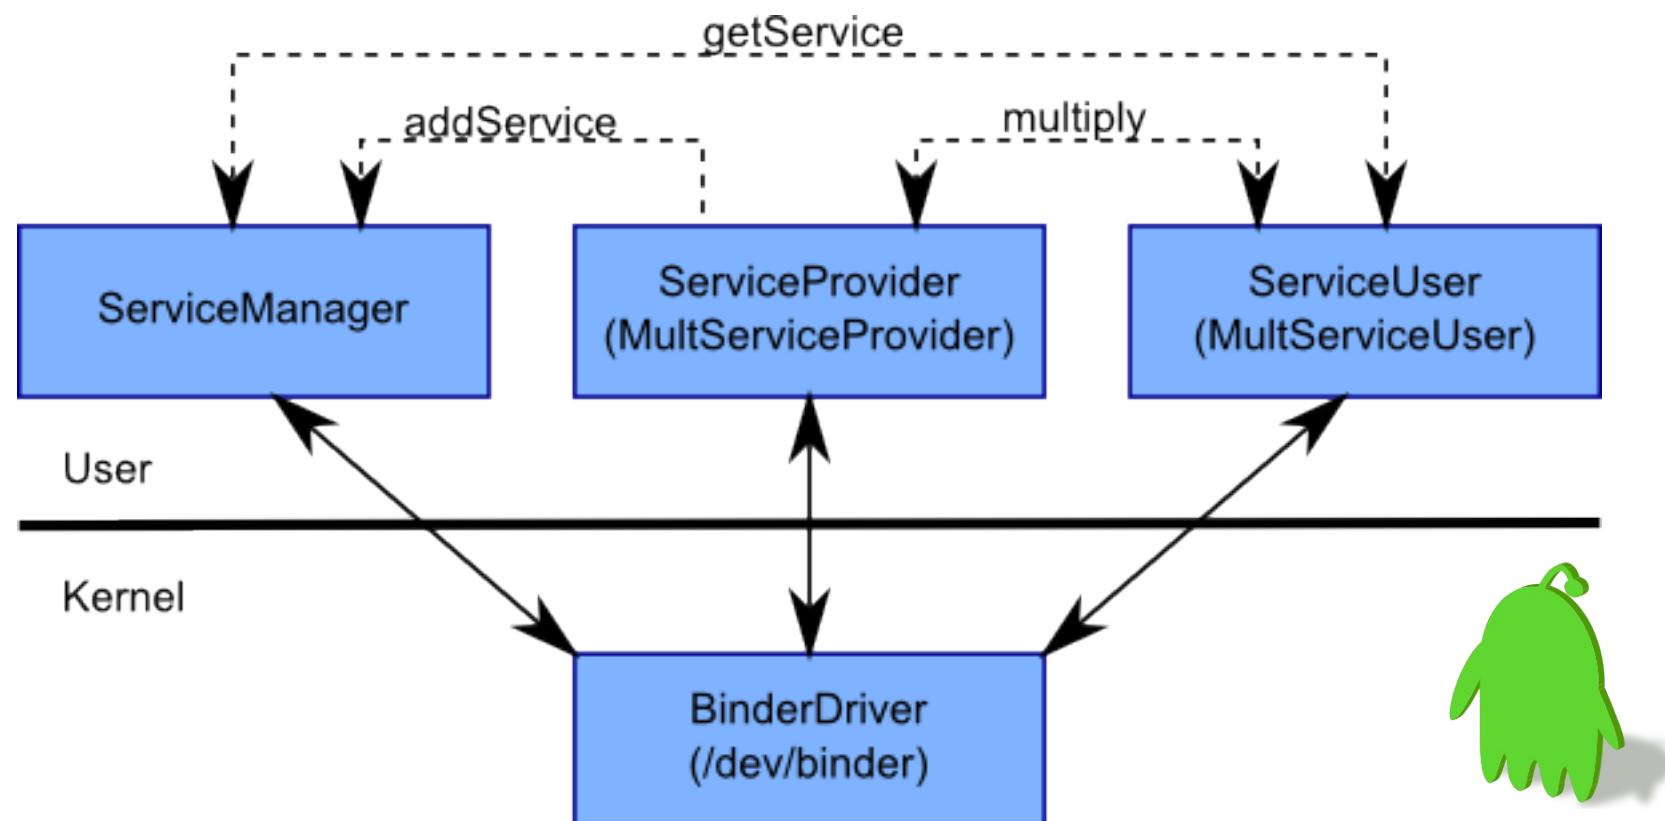
If one process sends data to another process, it is called transaction.  
The data is called transaction data.

| Target              | Binder<br>Driver<br>Command | Cookie | Sender ID | Data:   |                     |             |                  |             |     |     |                    |               |
|---------------------|-----------------------------|--------|-----------|---|---------------------|-------------|------------------|-------------|-----|-----|--------------------|---------------|
|                     |                             |        |           | <table border="1"><tr><td>Target<br/>Command 0</td><td>Arguments 0</td></tr><tr><td>Target Command 1</td><td>Arguments 1</td></tr><tr><td>...</td><td>...</td></tr><tr><td>Target Command n-1</td><td>Arguments n-1</td></tr></table> | Target<br>Command 0 | Arguments 0 | Target Command 1 | Arguments 1 | ... | ... | Target Command n-1 | Arguments n-1 |
| Target<br>Command 0 | Arguments 0                 |        |           |   |                     |             |                  |             |     |     |                    |               |
| Target Command 1    | Arguments 1                 |        |           |   |                     |             |                  |             |     |     |                    |               |
| ...                 | ...                         |        |           |   |                     |             |                  |             |     |     |                    |               |
| Target Command n-1  | Arguments n-1               |        |           |   |                     |             |                  |             |     |     |                    |               |



# Service Manager (SM)

- Special Binder node with known Binder address
- Client does not know the address of remote Binder
  - only Binder interface knows its own address
- Binder submits a name and its Binder token to SM
  - Client retrieves Binder address with service name from SM



# Get Service list from SM

```
$ adb shell service list
```

Found 71 services:

```
0  stub_isms: [com.android.internal.telephony.ISms]
1  stub_phone: [com.android.internal.telephony.ITelephony]
2  stub_iphonesubinfo:
    [com.android.internal.telephony.IPhoneSubInfo]
...
5  stub_telephony.registry:
    [com.android.internal.telephony.ITelephonyRegistry]
...
7  stub_activity: [android.app.IActivityManager]
...
9  phone: [com.android.internal.telephony.ITelephony]
...
56 activity: [android.app.IActivityManager]
...
64 SurfaceFlinger: [android.ui.ISurfaceComposer]
...
```



# Call remote method in ActivityManager

```
$ adb shell service list
```

...

56 activity: [android.app.IActivityManager]

...

```
$ adb service call activity 1598968902
```

Result: Parcel(

|             |                                     |                     |
|-------------|-------------------------------------|---------------------|
| 0x00000000: | 0000001c 006e0061 00720064 0069006f | '....a.n.d.r.o.i.'  |
| 0x00000010: | 002e0064 00700061 002e0070 00410049 | 'd...a.p.p...I.A.'  |
| 0x00000020: | 00740063 00760069 00740069 004d0079 | 'c.t.i.v.i.t.y.M.'  |
| 0x00000030: | 006e0061 00670061 00720065 00000000 | 'a.n.a.g.e.r.....') |

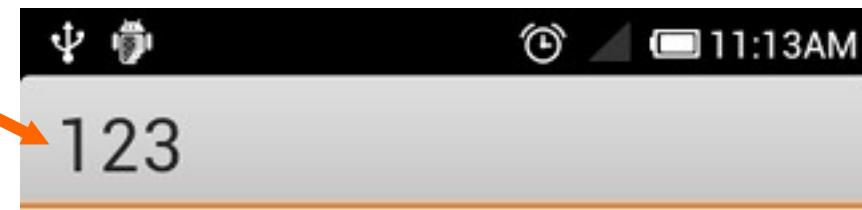
```
public abstract interface IBinder {  
    ...  
    field public static final int INTERFACE_TRANSACTION  
        = 1598968902; // 0x5f4e5446  
    ...  
}
```

Source: frameworks/base/api/current.txt

# Interact with Android Service

```
$ adb service call phone 1 s16 "123"
```

Result: Parcel(00000000 '....')



```
interface ITelephony {  
    /* Dial a number. This doesn't place the call. It displays  
     * the Dialer screen. */  
    void dial(String number);
```

Source: frameworks/base/  
telephony/java/com/android/internal/telephony/ITelephony.aidl

```
service call SERVICE CODE [i32 INT | s16 STR] ...
```

Options:

i32: Write the integer INT into the send parcel.

s16: Write the UTF-16 string STR into the send parcel.



```
$ adb shell service list
```

Found 71 services:

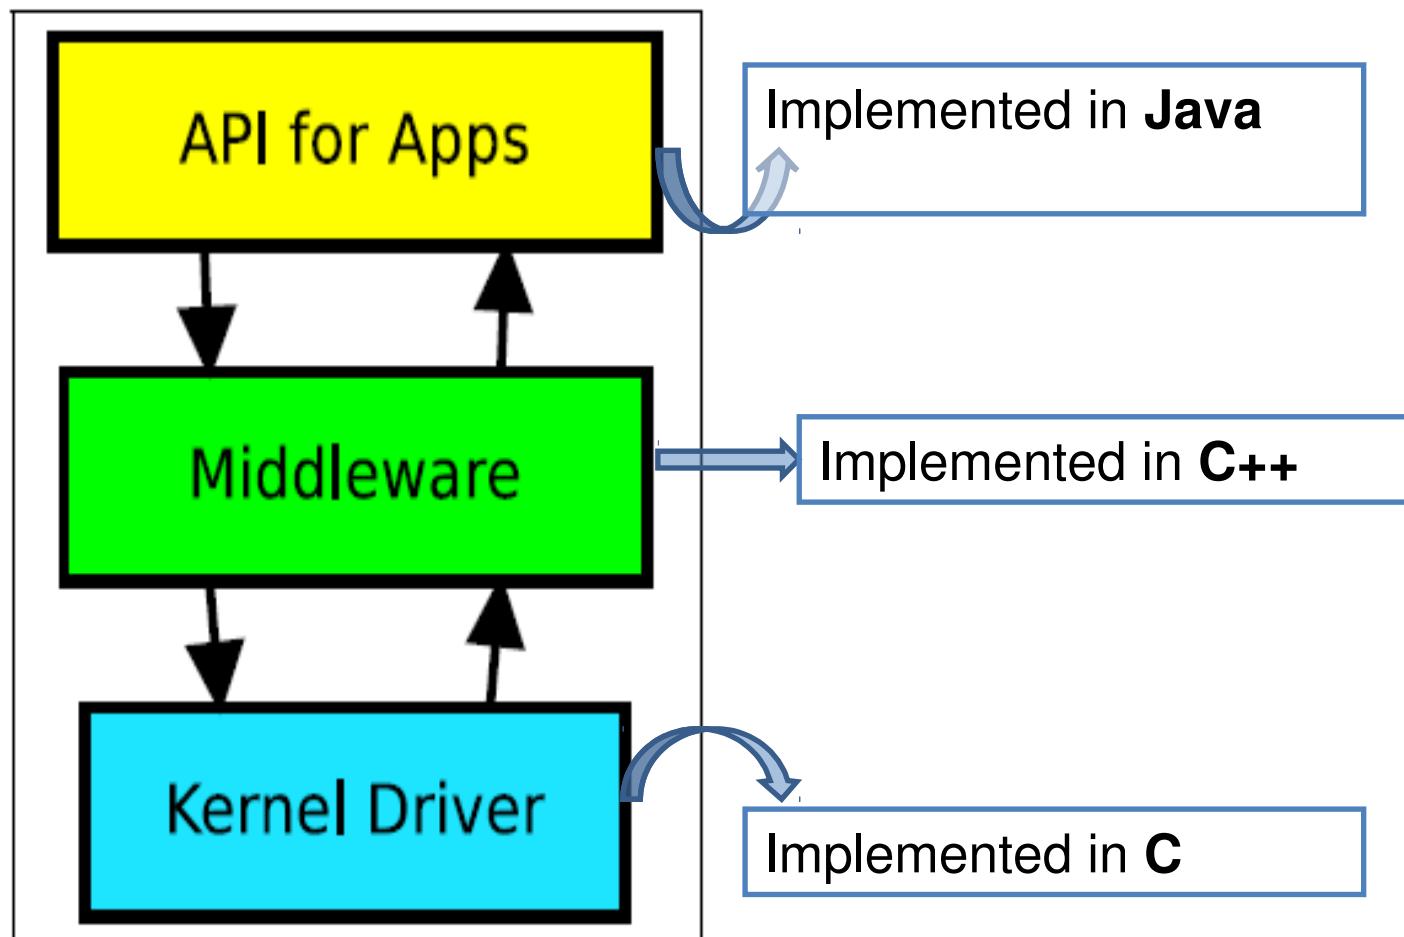
...

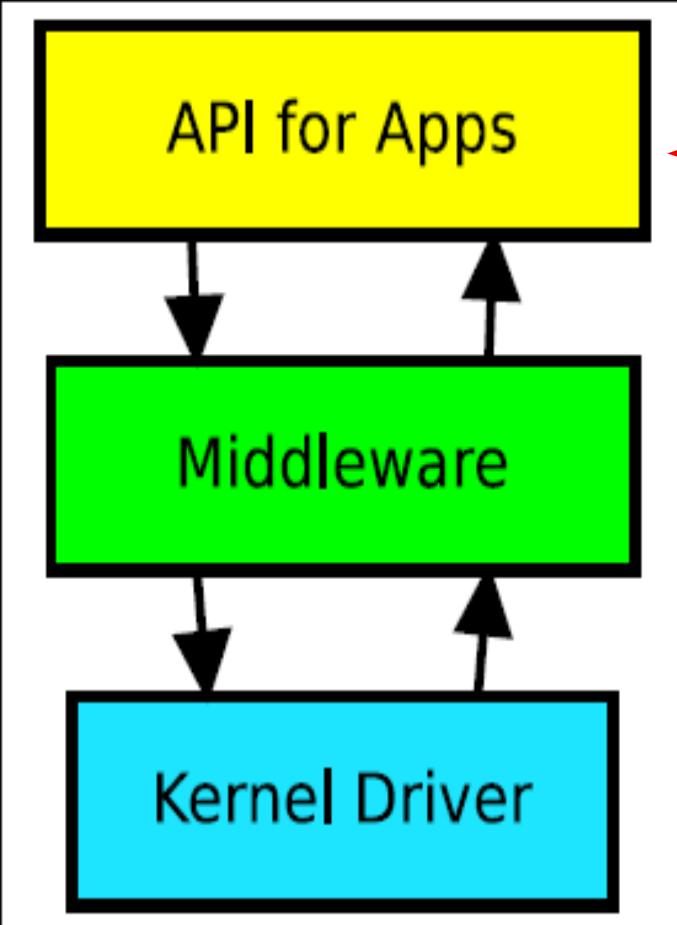
9 phone: [com.android.internal.telephony.ITelephony]

Phone Application appears in foreground.  
parameter "1" → dial()  
s16 "123" → String("123")



# Implementation Layers of Binder





## API Layer

- **AIDL** (Android Interface Definition Language)
  - Ease the implementation of Android remote services
  - Defines an interface with methods of remote services
  - AIDL parser generates Java class
    - Proxy class for Client
    - Stub class for Service
- **Java API Wrapper**
  - Introduce facilities to the binder
  - Wraps the middleware layer



- Data Types
  - Java Primitives
  - Containers
    - String, List, Map, CharSequence
    - List<>
    - Multidimensional Array
  - Parcelable
  - Interface Reference
- Direction - in, out, inout
- oneway
  - android.os.IBinder.FLAG\_ONEWAY



# AIDL Compiler

- Full-fledged Java(-only) Support
- Stub and Proxy Generator

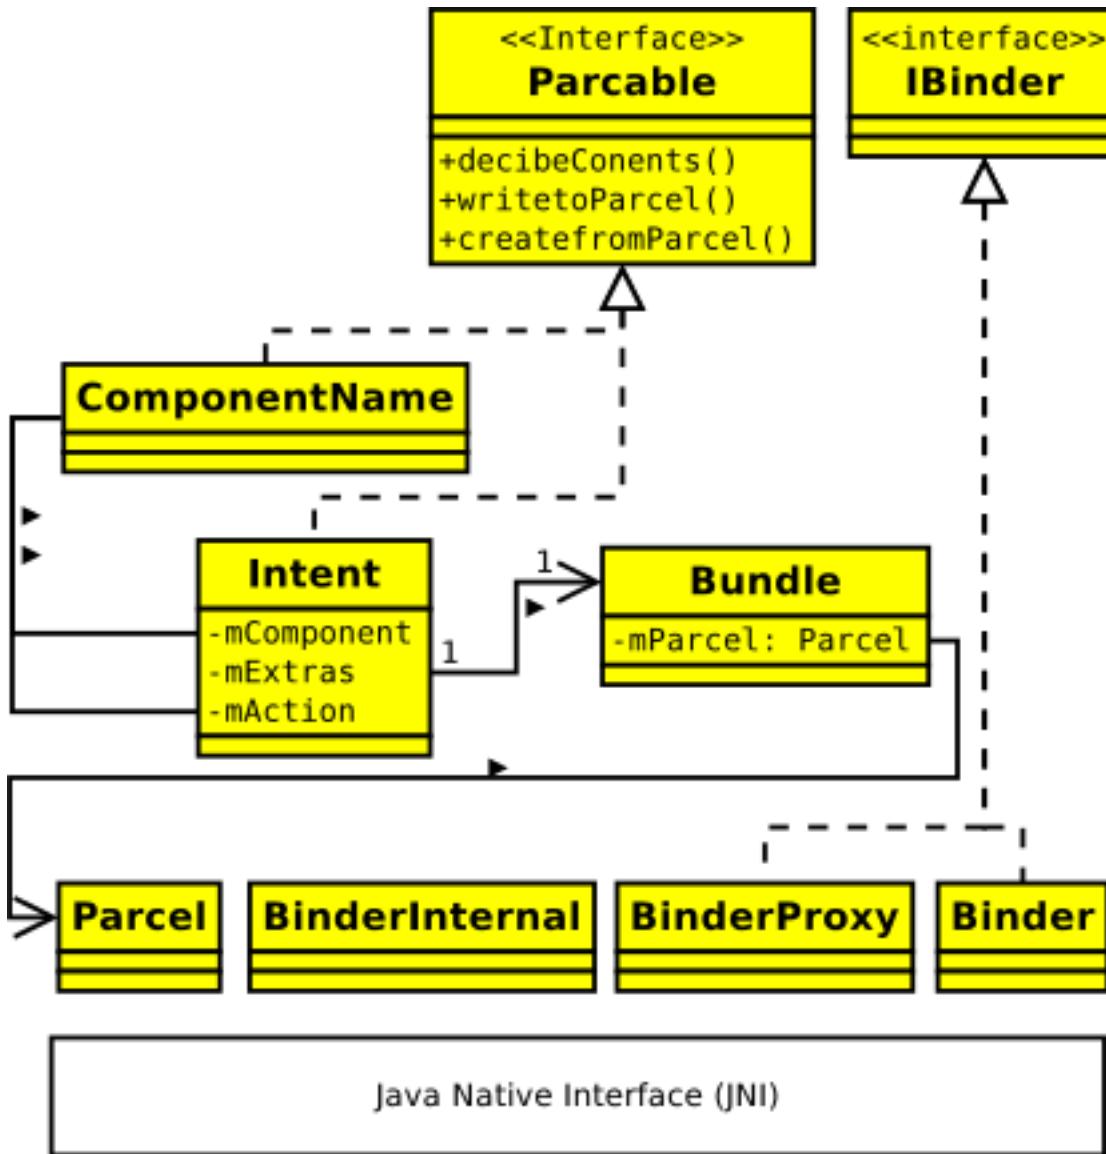
```
// Interface  
interface IRemoteService {  
    void ping();  
}
```

The diagram illustrates the AIDL compilation process. It starts with an 'Interface' code block at the top left, which is connected by a red arrow pointing down to a 'Client' code block. Another red arrow points from the 'Client' code back up to the 'Interface' code. The 'Client' code contains code to create a proxy object: `IRemoteService mService = IRemoteService.Stub.asInterface(service);`. The 'Client' code is enclosed in an orange box.

```
IRemoteService mService =  
    IRemoteService.Stub.asInterface(service);
```

The diagram illustrates the AIDL compilation process. It starts with an 'Interface' code block at the top left, which is connected by a red arrow pointing down to a 'Server' code block. Another red arrow points from the 'Server' code back up to the 'Interface' code. The 'Server' code contains the implementation of the service: `public class RemoteService extends Service { ... }`. The 'Server' code is enclosed in an orange box.

```
public class RemoteService extends Service {  
    public IBinder onBind(Intent intent) { return mBinder; }  
    private final IRemoteService.Stub mBinder =  
        new IRemoteService.Stub() {  
            public void ping() { // Nothing }  
        };  
}
```



Android\_util\_binder.cpp



# Parcels and Marshalling

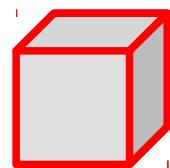
- Simple inter process messaging system
- In an object oriented view, the transaction data is called parcel.
- The procedure of building a parcel is called **marshalling** an object.
- The procedure of rebuilding a object from a parcel is called **unmarshalling** an object.



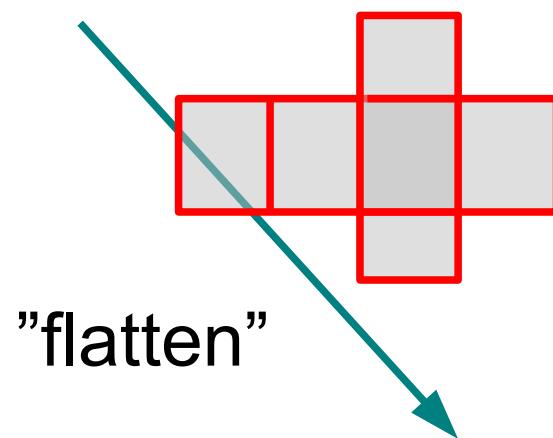
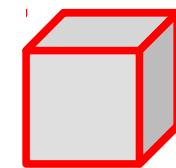
- Marshalling – The transferring of data across process boundaries
  - Represented in native binary encoding
- Mostly handled by AIDL-generated code
- Extensible – Parcelable



android.os.Parcel



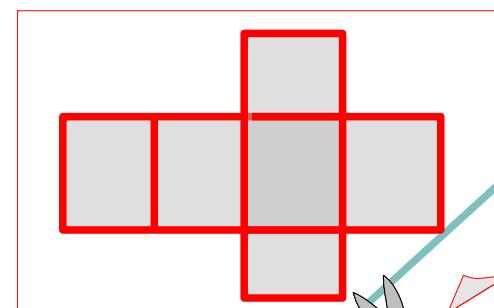
Delivering arguments of method



"flatten"



transmit

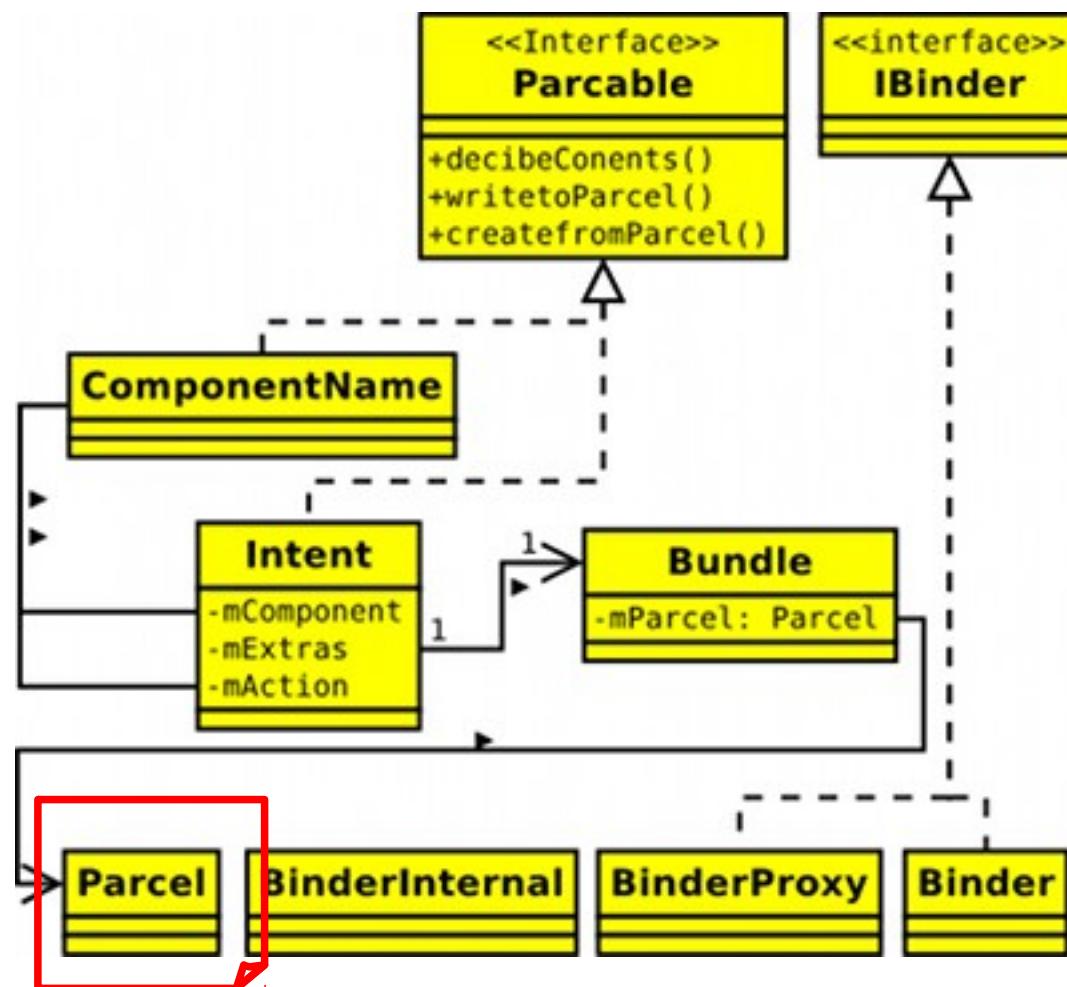


"unflatten"



# Parcel Definition

- Container for a message (data and object references) that can be sent through an IBinder.
  - A Parcel can contain both flattened data that will be unflattened on the other side of the IPC (using the various methods here for writing specific types, or the general Parcelable interface), and references to live IBinder objects that will result in the other side receiving a proxy IBinder connected with the original IBinder in the Parcel.



# Representation of Parcel

- Parcel is not for general-purpose serialization
  - This class (and the corresponding Parcelable API for placing arbitrary objects into a Parcel) is designed as a high-performance IPC transport.
  - Not appropriate to place any Parcel data into persistent storage
- Functions for writing/reading primitive data types:
  - `writeByte(byte)` / `readByte()`
  - `writeDouble(double)` / `readDouble()`
  - `writeFloat(float)` / `readFloat()`
  - `writeInt(int)` / `readInt()`
  - `writeLong(long)` / `readLong()`
  - `writeString(String)` / `readString()`



# Parcelable

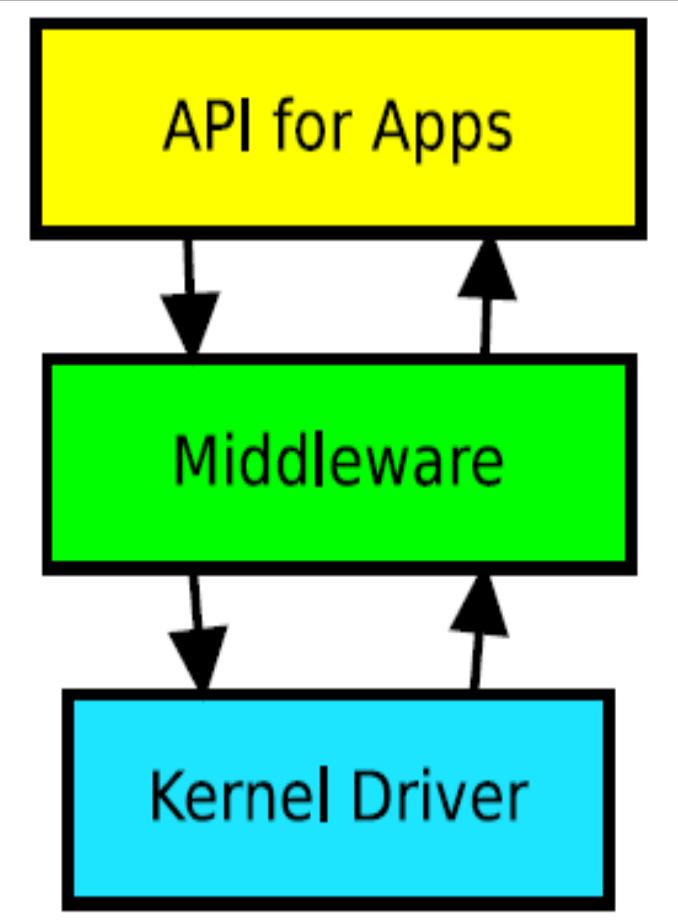
- The Parcelable protocol provides an extremely efficient (but low-level) protocol for objects to write and read themselves from Parcels.
- Use the direct methods to write/read
  - `writeParcelable(Parcel, int)`
  - `readParcelable(ClassLoader)`
  - `writeParcelableArray(T[], int)`
  - `readParcelableArray(ClassLoader)`
- These methods write both the class type and its data to the Parcel, allowing that class to be reconstructed from the appropriate class loader when later reading.



# Bundles

- A special type-safe container, called Bundle, is available for key/value maps of heterogeneous values.
- This has many optimizations for improved performance when reading and writing data, and its type-safe API avoids difficult to debug type errors when finally marshalling the data contents into a Parcel.



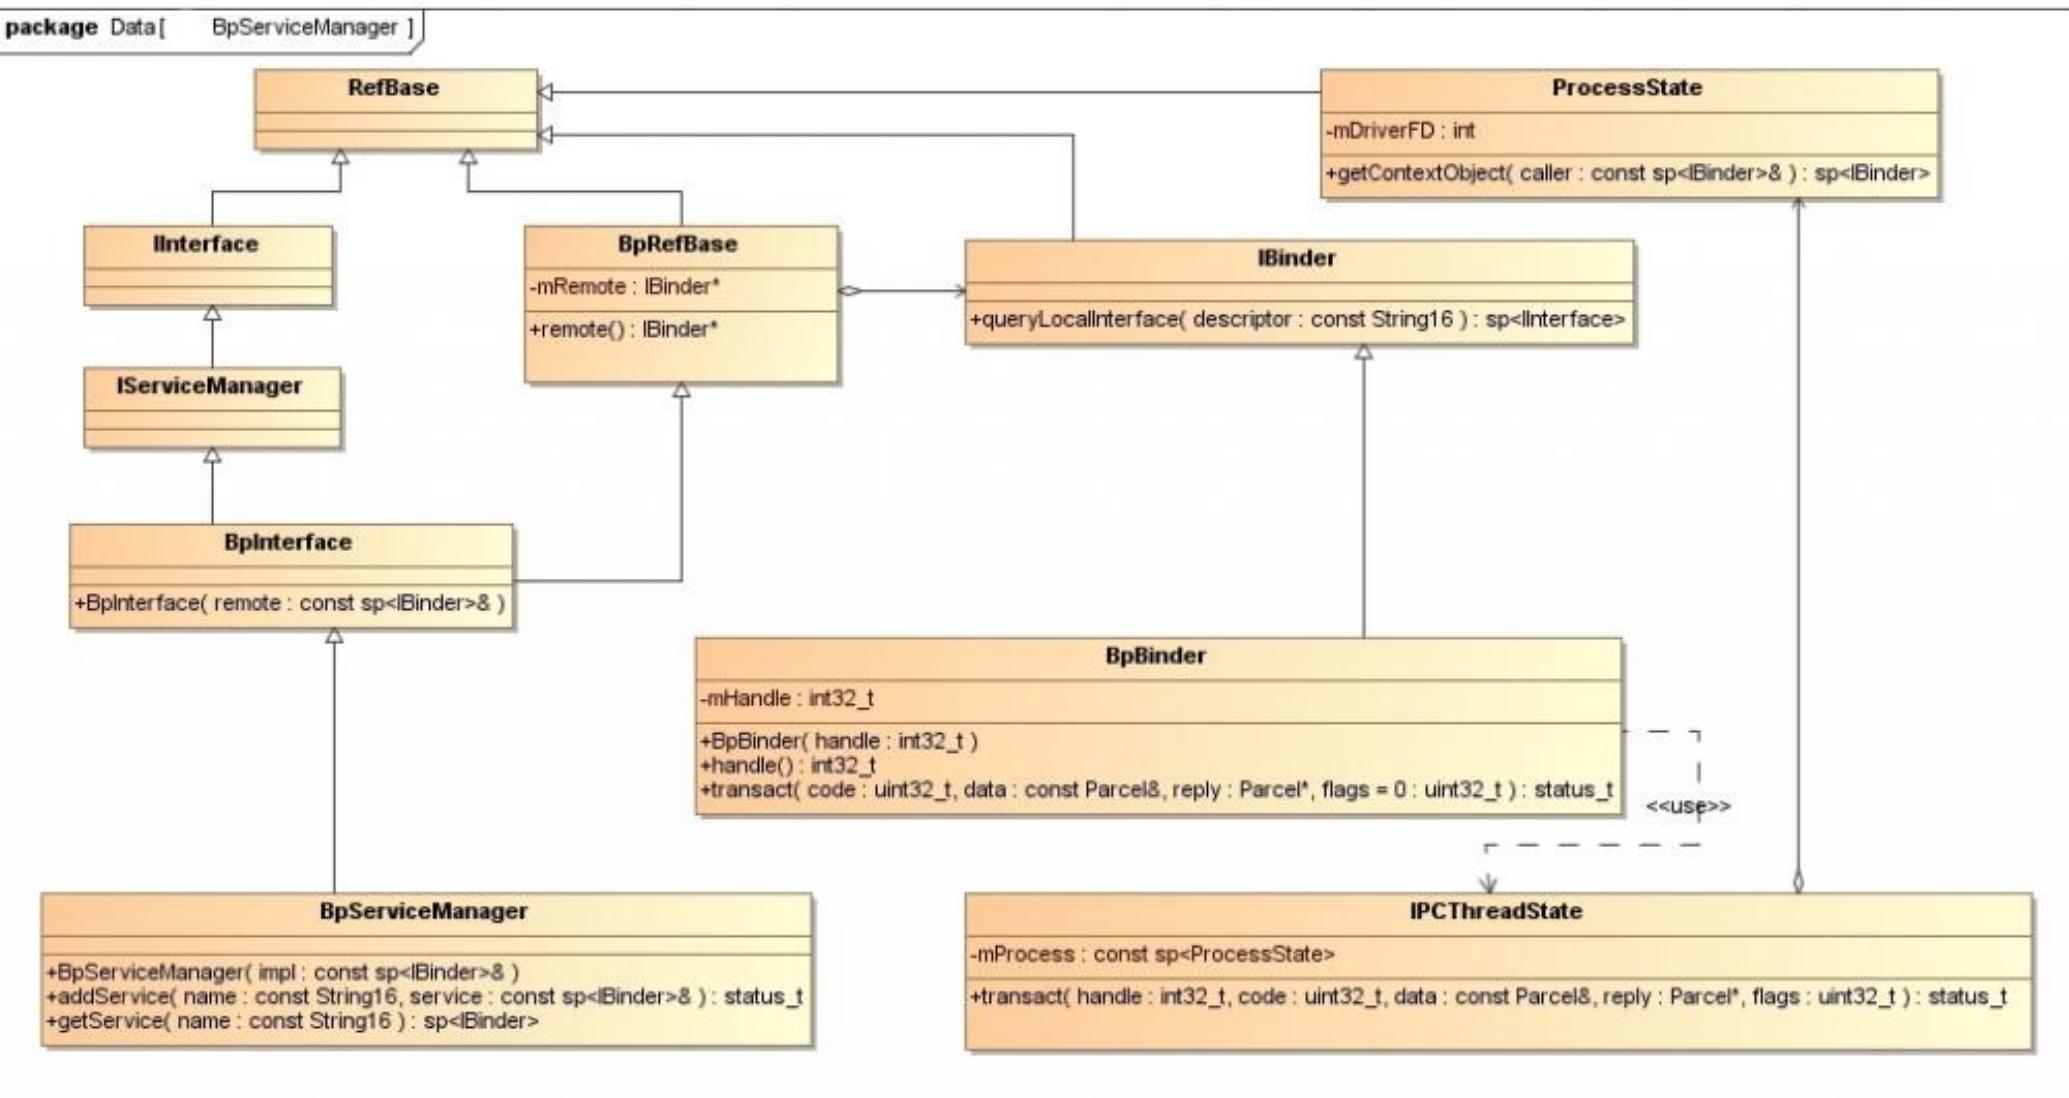


## Middleware Layer

- Implements the user space facilities of the Binder framework in C++
- Implements structures and methods to spawn and manage new threads
- Marshalling and unmarshalling of specific data
- Provides interaction with the Binder kernel driver



- frameworks/base/include/binder/IServiceManager.h  
`sp<IServiceManager> defaultServiceManager()`
- frameworks/base/include/binder/IInterface.h  
 template **BpInterface**



API for Apps

Middleware

Kernel Driver

# Kernel Driver Layer

- Binder Driver supports the file operations open, mmap, release, poll and the system call ioctl
  - ioctl arguments
    - Binder driver command code
    - Data buffer

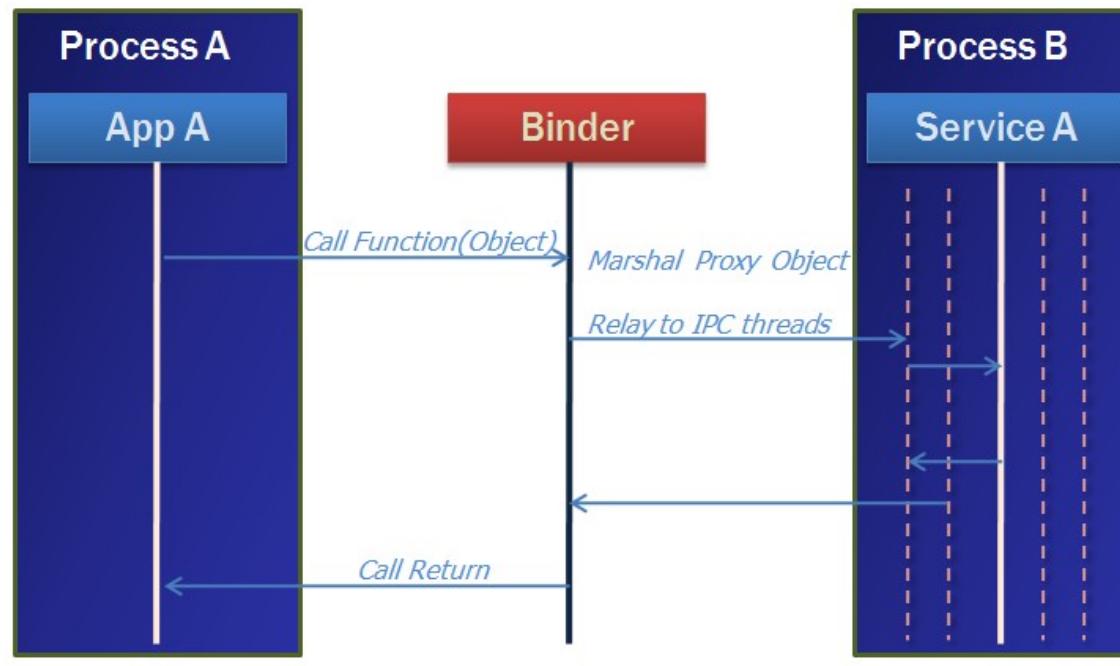
Command codes

- `BINDER_WRITE_READ`
- `BINDER_SET_MAX_THREADS`
- `BINDER_SET_CONTEXT_MGR`
- `BINDER_THREAD_EXIT`
- `BINDER_VERSION`



# Binder Driver

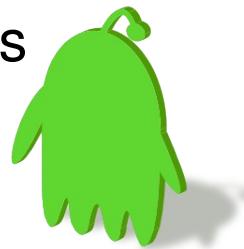
- Multi-thread aware
  - Have internal status per thread
  - Compare to UNIX socket: sockets have internal status per file descriptor (FD)



# Binder Driver

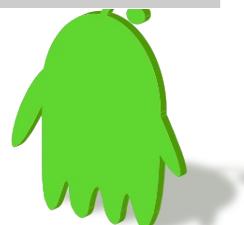


- A pool of threads is associated to each service application to process incoming IPC
- Binder performs mapping of object between two processes.
- Binder uses an object reference as an address in a process's memory space.
- Synchronous call, reference counting



# Binder is different from UNIX socket

|                        | socket           | binder  |
|------------------------|------------------|---|
| internal status        | associated to FD | associated to PID<br>(FD can be shared among threads in the same process) |
| read & write operation | stream I/O       | done at once by <b>ioctl</b>  |
| network transparency   | Yes              | No<br>expected local only   |



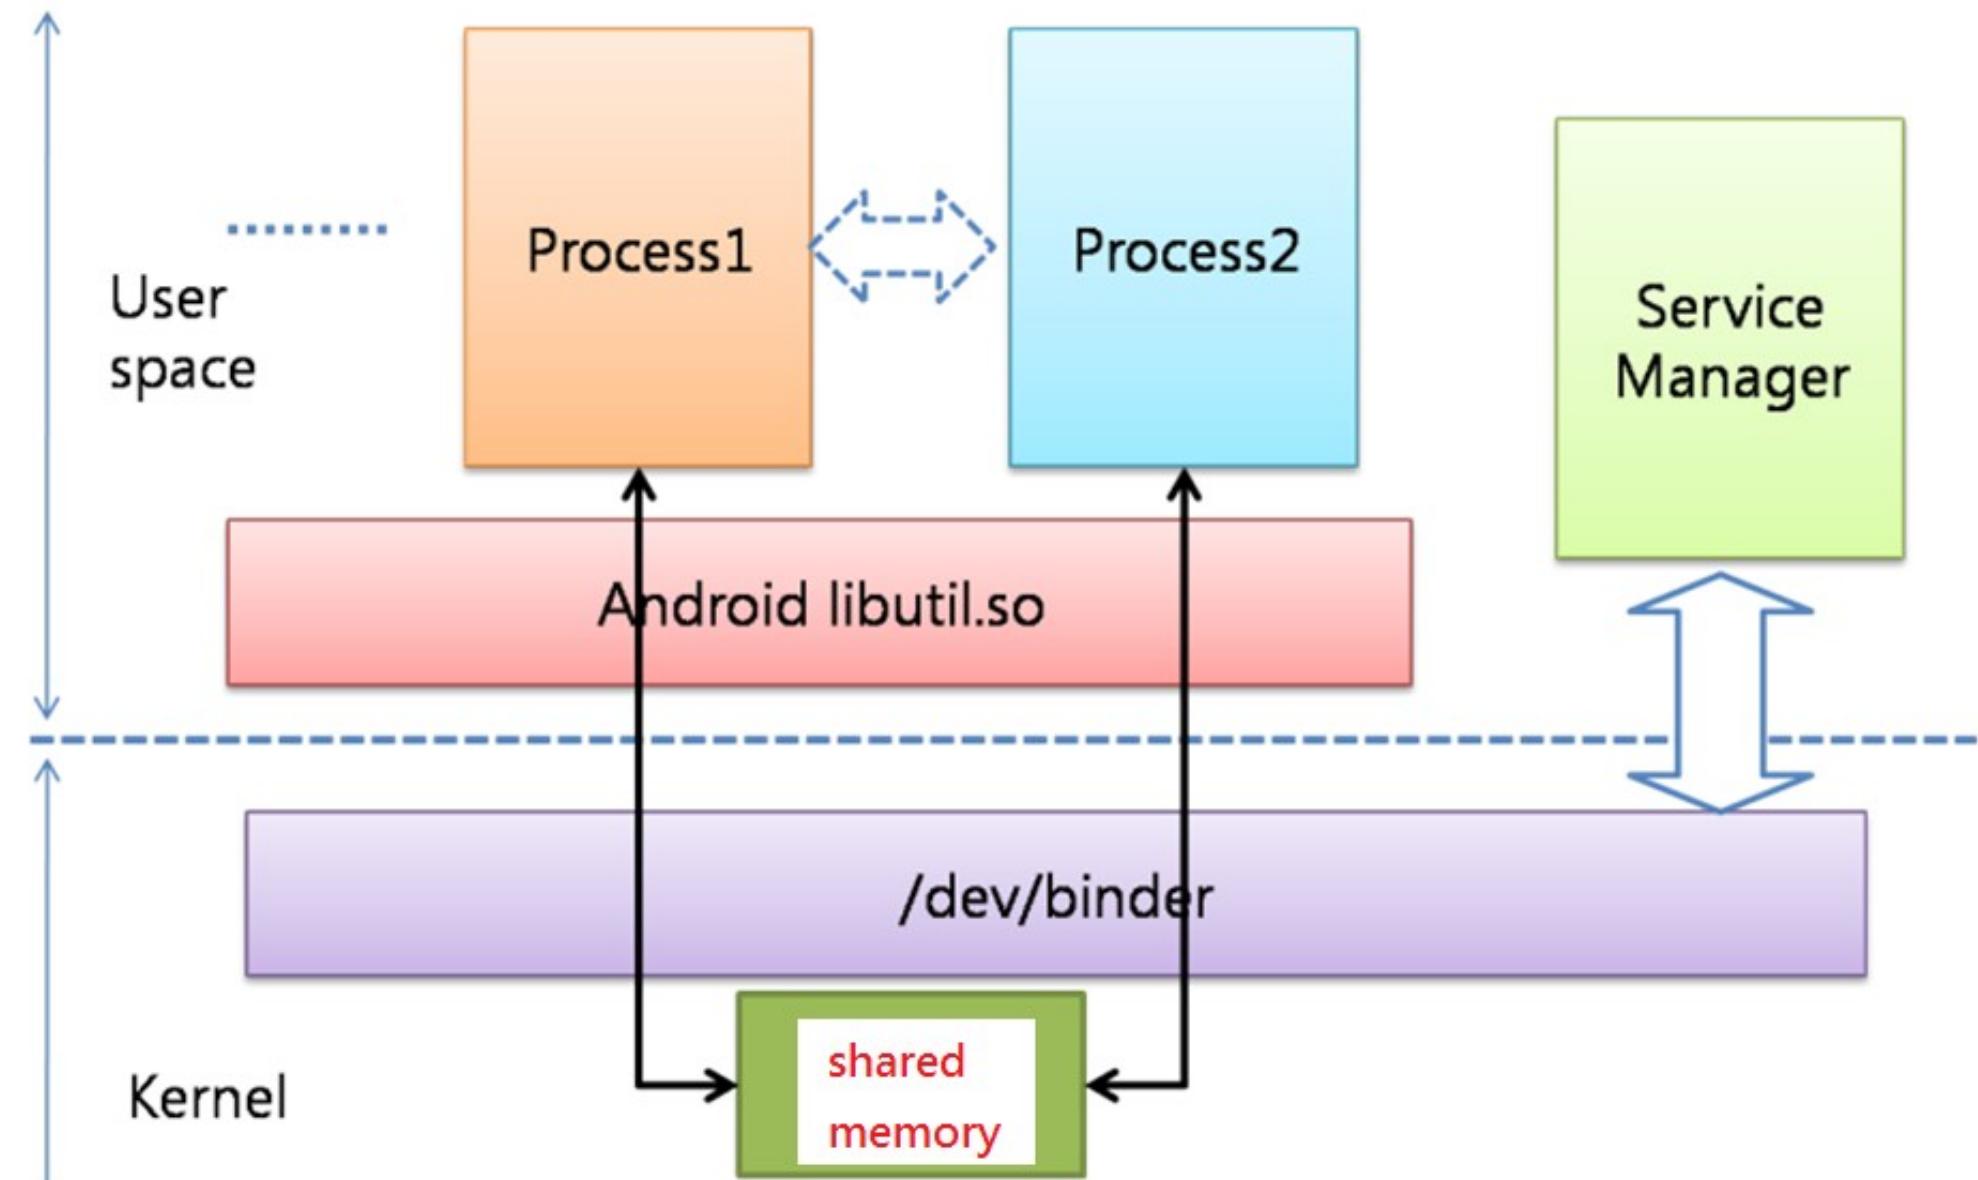
```
$ adb cat /sys/devices/virtual/misc/binder/uevent
```

Binder

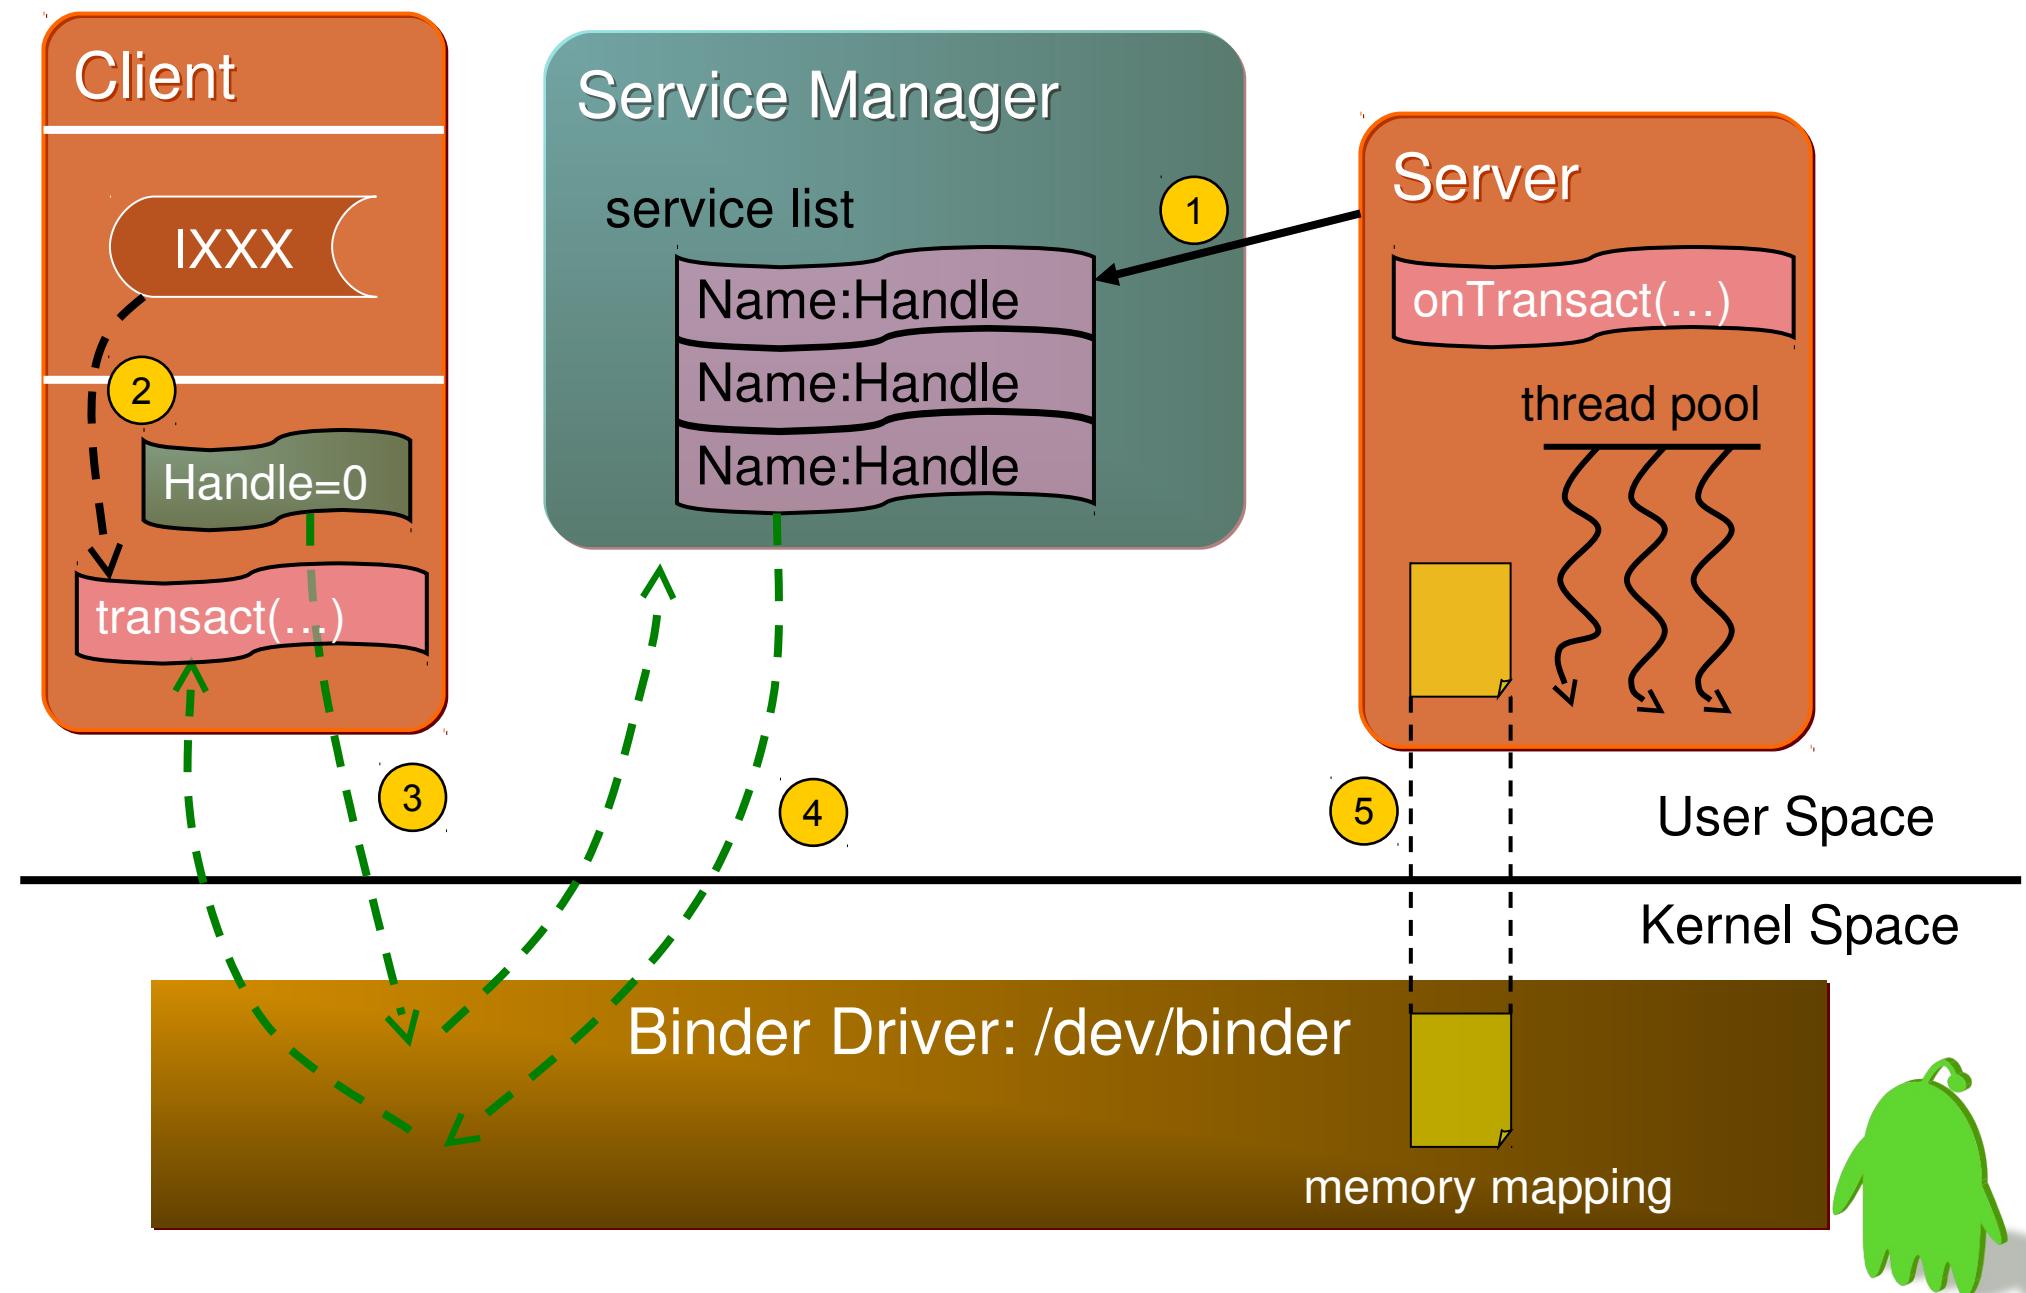
MAJOR=10

MINOR=47

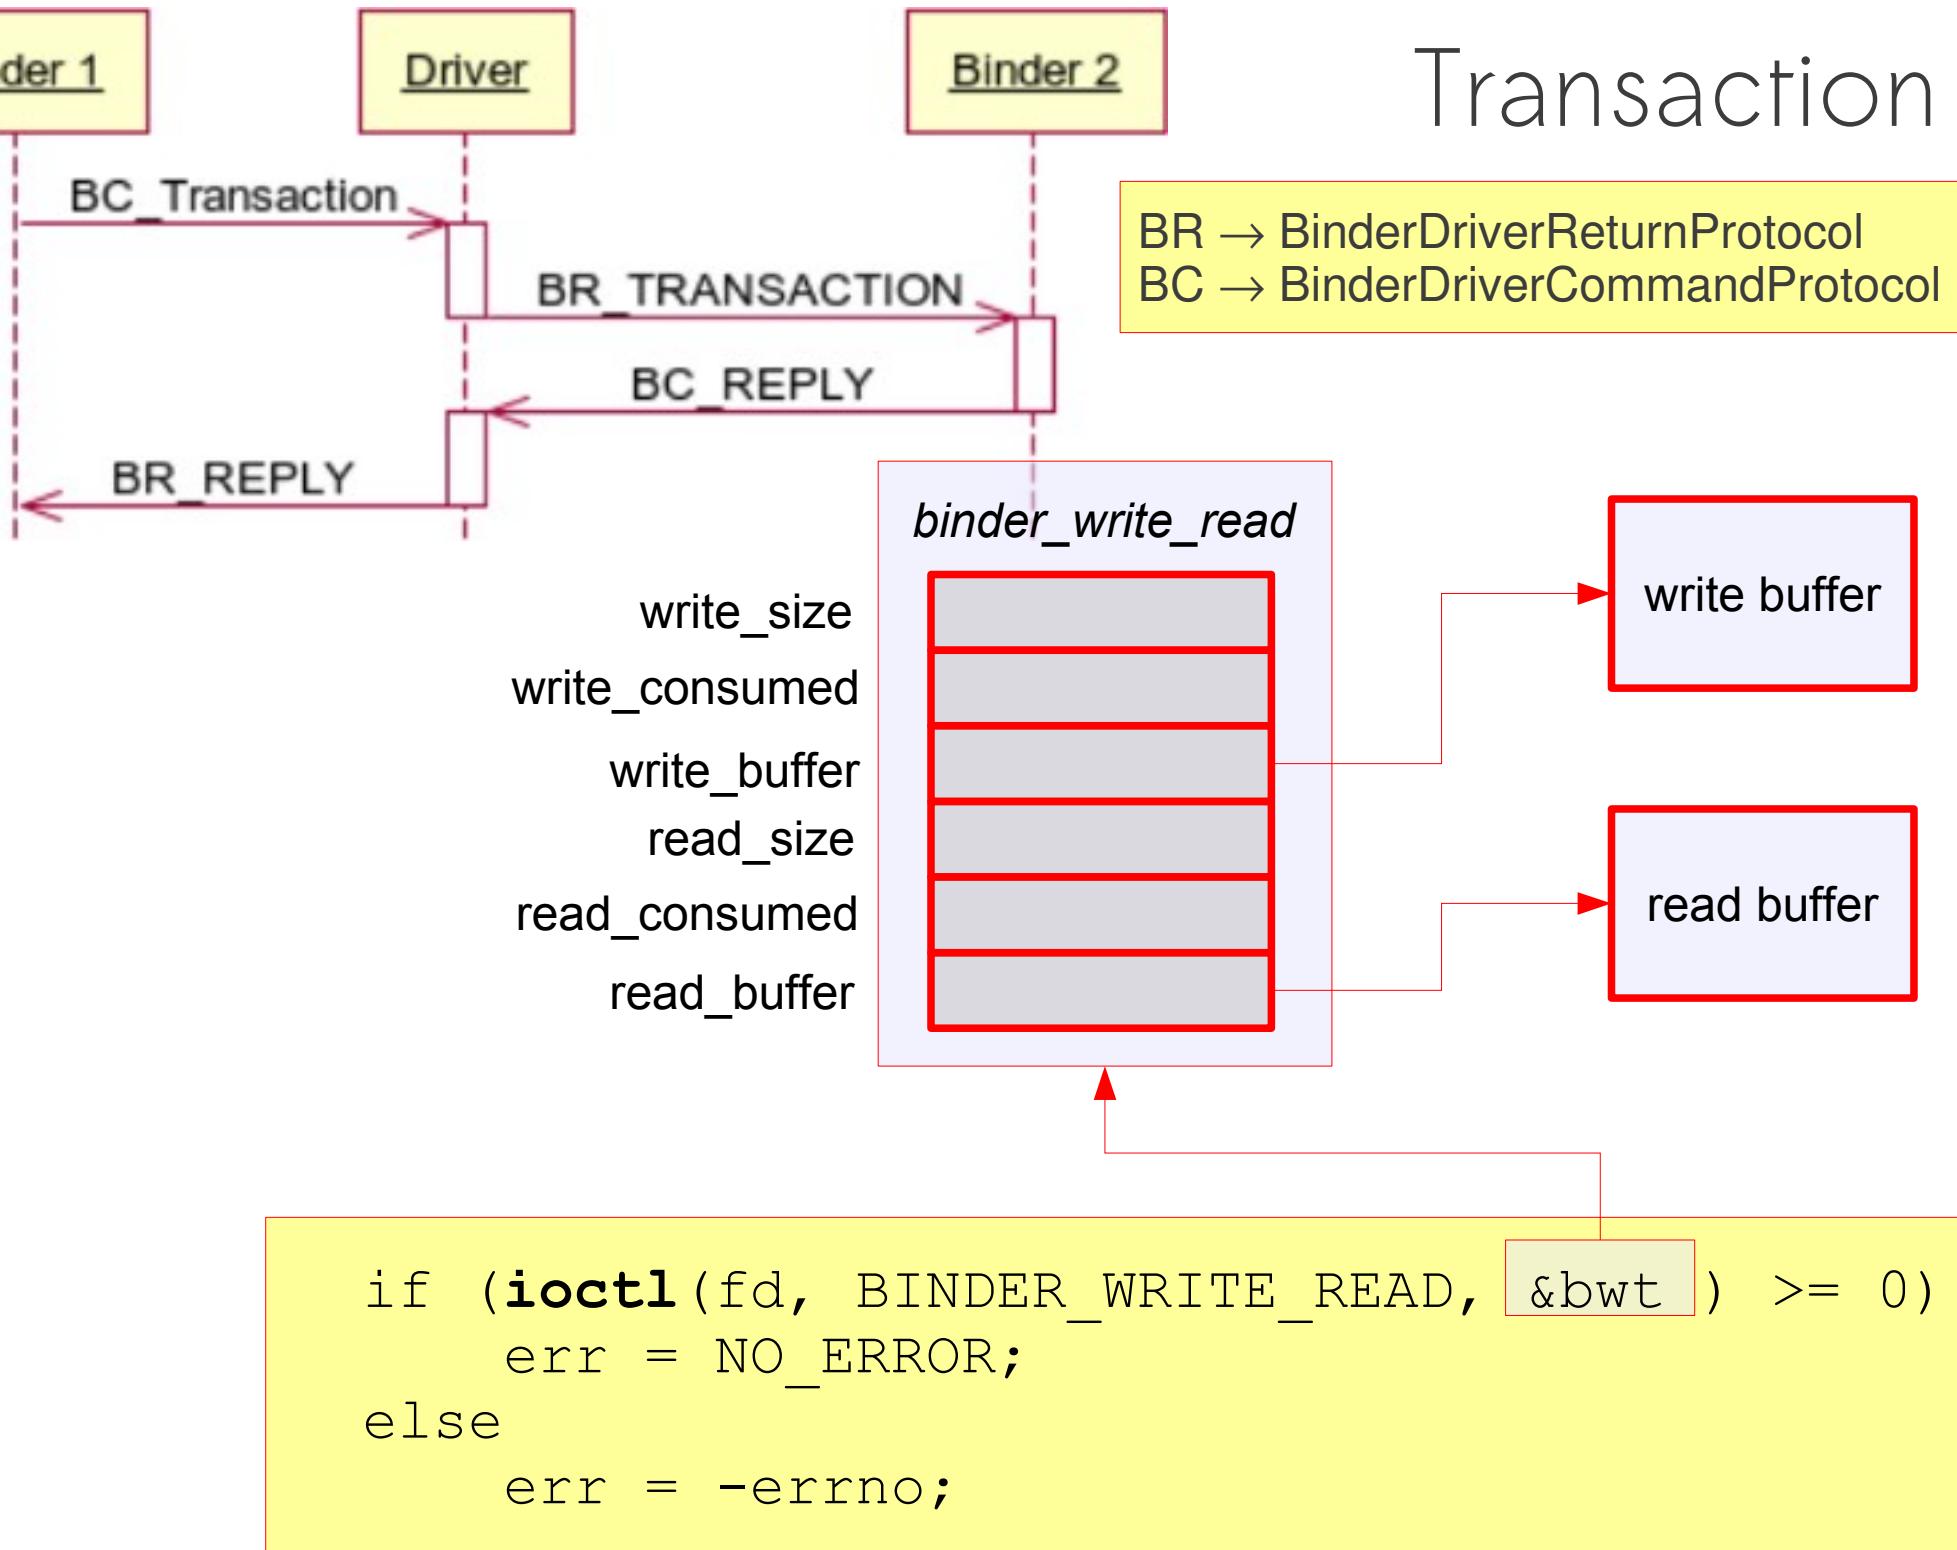
DEVNAME=binder



# from SM to Binder Driver

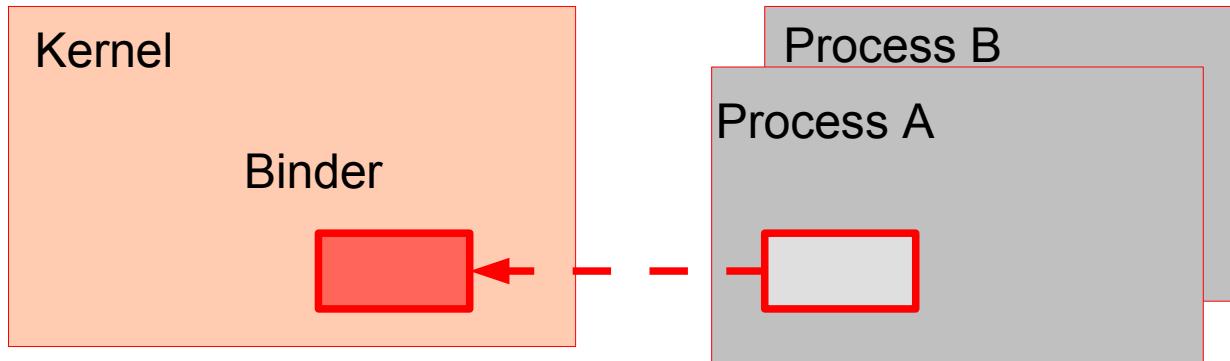


# Transaction

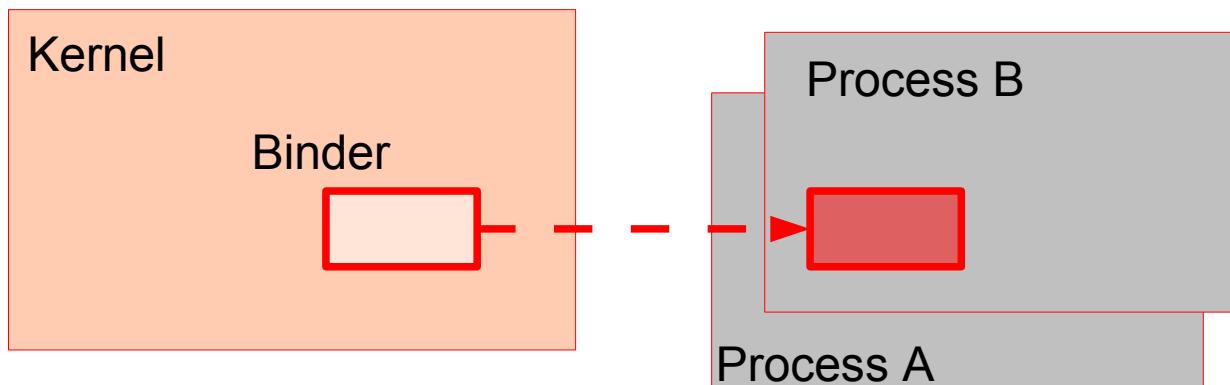


# Transaction of Binder

Process A and B have different memory space.  
They can not see each other.



Copy memory by **copy\_from\_user**  
Then, wake up process B



Copy memory by **copy\_to\_user**

Internally, Android uses Binder for graphics data transaction across processes.  
It is fairly efficient.



# Limitation of Binder IPC

- Binders are used to communicate over process boundaries since different processes don't share a common VM context
  - no more direct access to each others Objects (memory).
- Binders are not ideal for transferring large data streams (like audio/video) since every object has to be converted to (and back from) a Parcel.



# Binder Performance

- Good
  - Compact method index
  - Native binary marshalling
  - Support of ashmem shortcut
  - No GUID
- Bad
  - Dalvik Parcel overhead
  - ioctl() path is not optimal
  - Interface name overhead
  - Global lock



# Binder Security

- Binder's Security Features
  - Securely Determined Client Identity
    - Binder.getCallingUid(), Binder.getCallingPid()
    - Similar to Unix Domain Socket  
`getsockopt(..., SO_PEERCREDS, ...)`
  - Interface Reference Security
    - Client cannot guess Interface Reference
- Service Manager
  - Directory Service for System Services
- Server should check client permission

`Context.checkSelfPermission(permission, pid, uid)`



# Binder sample program

- Build binder benchmark program

```
cd system/extras/tests/binder/benchmarks
```

```
mm
```

```
adb push \  
    ../../../../out/target/product/crespo/data/nativebenchmark/binderAddInts \  
    /data/local/
```

- Execute

```
adb shell
```

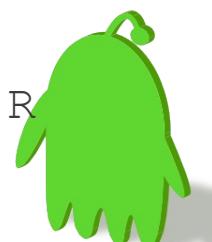
```
su
```

```
/data/local/binderAddInts -d 5 -n 5 &
```

```
ps
```

```
...
```

|      |                           |       |      |     |            |          |   |
|------|---------------------------|-------|------|-----|------------|----------|---|
| root | 17133                     | 16754 | 4568 | 860 | ffffffffff | 400e6284 | S |
|      | /data/local/binderAddInts |       |      |     |            |          |   |
| root | 17135                     | 17133 | 2520 | 616 | 00000000   | 400e5cb0 | R |
|      | /data/local/binderAddInts |       |      |     |            |          |   |



# Binder sample program

- Execute

```
/data/local/binderAddInts -d 5 -n 5 &
```

```
ps
```

```
...
```

```
root      17133 16754 4568    860      ffffffff 400e6284 S  
/data/local/binderAddInts
```

```
root      17135 17133 2520    616      00000000 400e5cb0 R  
/data/local/binderAddInts
```

```
cat /sys/kernel/debug/binder/transaction_log
```

```
transaction_log:3439847: call from 17133:17133 to 72:0 node  
1 handle 0 size 124:4
```

```
transaction_log:3439850: reply from 72:72 to 17133:17133 node  
0 handle 0 size 4:0
```

```
transaction_log:3439855: call from 17135:17135 to 17133:0  
node 3439848 handle 1 size 8:0
```

```
...
```



# Binder sysfs entries

- **adb shell ls /sys/kernel/debug/binder**

failed\_transaction\_log

proc

state

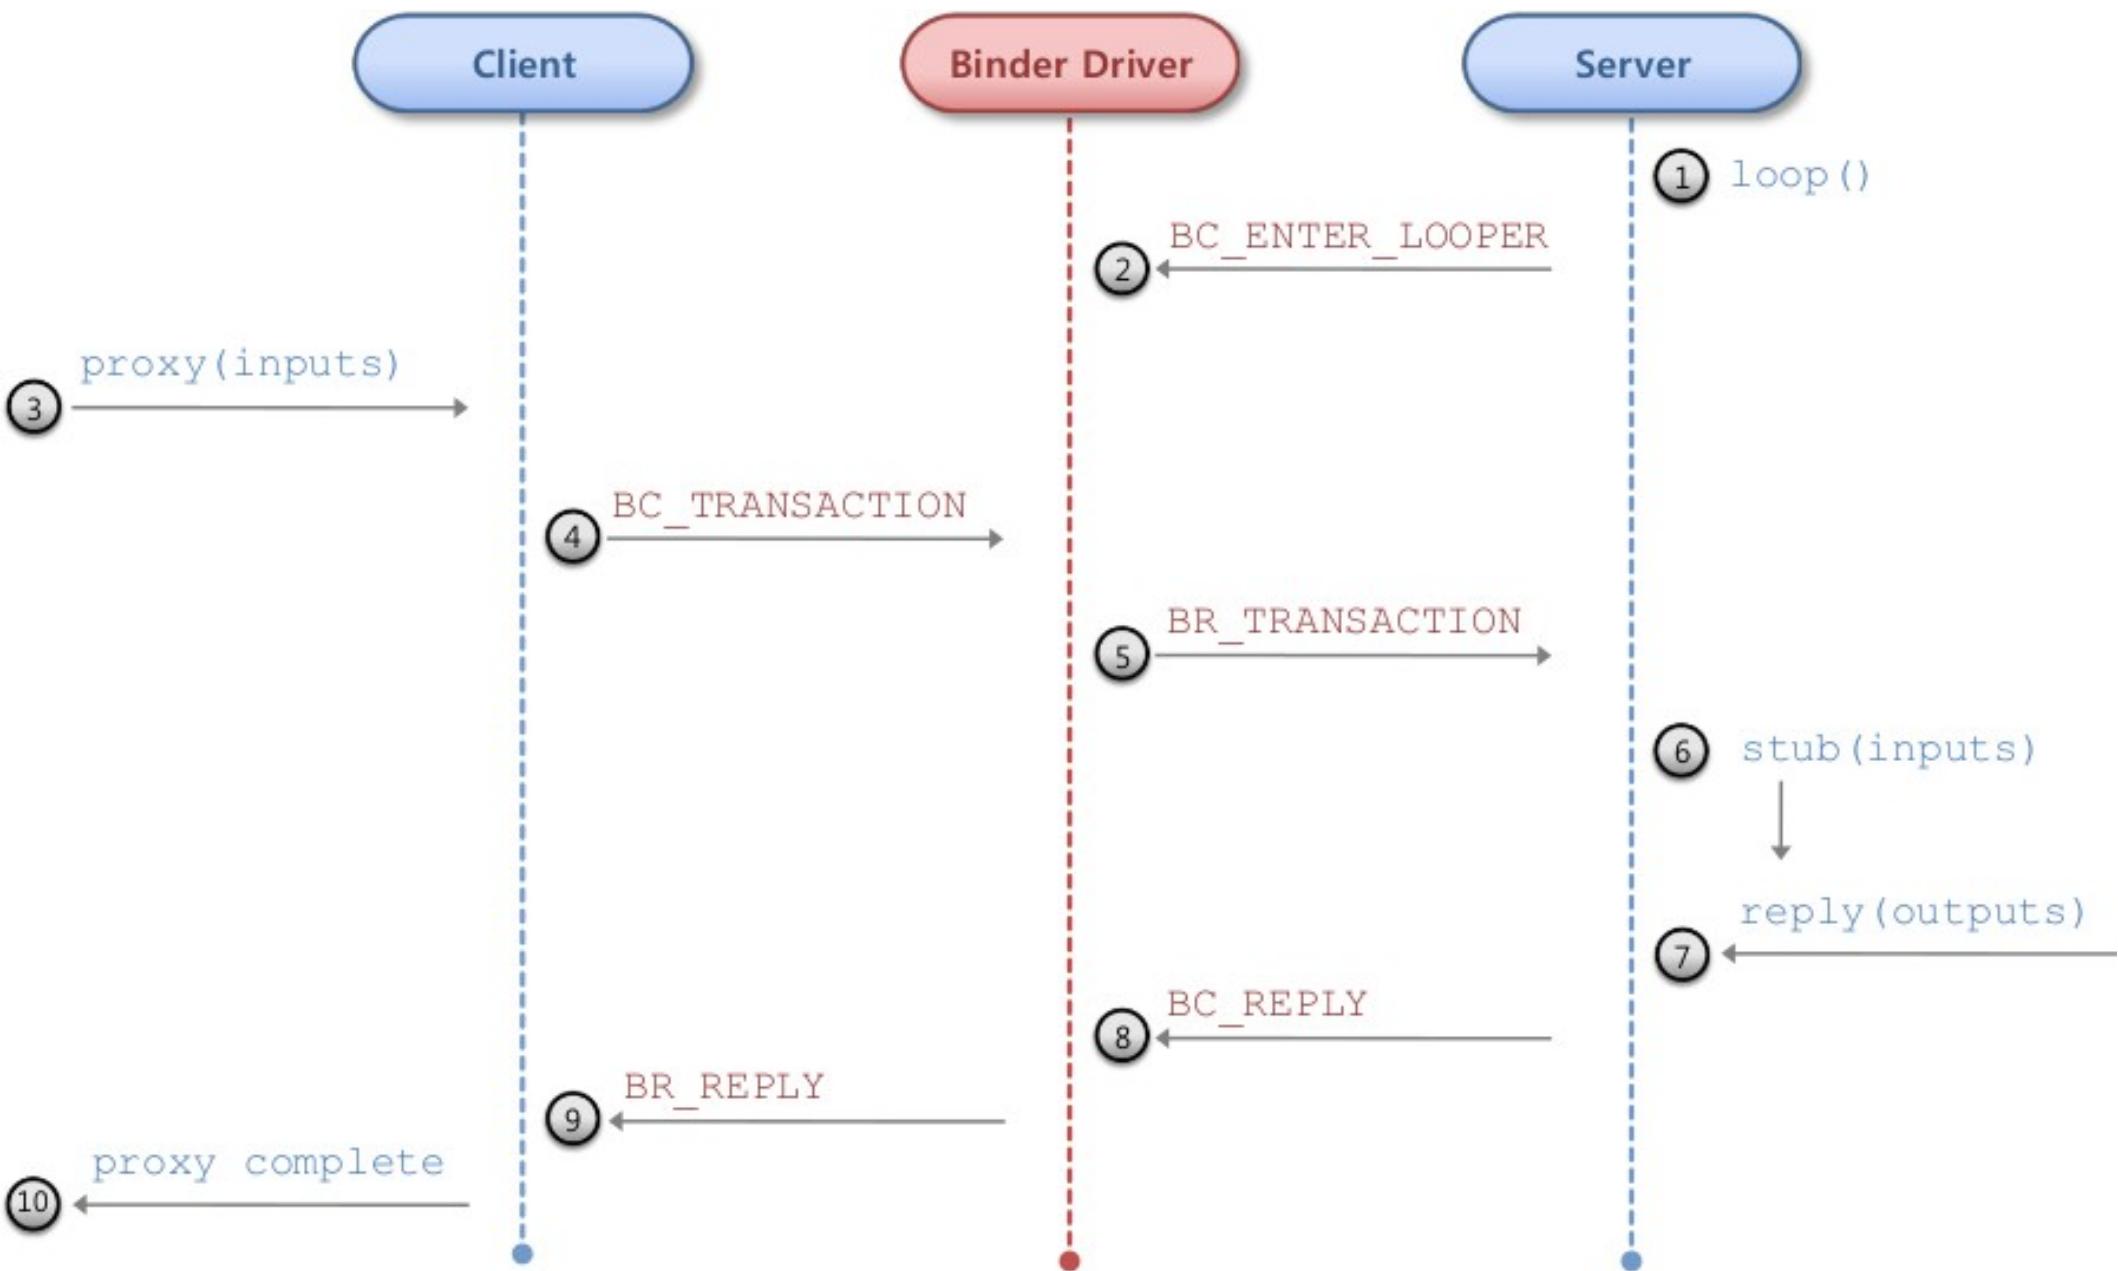
stats

transaction\_log

transactions



# Remote Procedure Call



# BINDER\_WRITE\_READ

```
struct binder_write_read {
    long write_size; /* bytes to write */
    long write_consumed; /* bytes consumed by driver */
    unsigned long write_buffer; /* pointer to write buffer */

    long read_size; /* bytes to read */
    long read_consumed; /* bytes consumed by driver */
    unsigned long read_buffer; /* pointer to read buffer */
};

#include <sys/ioctl.h>
#include <linux/binder.h>

int binder_write(int fd, void *data, long len) {
    struct binder_write_read bwr;

    bwr.write_size = len;
    bwr.write_consumed = 0;
    bwr.write_buffer = (unsigned) data;
    bwr.read_size = 0;
    bwr.read_consumed = 0;
    bwr.read_buffer = 0;
    return ioctl(fd, BINDER_WRITE_READ, &bwr);
}
```

The diagram illustrates the mapping of the `binder_write_read` structure fields to the parameters passed to the `ioctl` function. The `binder_write` function creates a `binder_write_read` structure and passes it to `ioctl` with the command `BINDER_WRITE_READ`. The structure fields are mapped as follows:

- `write_size` and `write_consumed` are mapped to the first `BC_*` parameter.
- `read_size` and `read_consumed` are mapped to the first `BR_*` parameter.



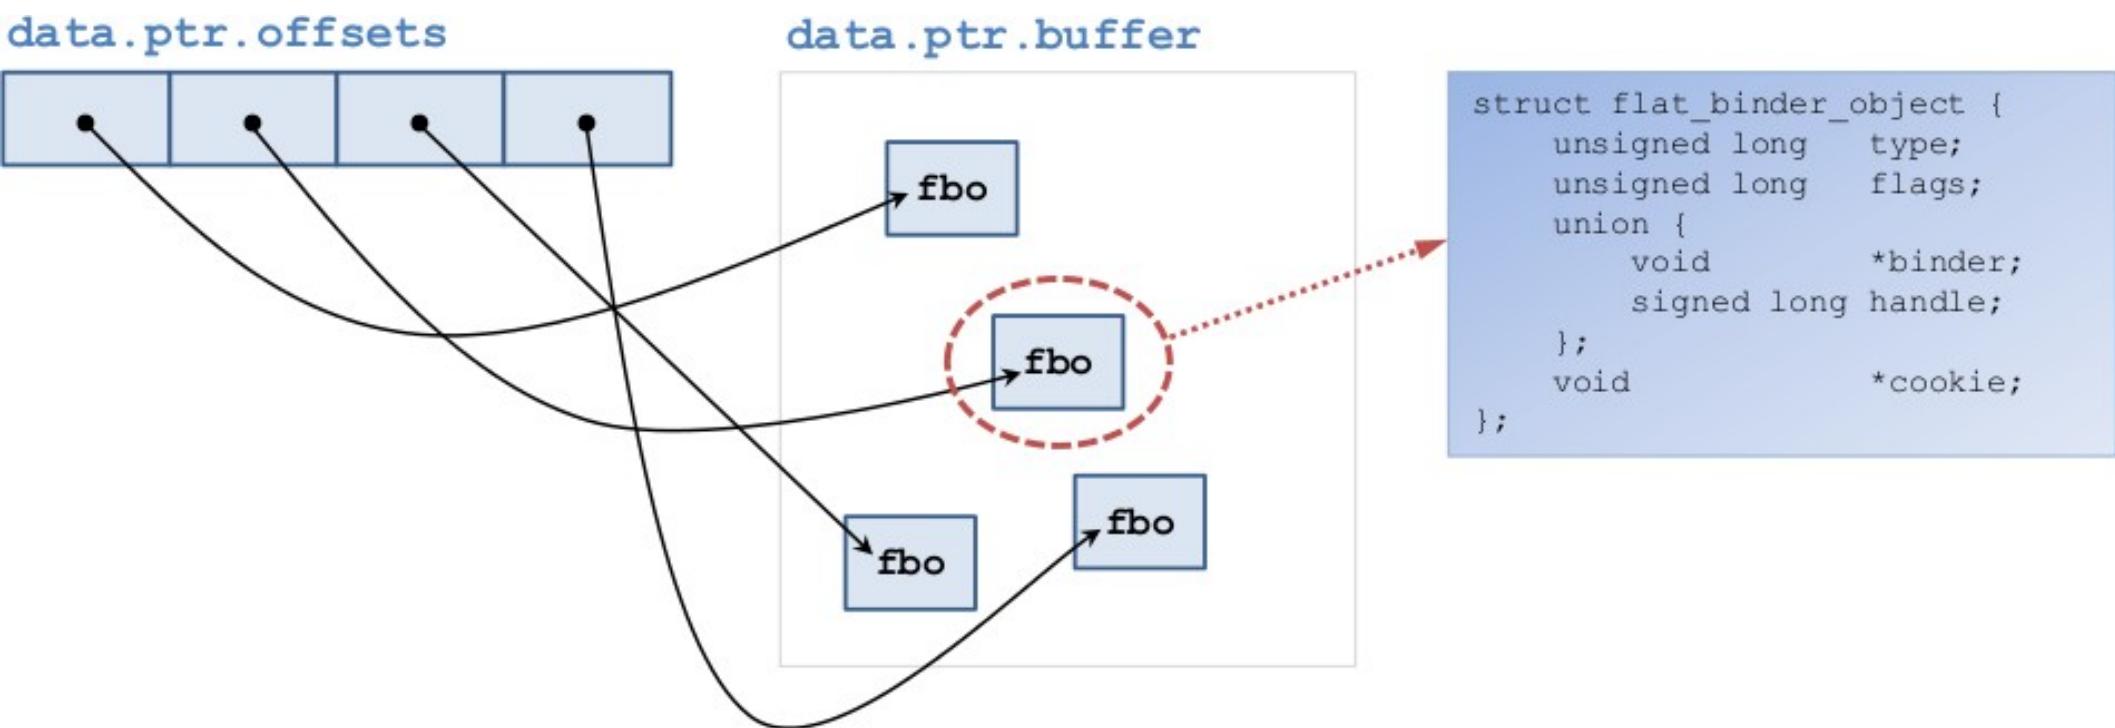
- Target Method
  - handle : Remote Interface
  - ptr & cookie : Local Interface
  - code : Method ID
- Parcel - Input/Output Parameters
  - data.ptr.buffer
  - data\_size
- Object Reference Management
  - data.ptr.offsets
  - offsets\_size
- Security
  - sender\_pid
  - sender\_euid
- No Transaction GUID
  - Transparent Recursion

## Binder Transaction

```
#define BC_TRANSACTION
#define BC_REPLY
#define BR_TRANSACTION
#define BR_REPLY

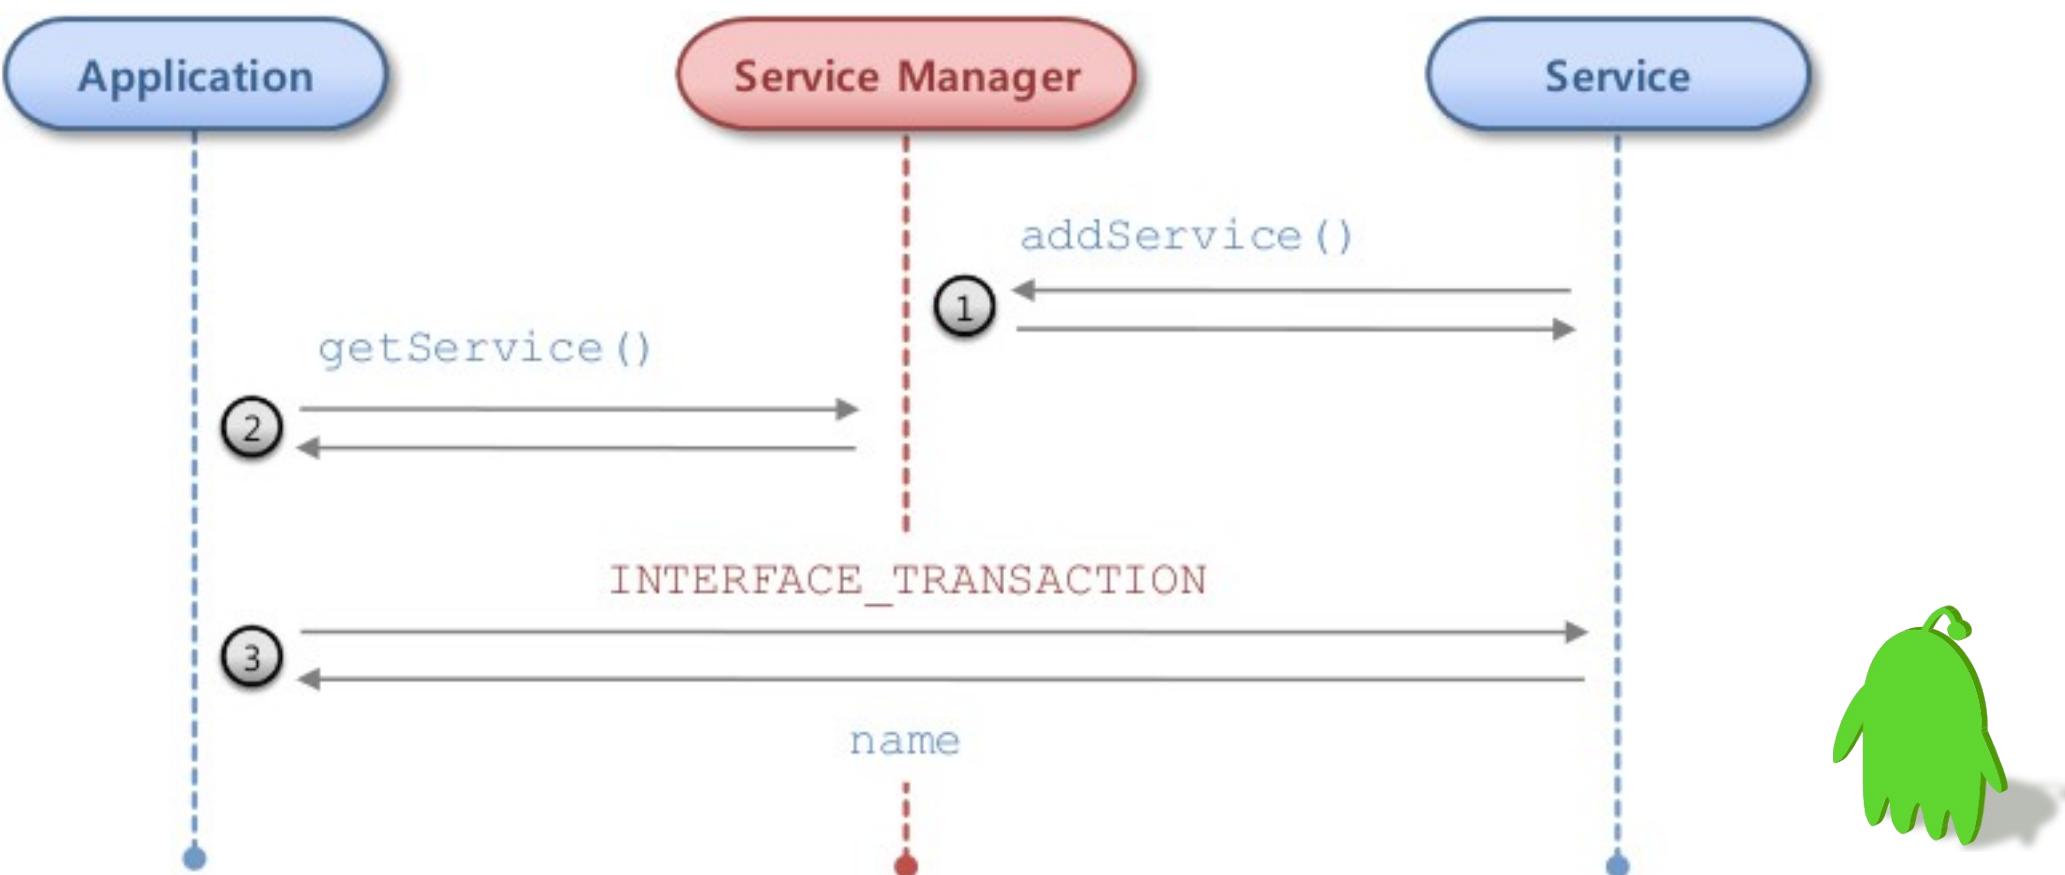
struct binder_transaction_data {
    union {
        size_t          handle;
        void*          *ptr;
    } target;
    void            *cookie;
    unsigned int     code;
    unsigned int     flags;
    pid_t           sender_pid;
    uid_t           sender_euid;
    size_t          data_size;
    size_t          offsets_size;
    union {
        struct {
            const void*   buffer;
            const void*   *offsets;
        } ptr;
        uint8_t        buf[8];
    } data;
};
```

# Object Reference Management



# Service Registration and Discovery

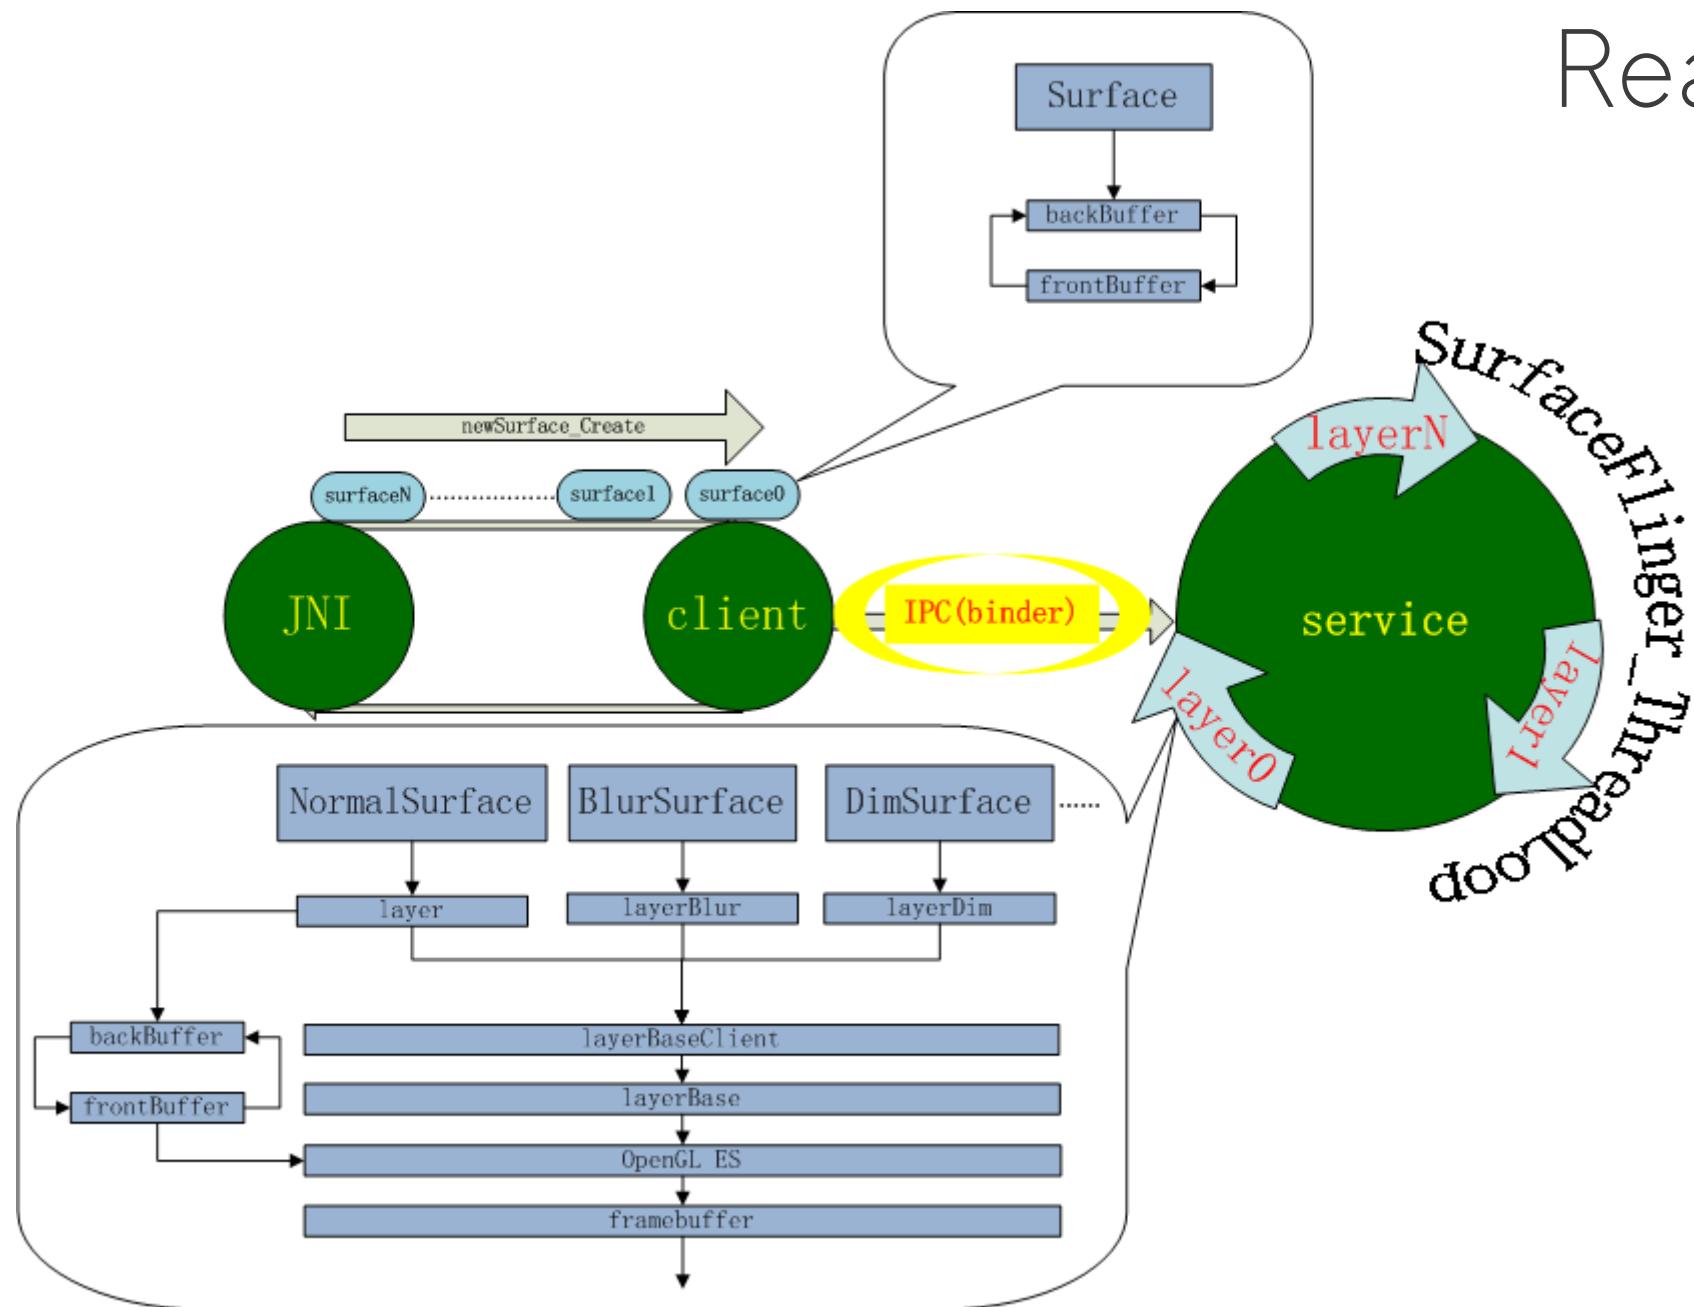
- System service is executed by `IServiceManager::addService()` calls.
  - Parameter: handle to Binder Driver
- Look up the name of specific service in Binder Driver Map
  - `IServiceManager::getService()` returns the handle of the found registered services
- `android.os.IBinder.INTERFACE_TRANSACTION`: the actual name



# Binder use case: Android Graphics



# Real Case



Binder IPC is used for communicating between Graphics client and server.  
Taken from <http://www.cnblogs.com/xl19862005/archive/2011/11/17/2215363.html>



# Surface

Source: frameworks/base/core/java/android/view/Surface.java

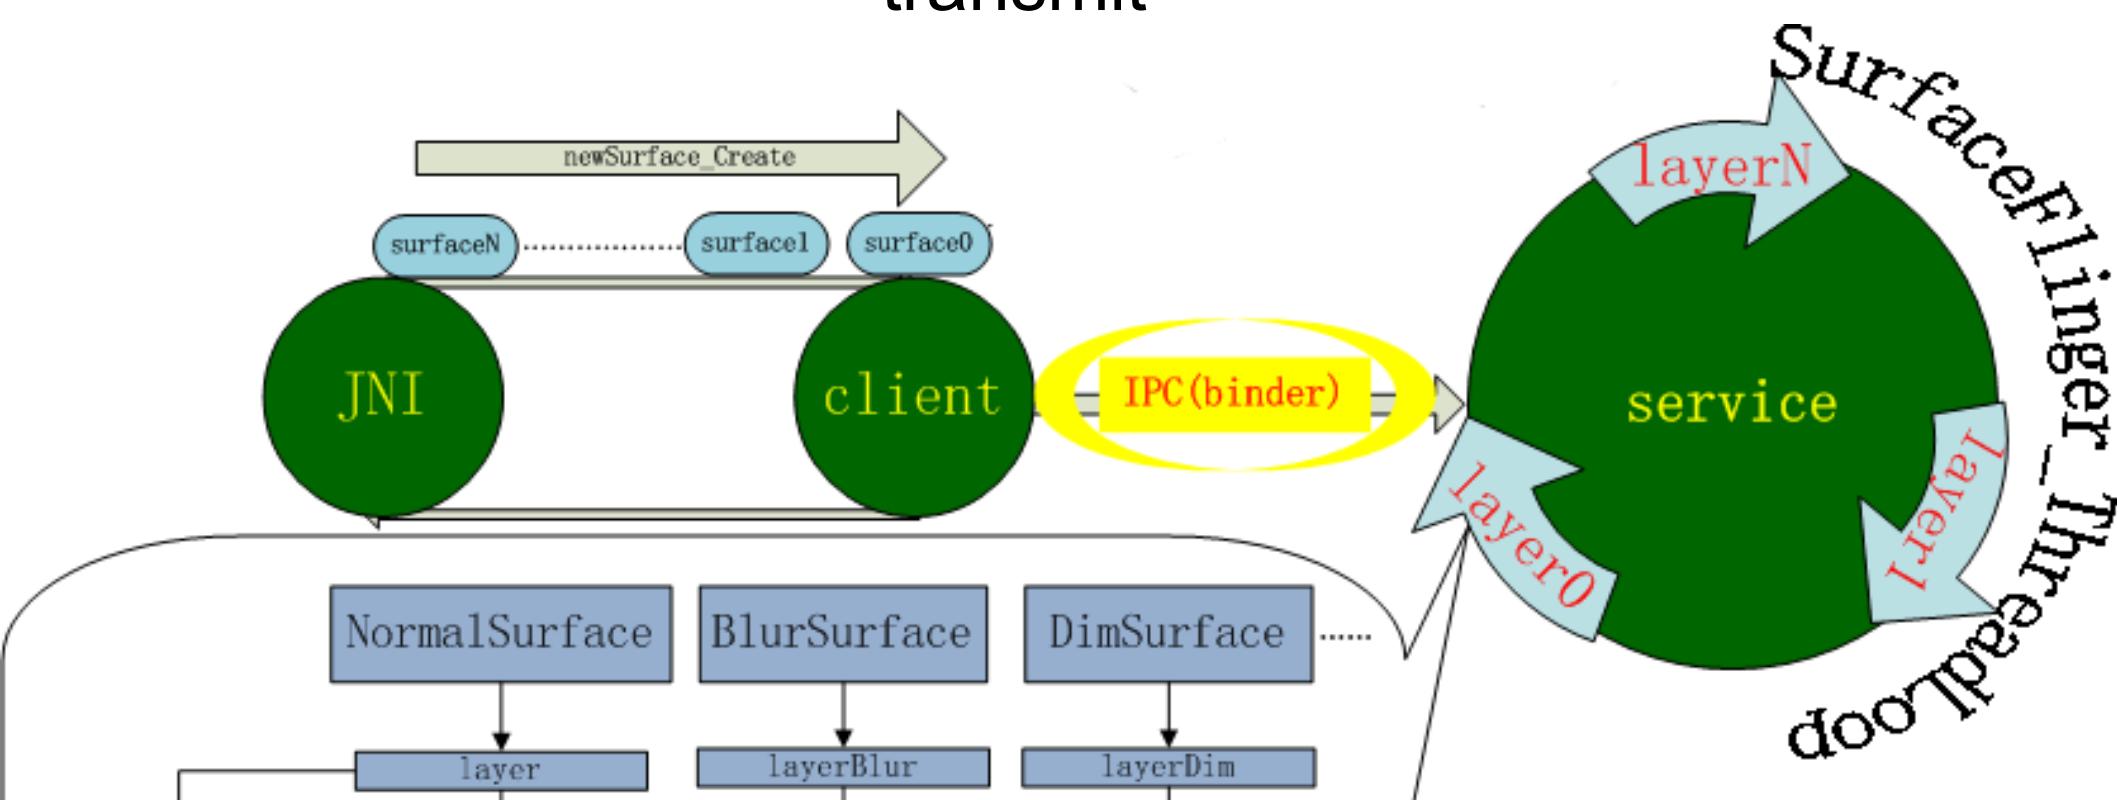
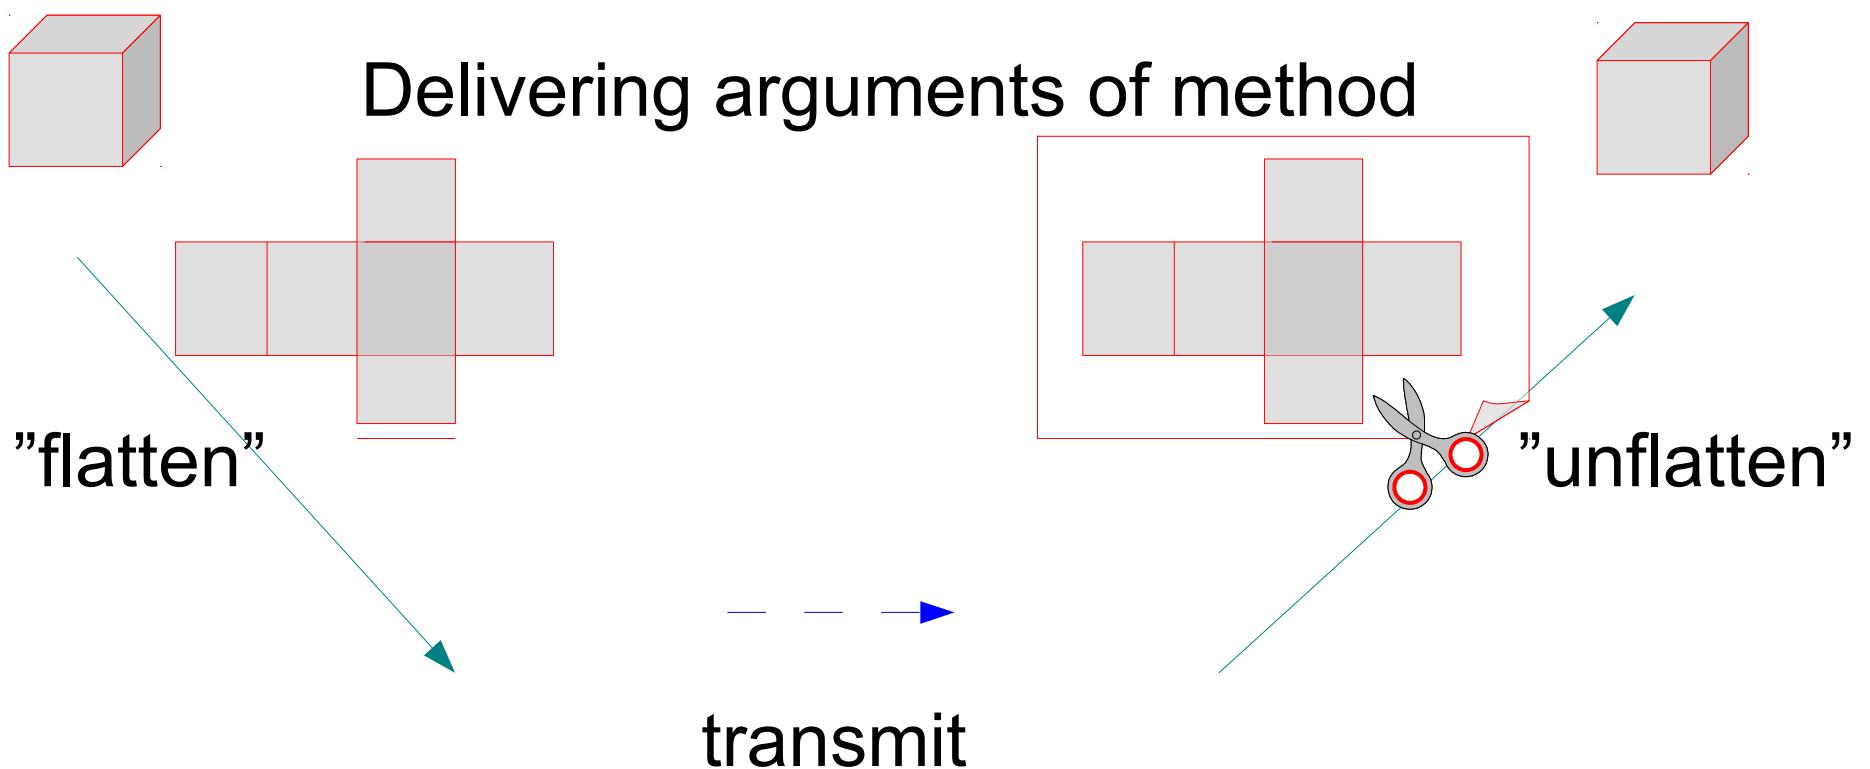
- **/\* Handle on to a raw buffer that is being managed by the screen compositor \*/**

```
public class Surface implements Parcelable {  
    public Surface() {  
        mCanvas = new CompatibleCanvas();  
    }  
    private class CompatibleCanvas  
        extends Canvas { /* ... */ }  
}
```

Surface instances can be written to and restored from a Parcel.



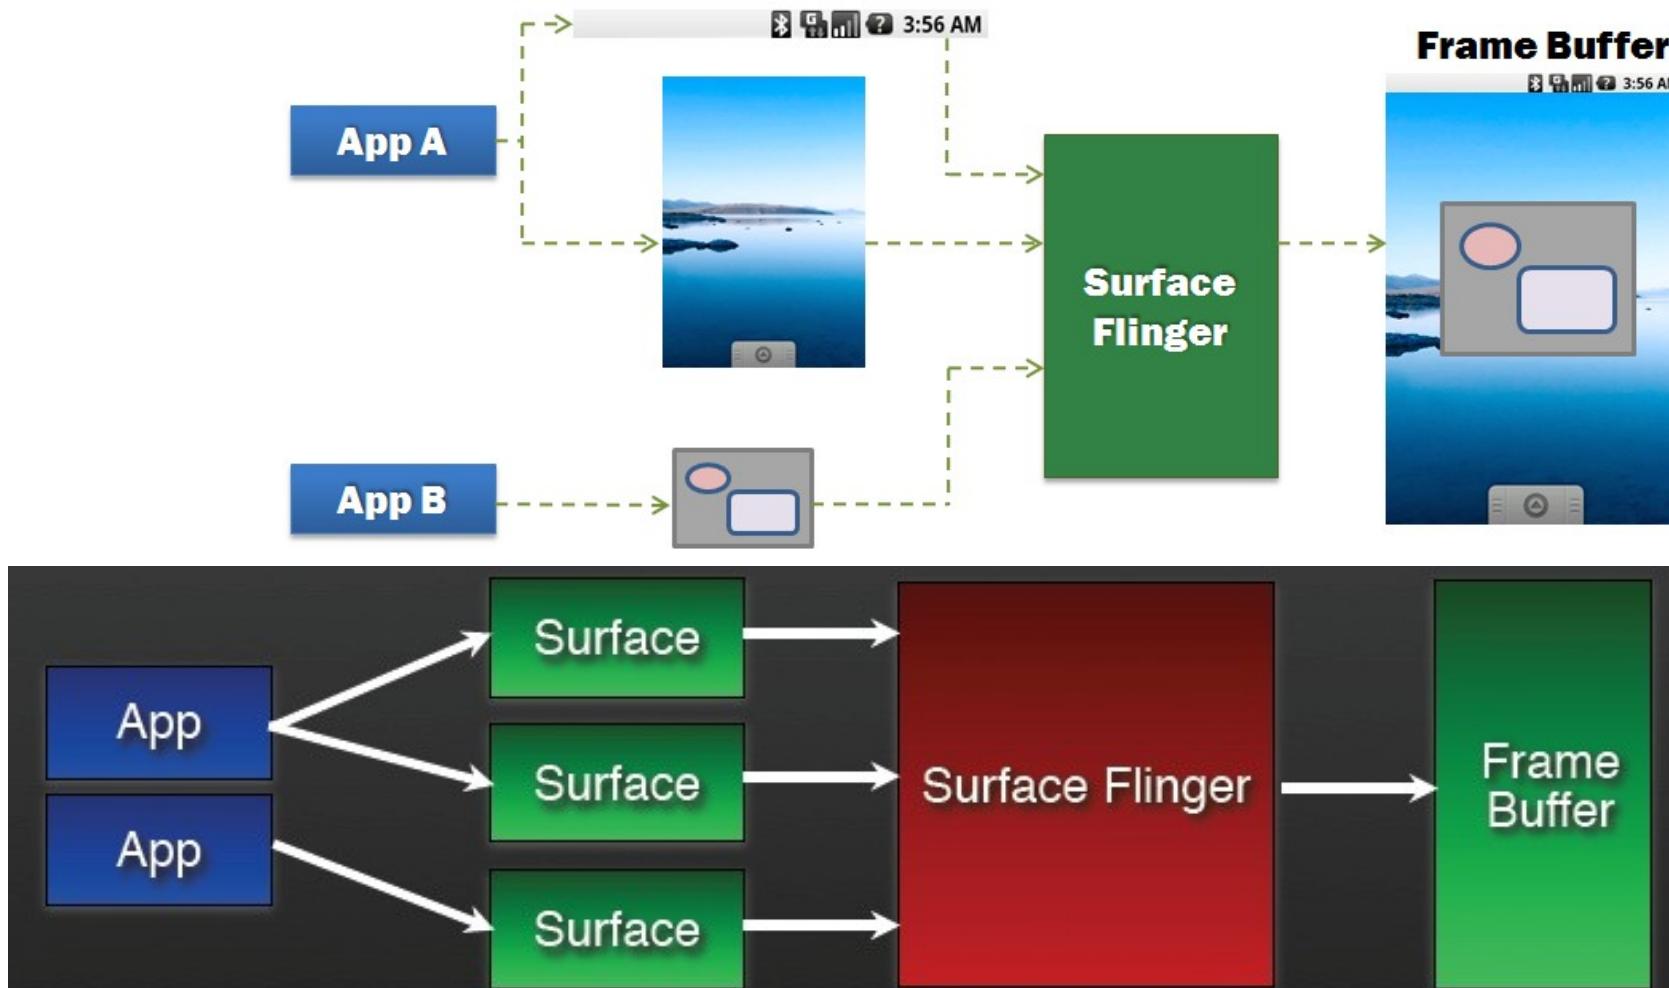
# Delivering arguments of method



# Android SurfaceFlinger

- Properties

- Can combine 2D/3D surfaces and surfaces from multiple applications
- Surfaces passed as buffers via Binder IPC calls
- Can use OpenGL ES and 2D hardware accelerator for its compositions
  - Double-buffering using page-flip



## System Server Process

Surface Flinger Service

CopyBit HAL

EGL / OpenGL API

Everything is  
around Binder

Linux Kernel

Driver

Driver

MALI Driver

## Application Process

Still Capture Application

Dalvik VM

Display JNI

OpenGL / EGL  
JNI

EGL / OpenGL  
API

B  
I  
N  
D  
E  
R  
I  
P  
C



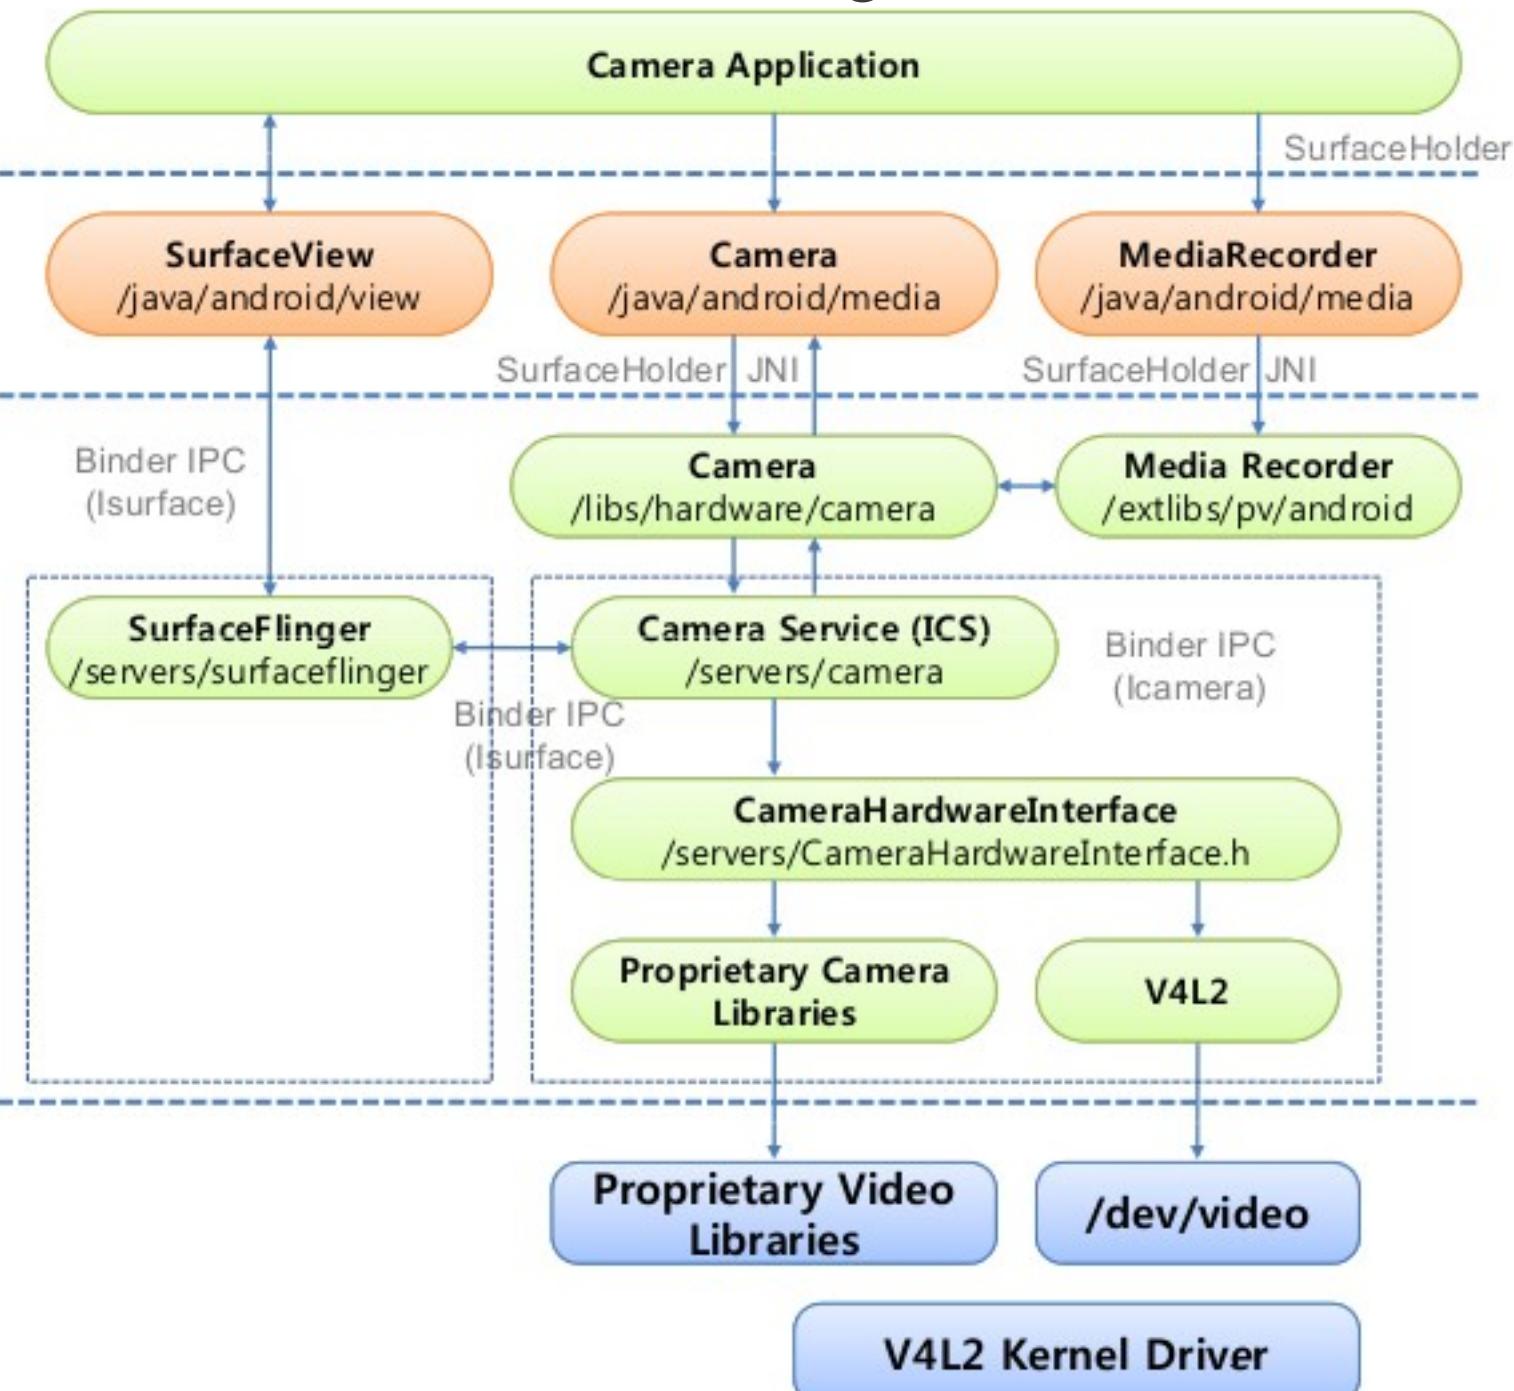
# Camera + SurfaceFlinger + Binder

Applications

Applications Framework

Native Libraries  
(user space HAL)

Linux Kernel



# Reference

- Inter-process communication of Android, Tetsuyuki Kobayashi
- 淺談 Android 系統進程間通信（IPC）機制 Binder 中的 Server 和 Client 獲得 Service Manager 接口之路  
<http://blog.goggb.com/?post=1580>
- Android Binder – Android Interprocess Communication, Thorsten Schreiber
- Design Patterns in the Android Framework, Prof. Sheng-De Wang





<http://0xlab.org>